

HBase Tutorial



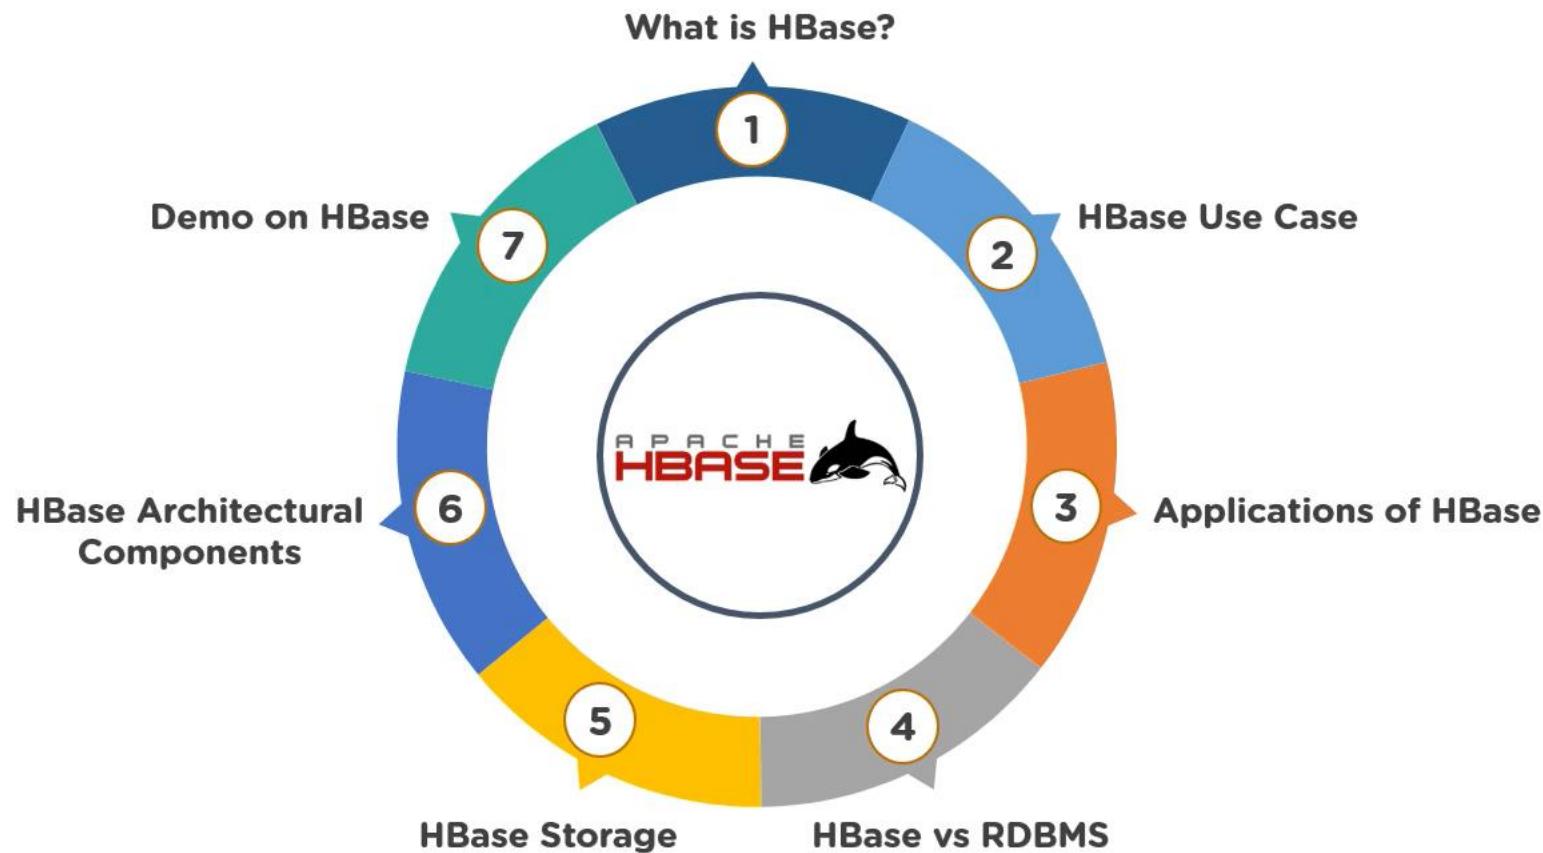
NoSQL

Database System



Bulk Data Load

What's in it for you?



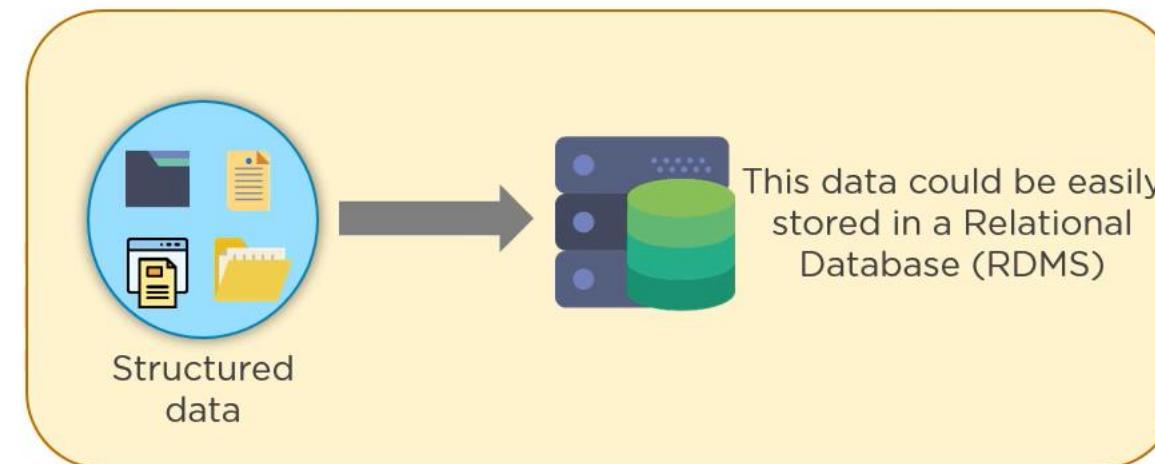
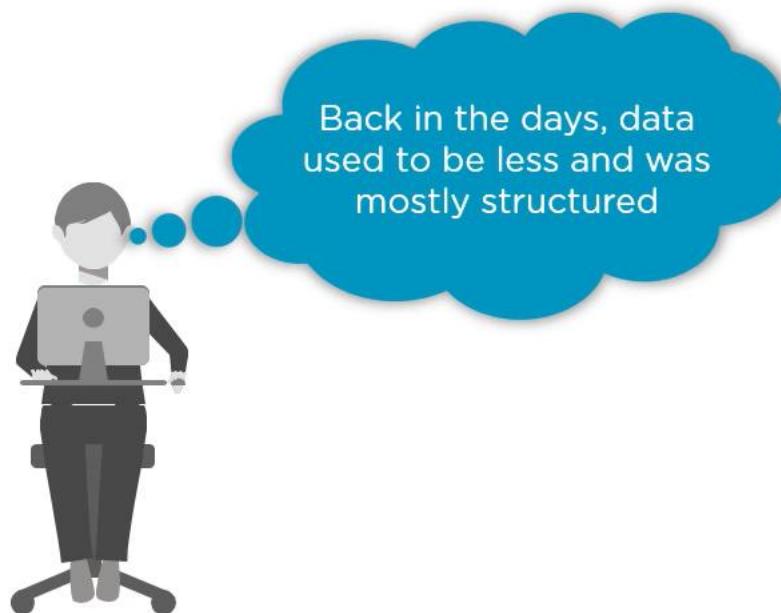
Introduction to HBase

NoSQL

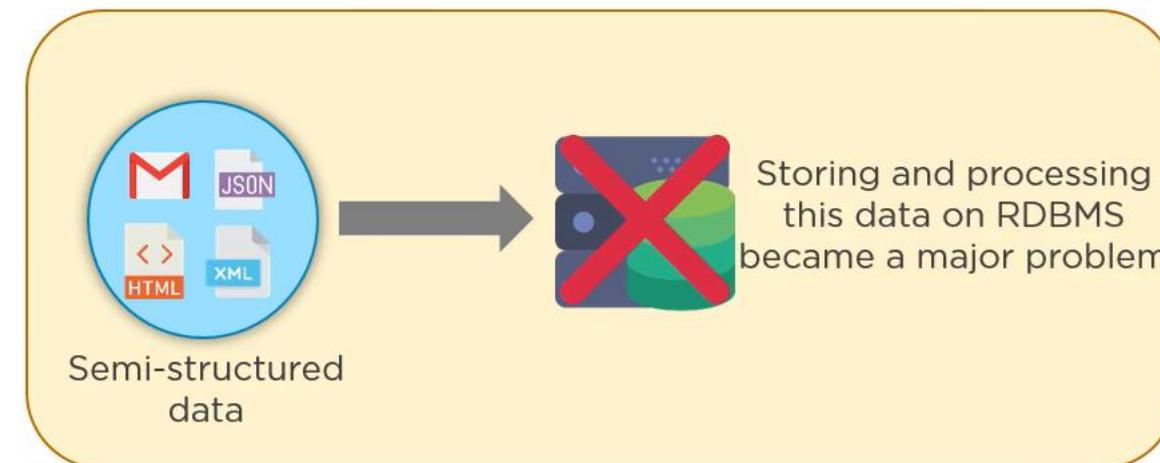
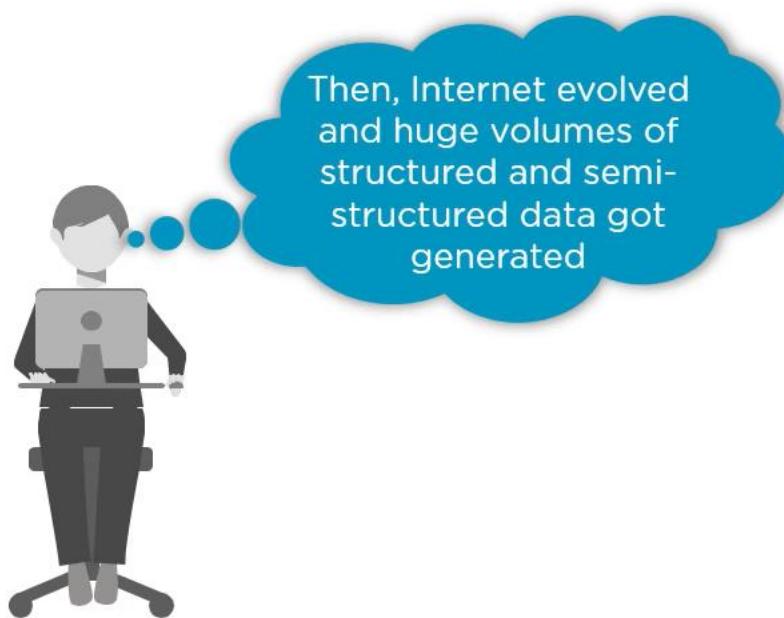
Database System



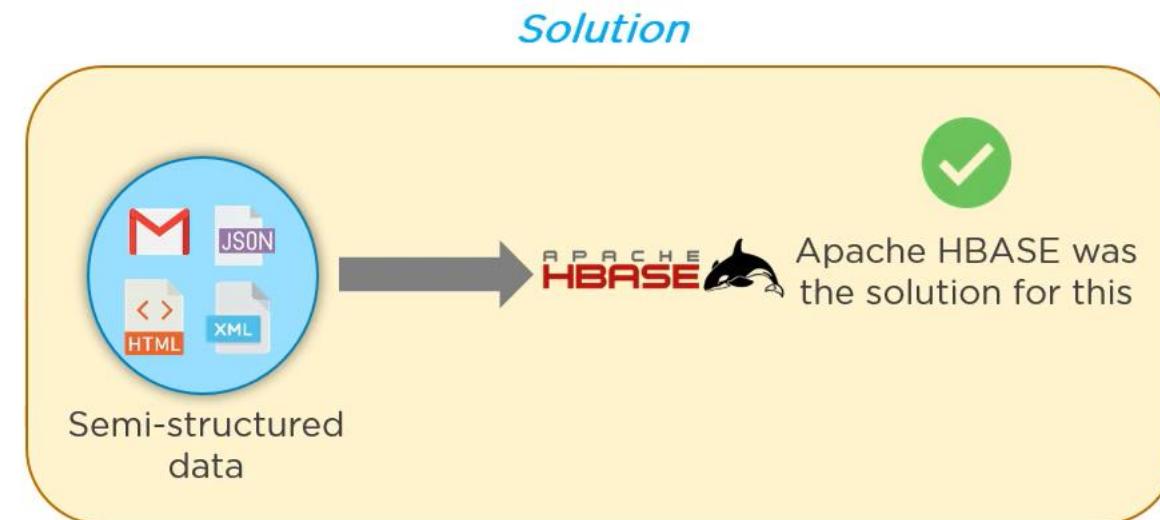
Introduction to HBase



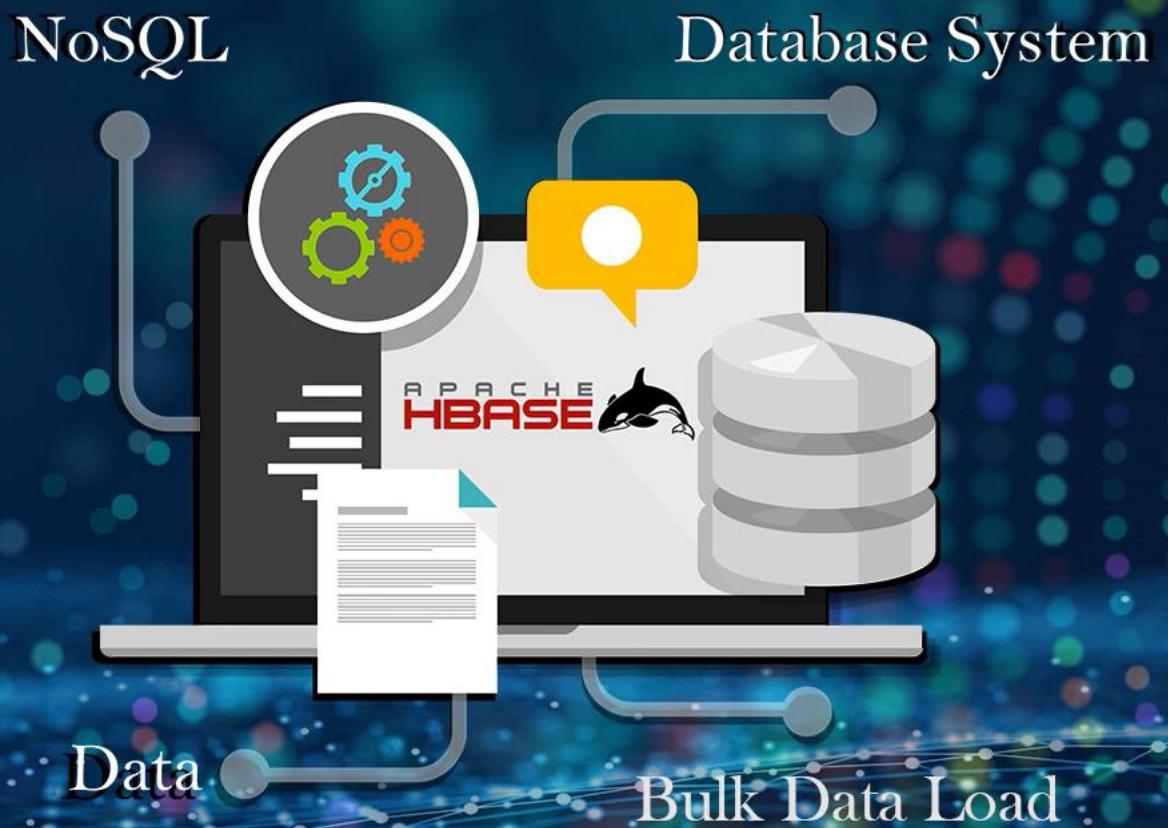
Introduction to HBase



Introduction to HBase



HBase History



HBase History

Nov 2006

Google released the paper on BigTable

1

Feb 2007

HBase prototype was created as a Hadoop contribution

2

Oct 2007

First usable HBase along with Hadoop 0.15.0 was released

3

Jan 2008

HBase became the subproject of Hadoop

4

Oct 2008 - Sep 2009

HBase 0.81.1, 0.19.0 and 0.20.0 was released between Oct 2008 - Sep 2009

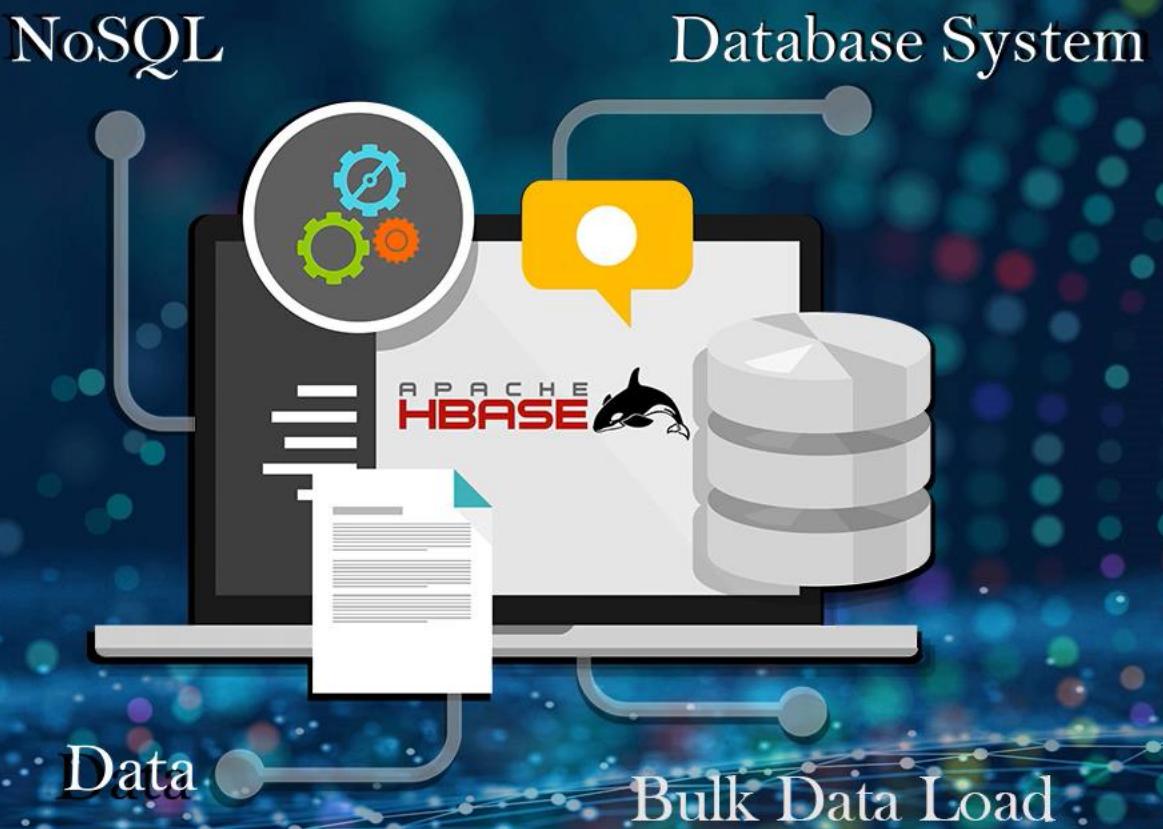
5

May 2010

HBase became Apache top-level project

6

What is HBase?



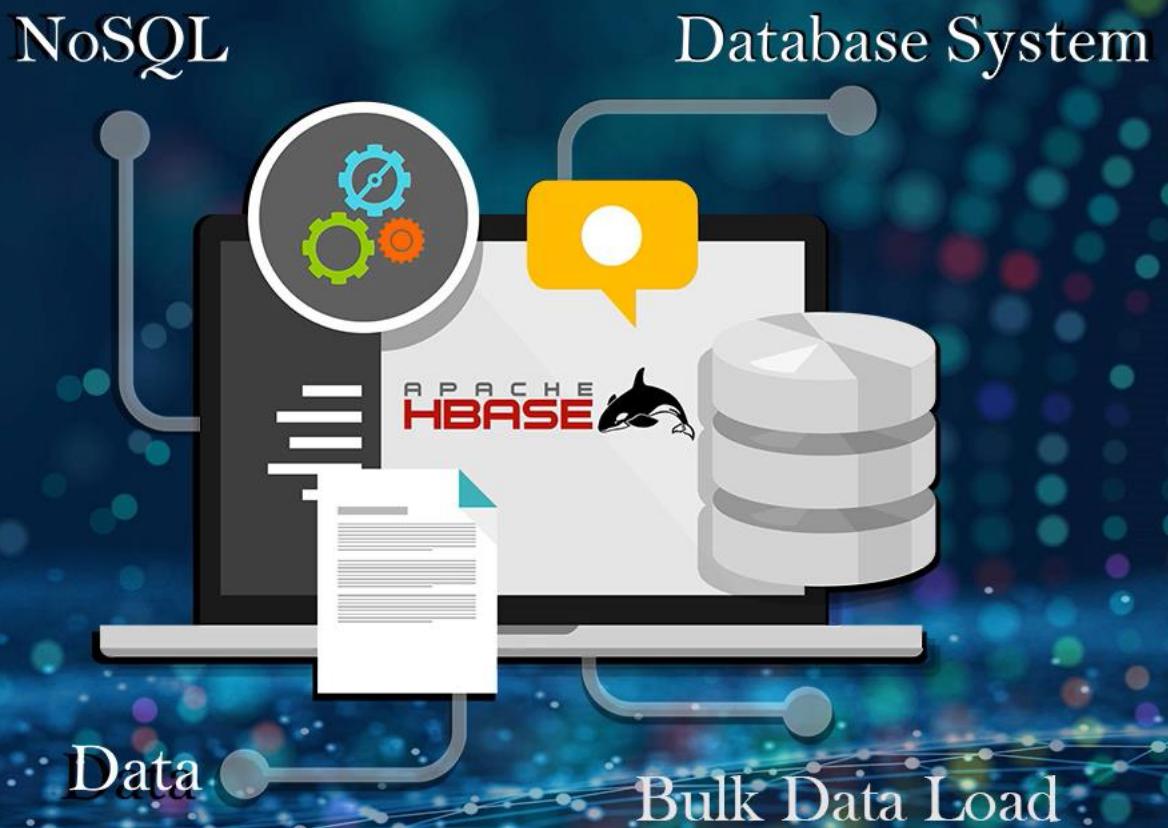
What is HBase?



HBase is a column oriented database management system derived from Google's NoSQL database [BigTable](#) that runs on top of HDFS

- 1 Open source project that is horizontally scalable
- 2 NoSQL database written in JAVA which performs faster querying
- 3 Well suited for sparse data sets
(can contain missing or NA values)

Companies using HBase



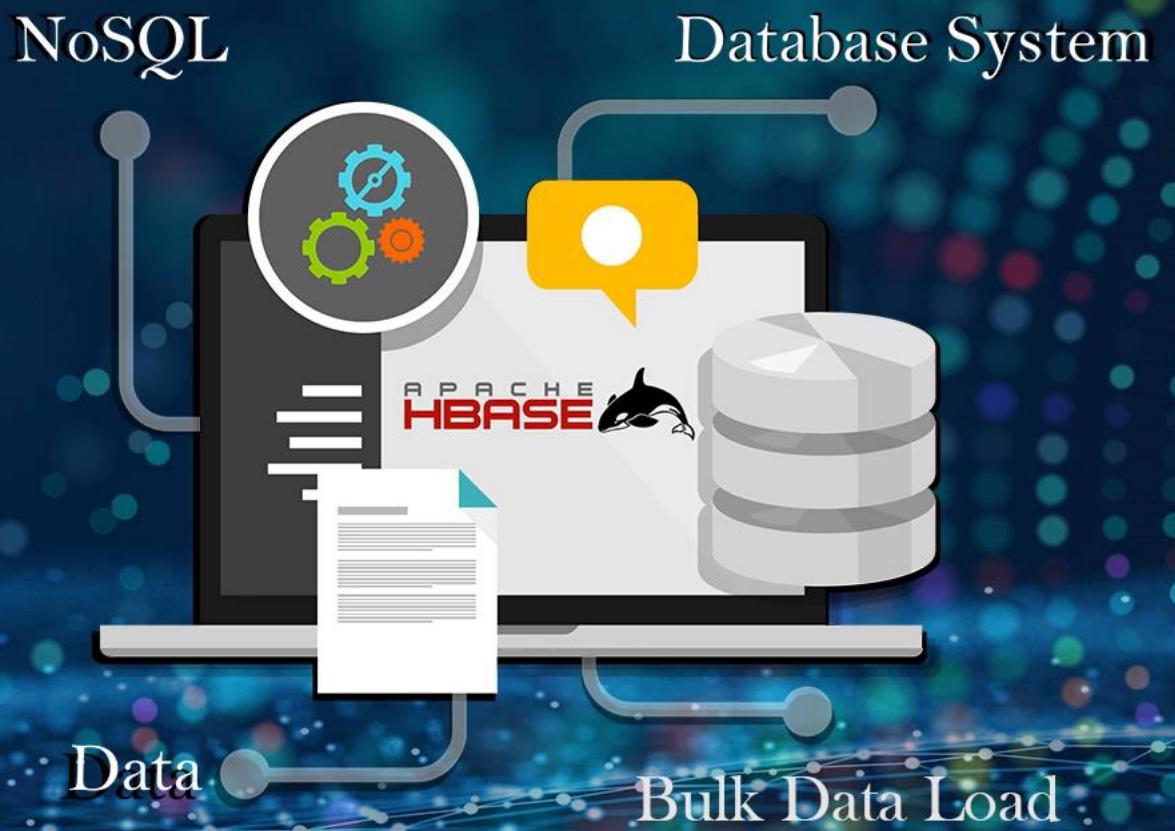
Companies using HBase



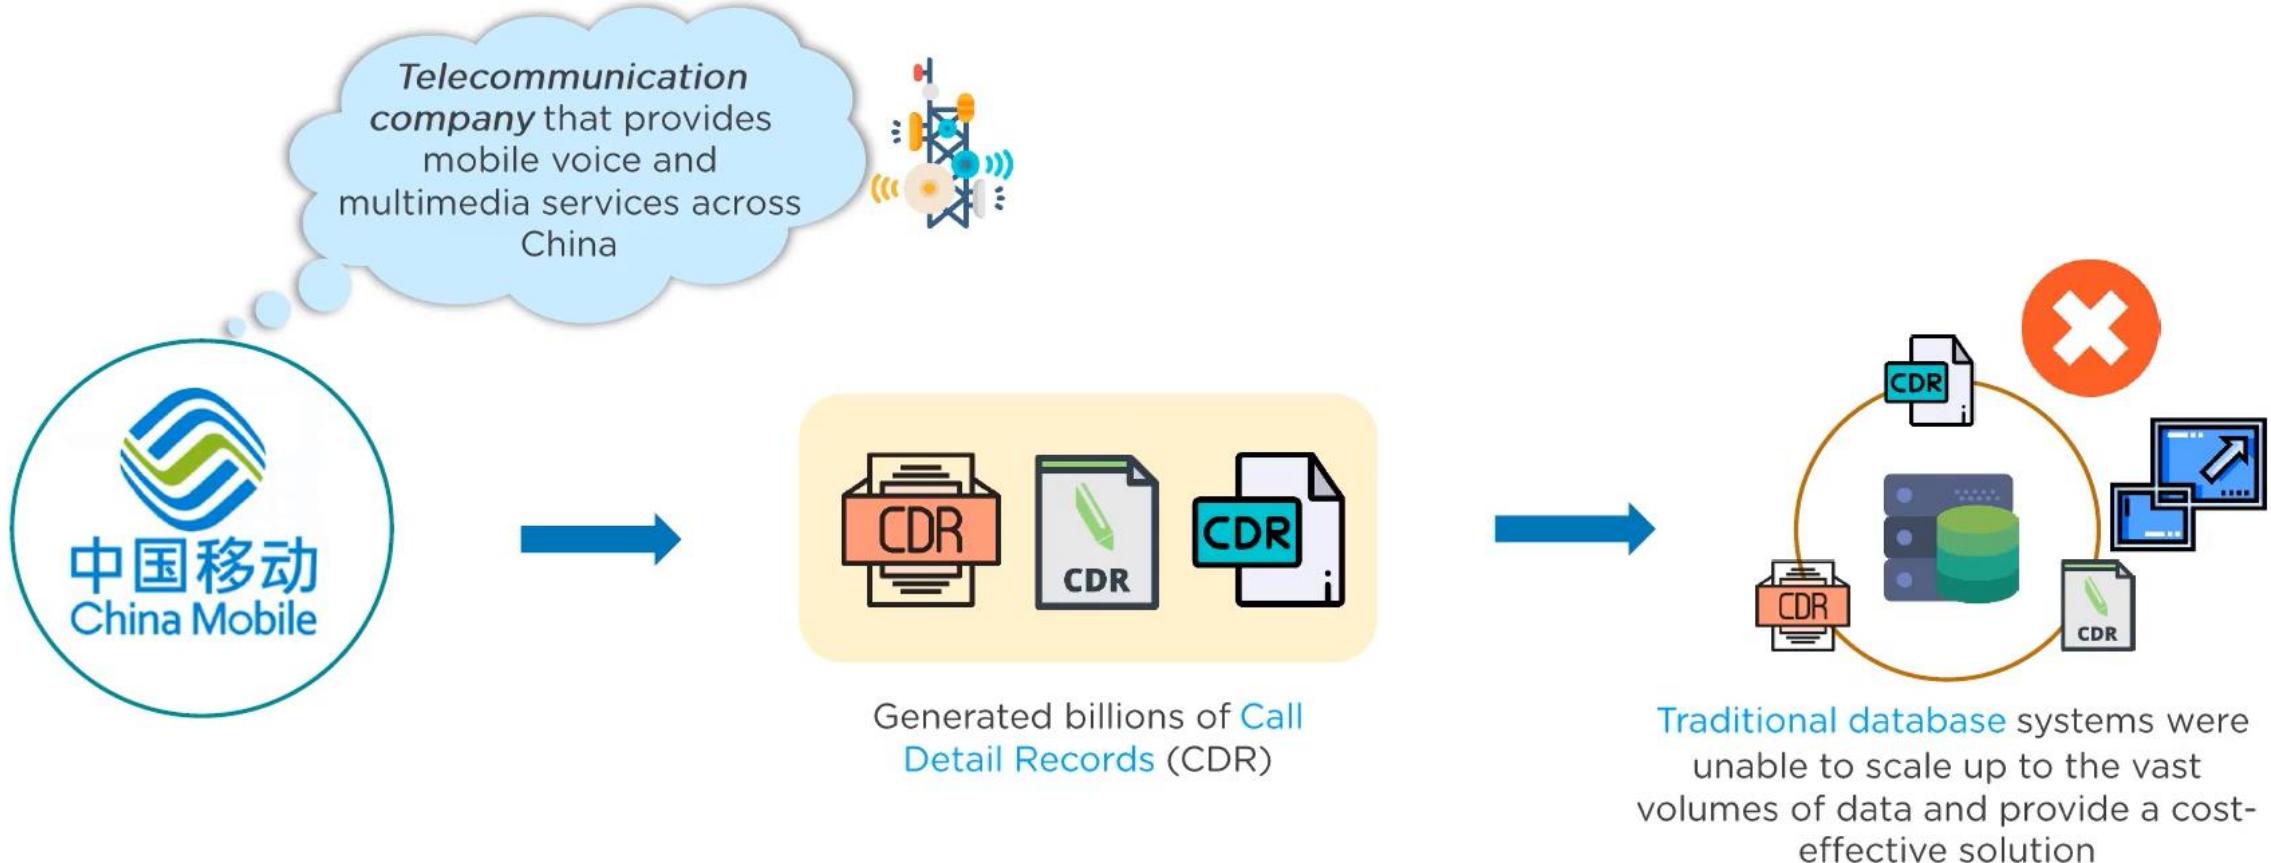
JPMORGAN
CHASE & CO.



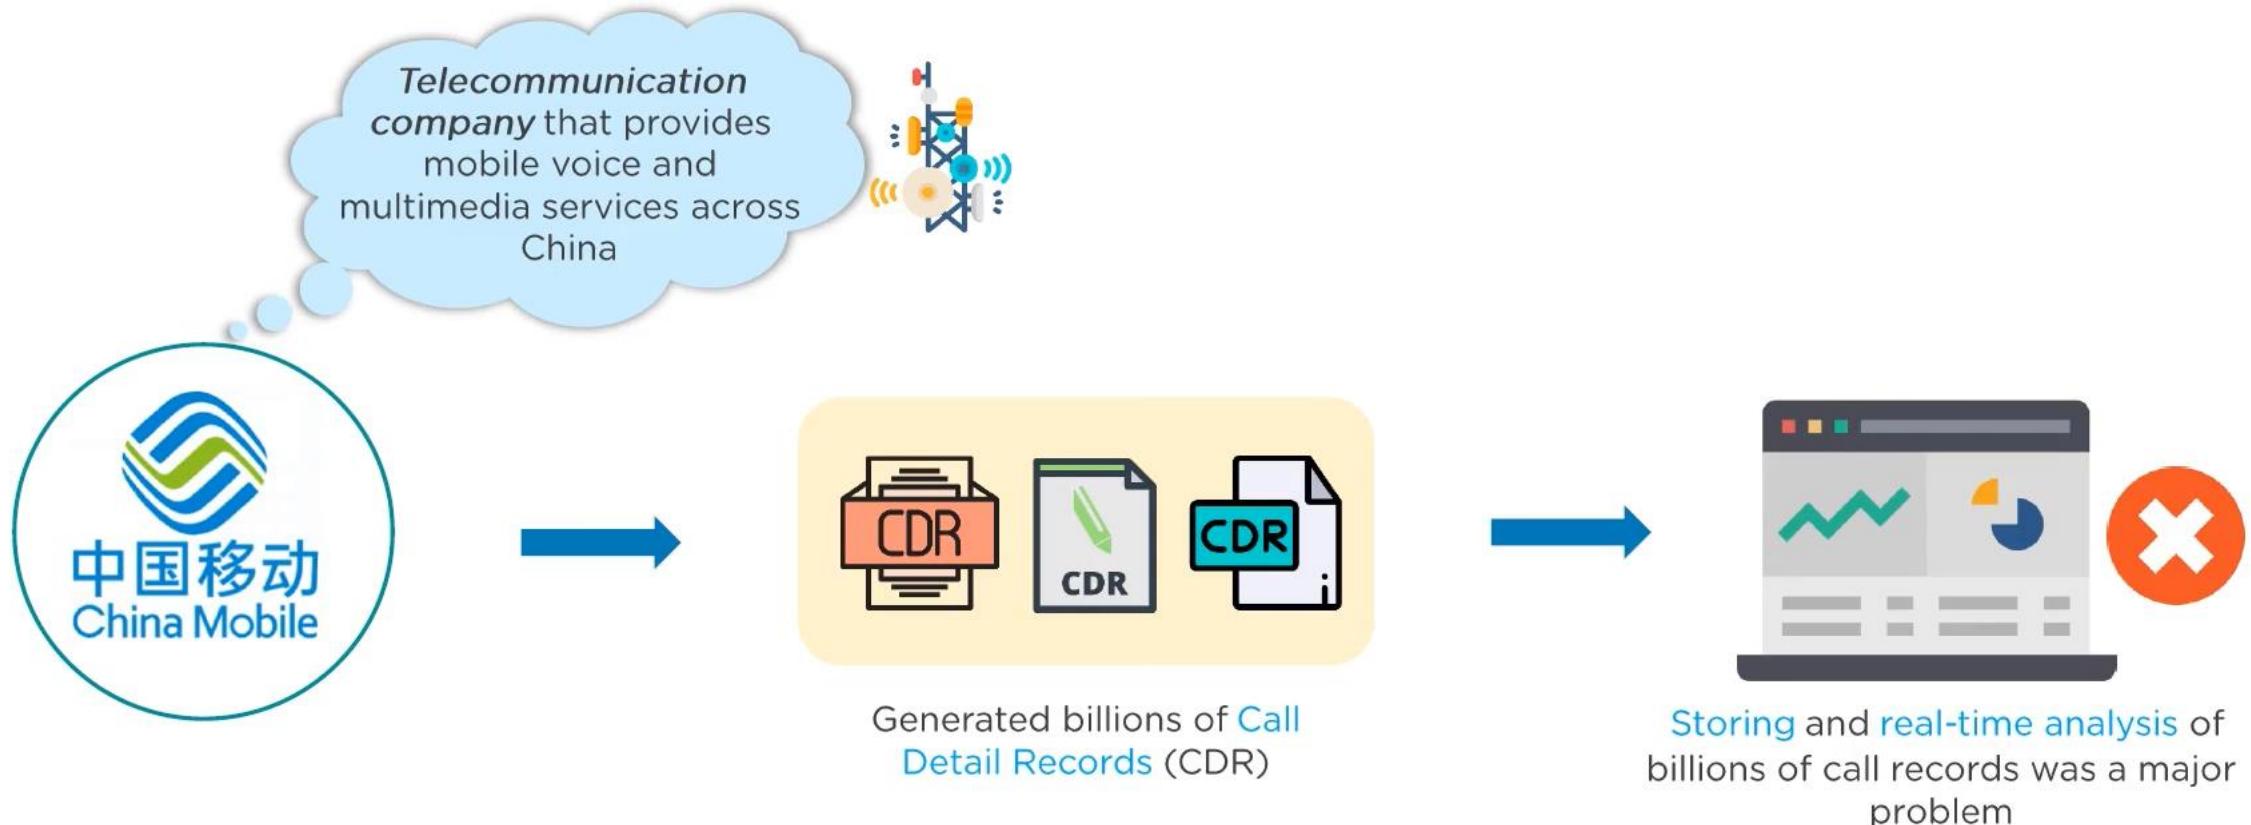
HBase Use Case



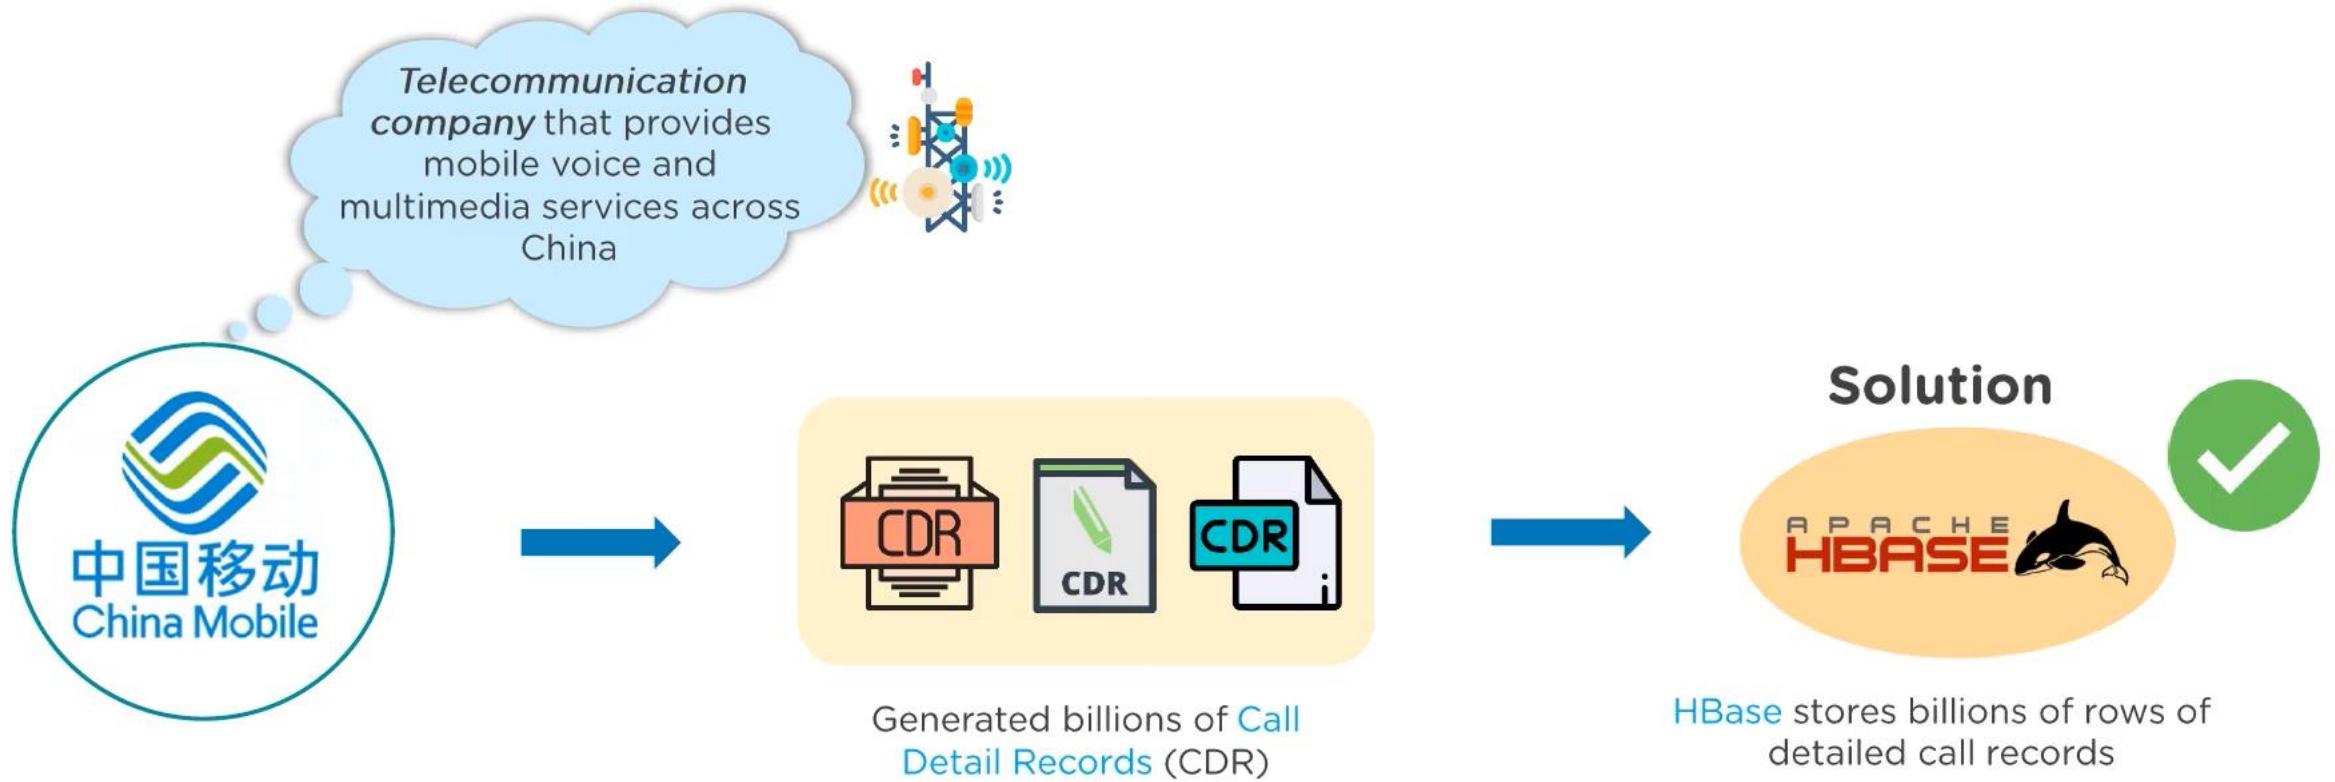
HBase Use Case



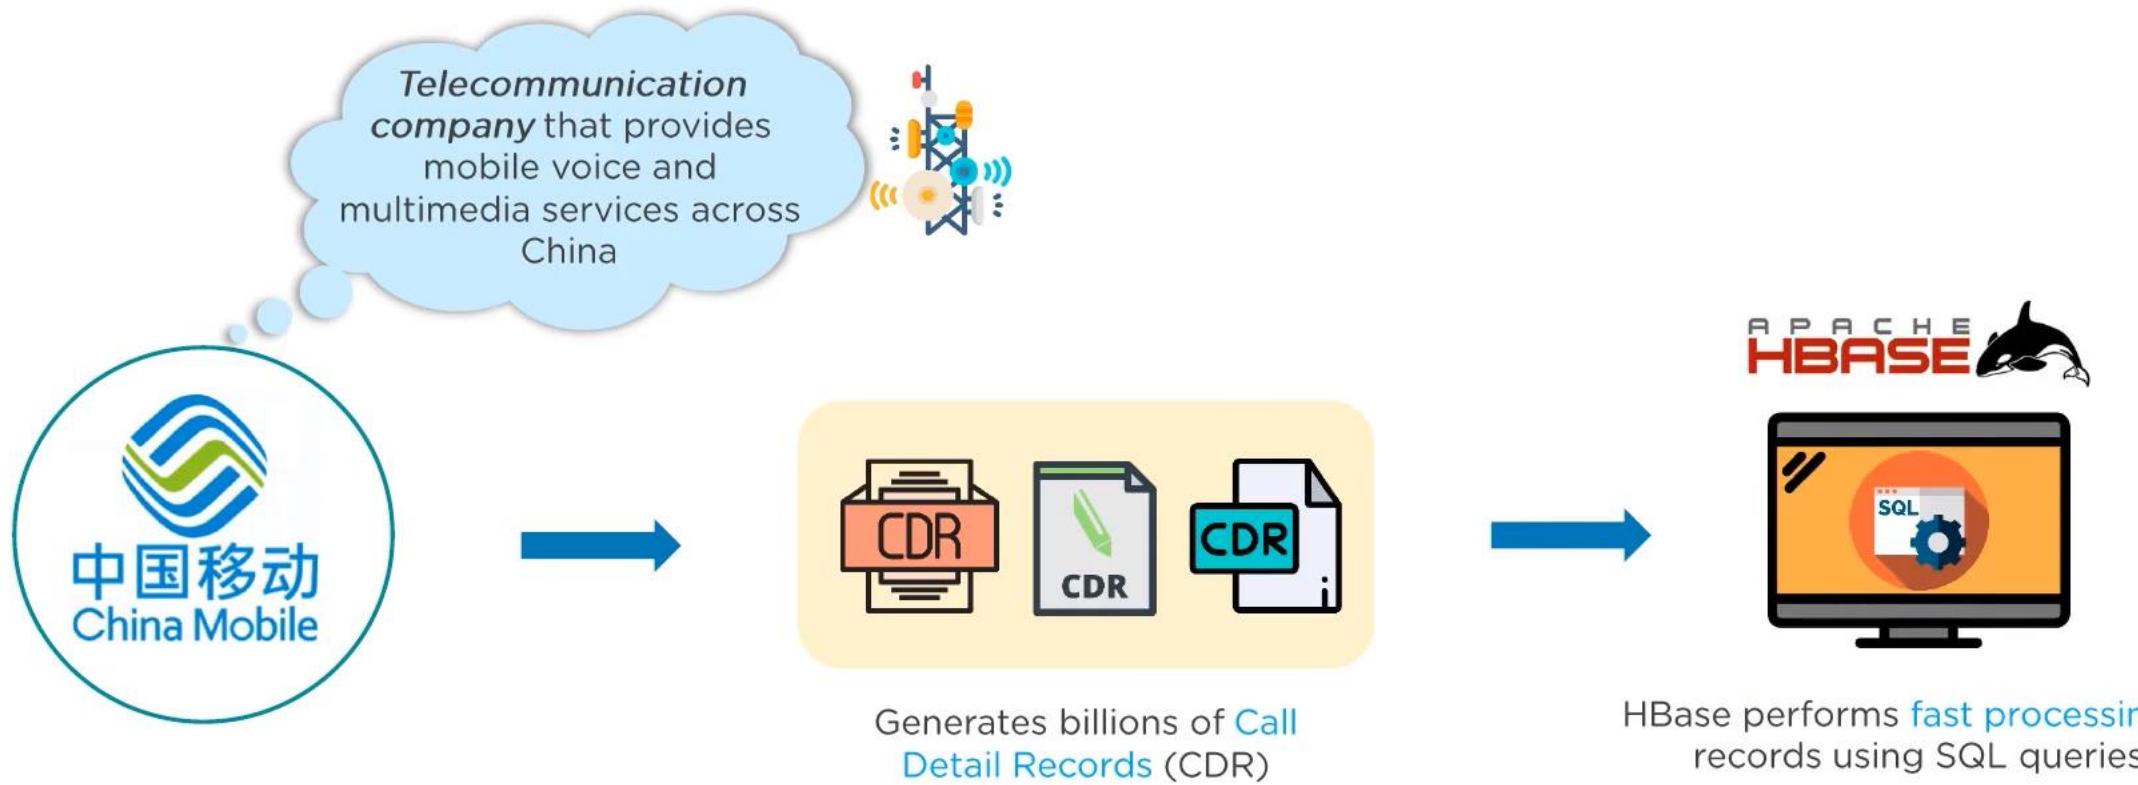
HBase Use Case



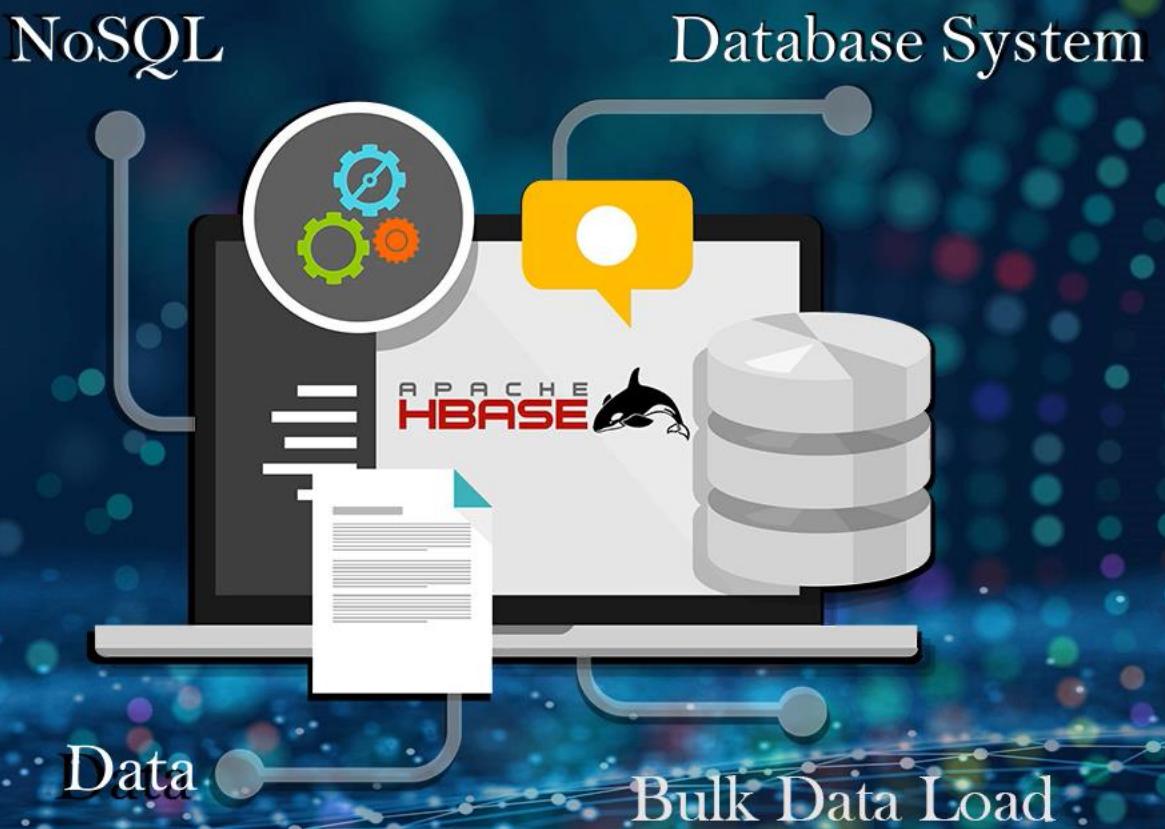
HBase Use Case



HBase Use Case



Applications of HBase



Applications of HBase



Medical

HBase is used for storing genome sequences

Storing disease history of people or an area



E-Commerce

HBase is used for storing logs about customer search history

Performs analytics and target advertisement for better business insights

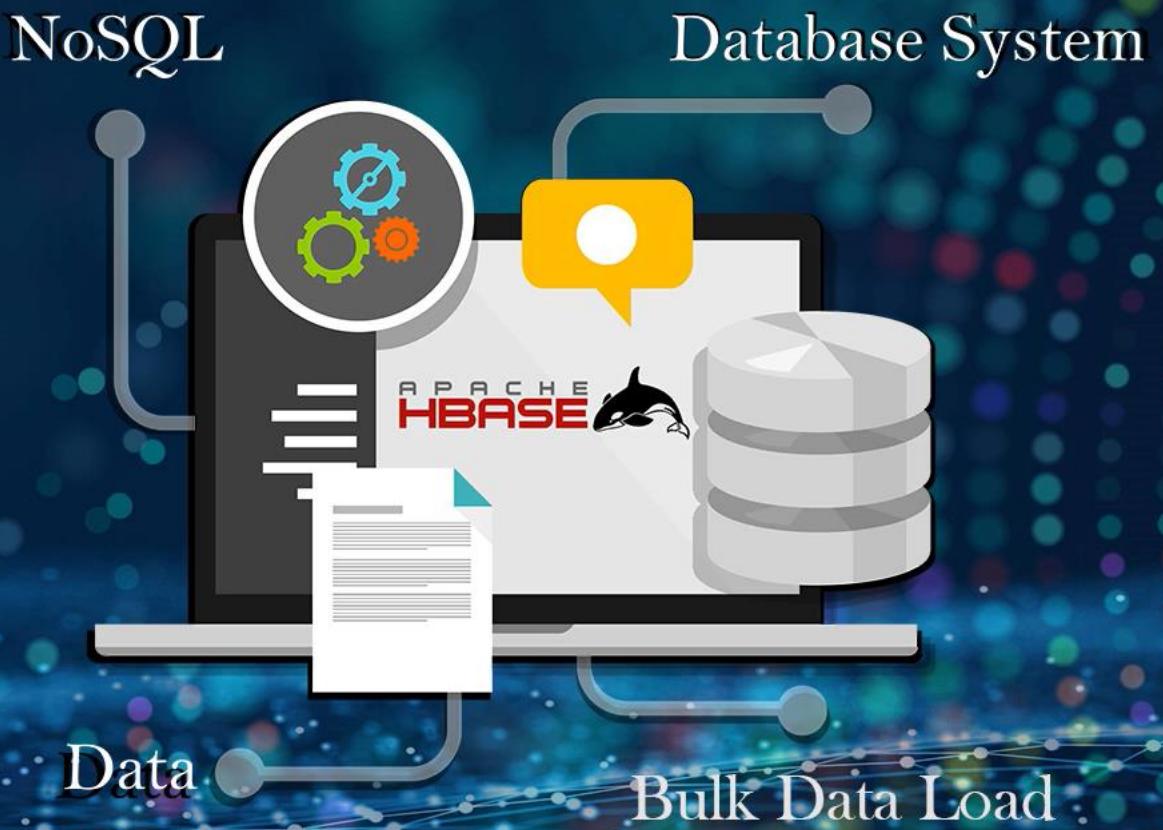


Sports

HBase stores match details and history of each match

Uses this data for better prediction

HBase vs RDBMS



HBase vs RDBMS

HBase

Does not have a fixed schema (schema-less). Defines only column families

Works well with structured and semi-structured data

It can have de-normalized data
(can contain missing or NA values)

Built for wide tables that can be scaled horizontally

RDBMS

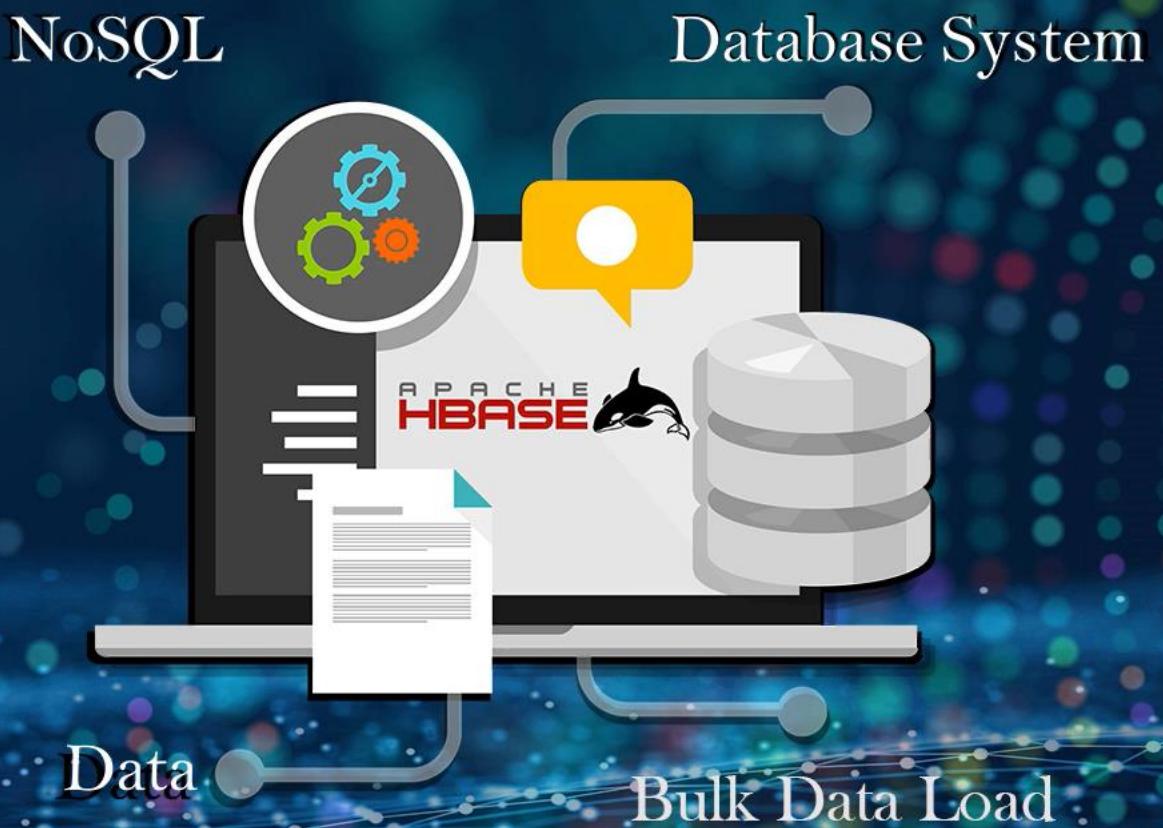
Has a fixed schema which describes the structure of the tables

Works well with structured data

RDBMS can store only normalized data

Built for thin tables that is hard to scale

Features of HBase



Features of HBase

Scalable



Automatic failure support



Consistent read and write



JAVA API for client access



Block cache and bloom filters



Data can be scaled across various nodes as it is stored in HDFS

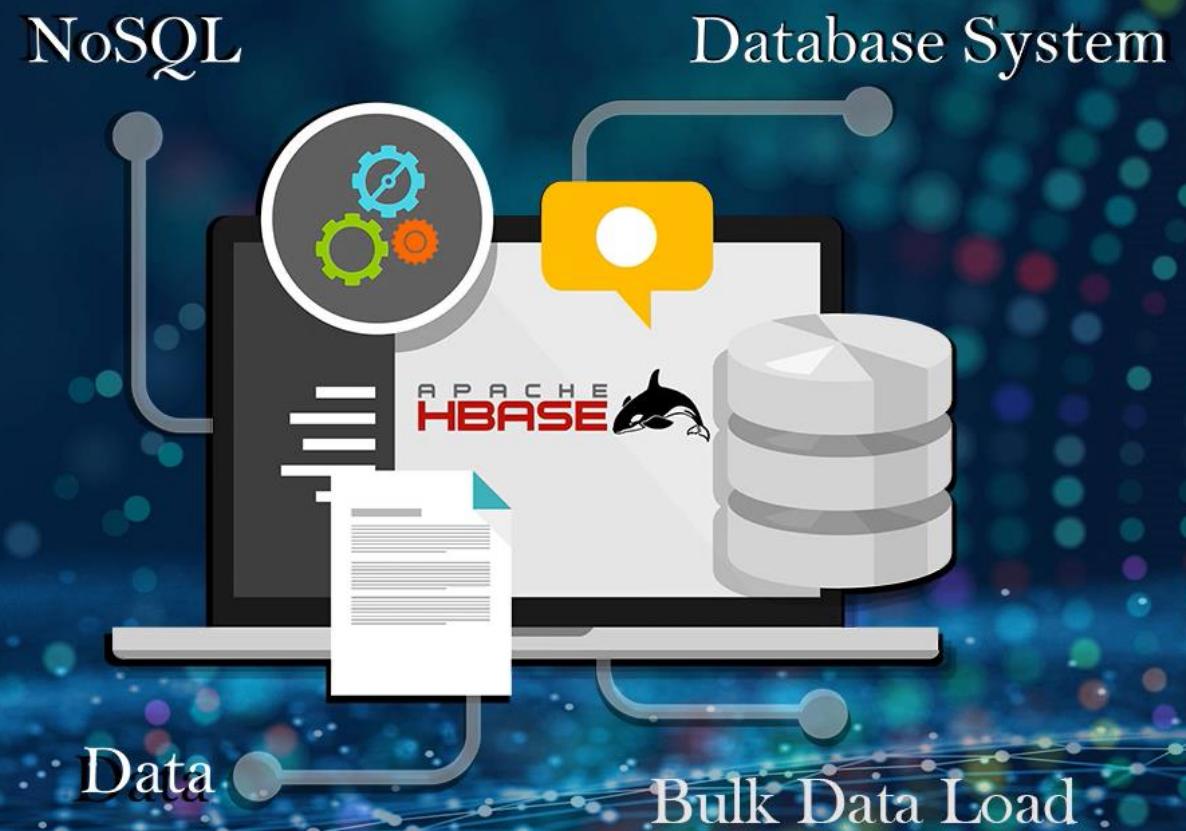
Write Ahead Log across clusters which provides automatic support against failure

HBase provides consistent read and write of data

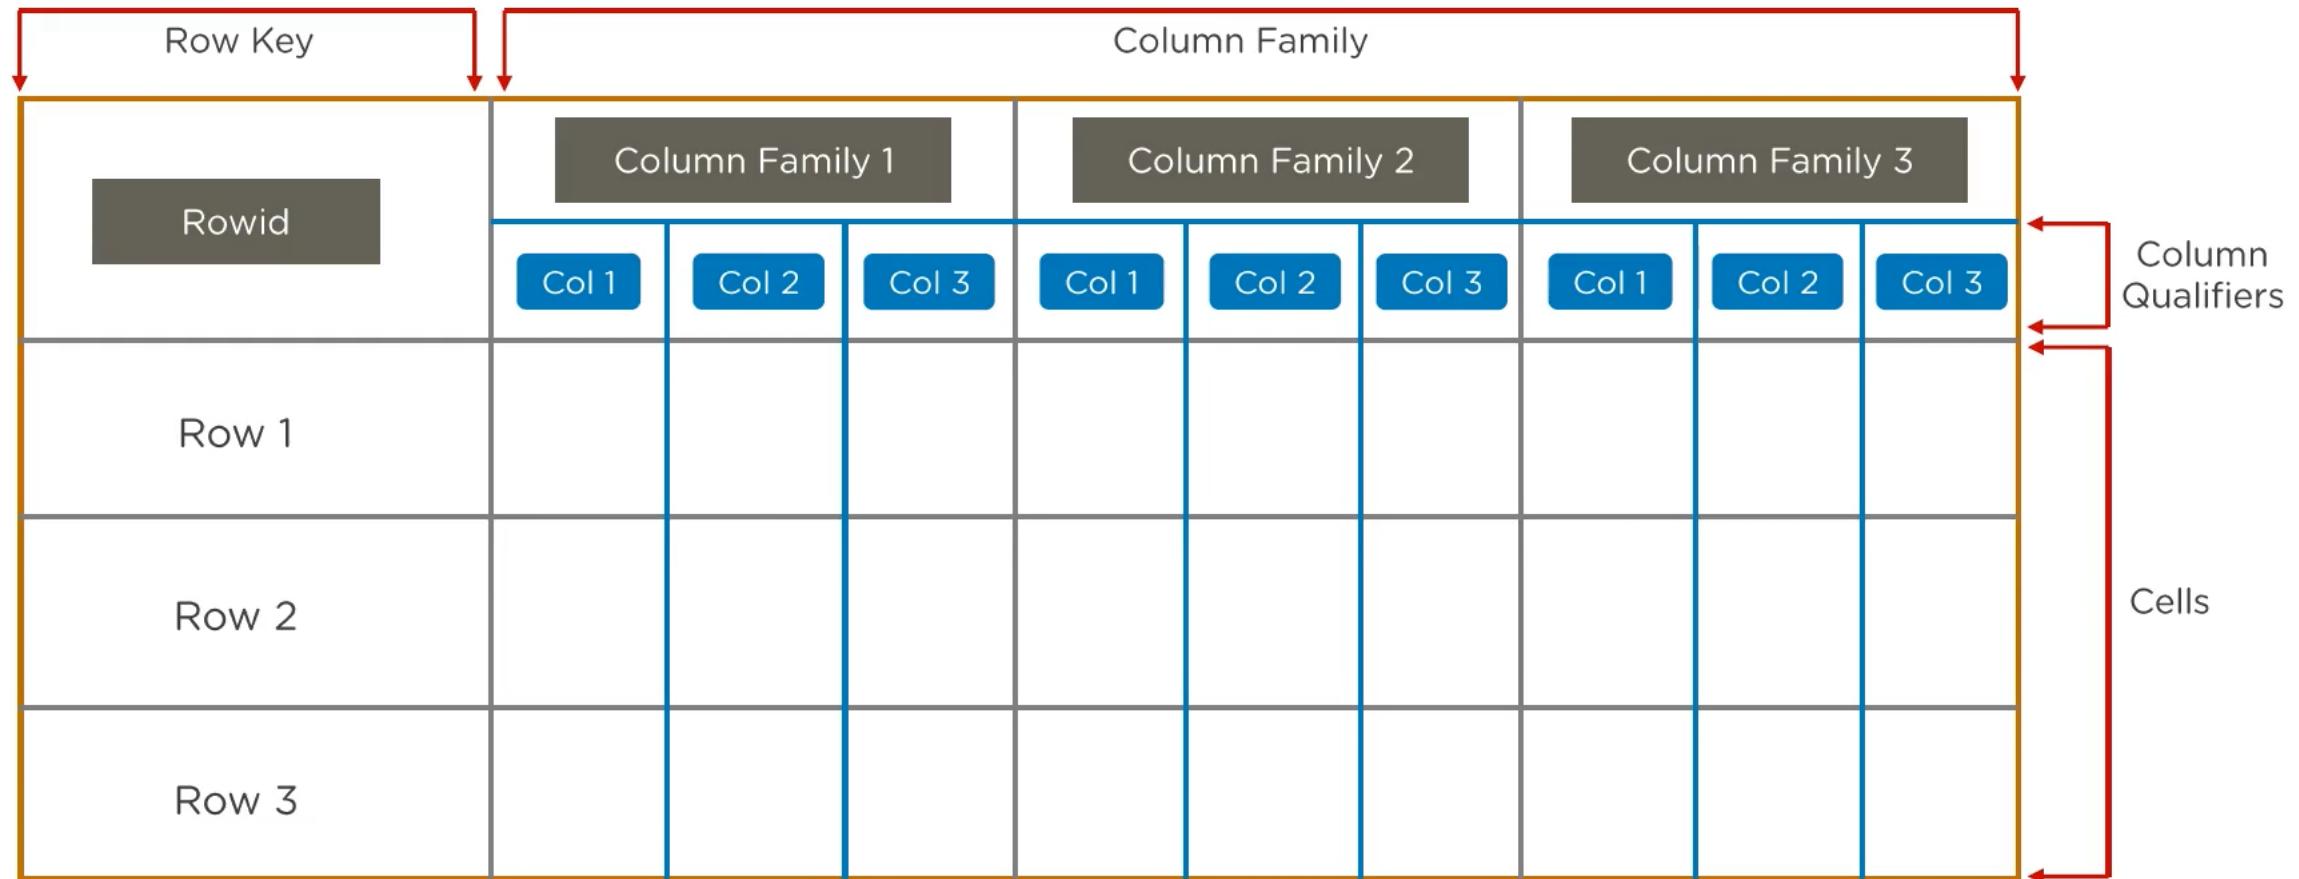
Provides easy to use JAVA API for clients

Supports block cache and bloom filters for high volume query optimization

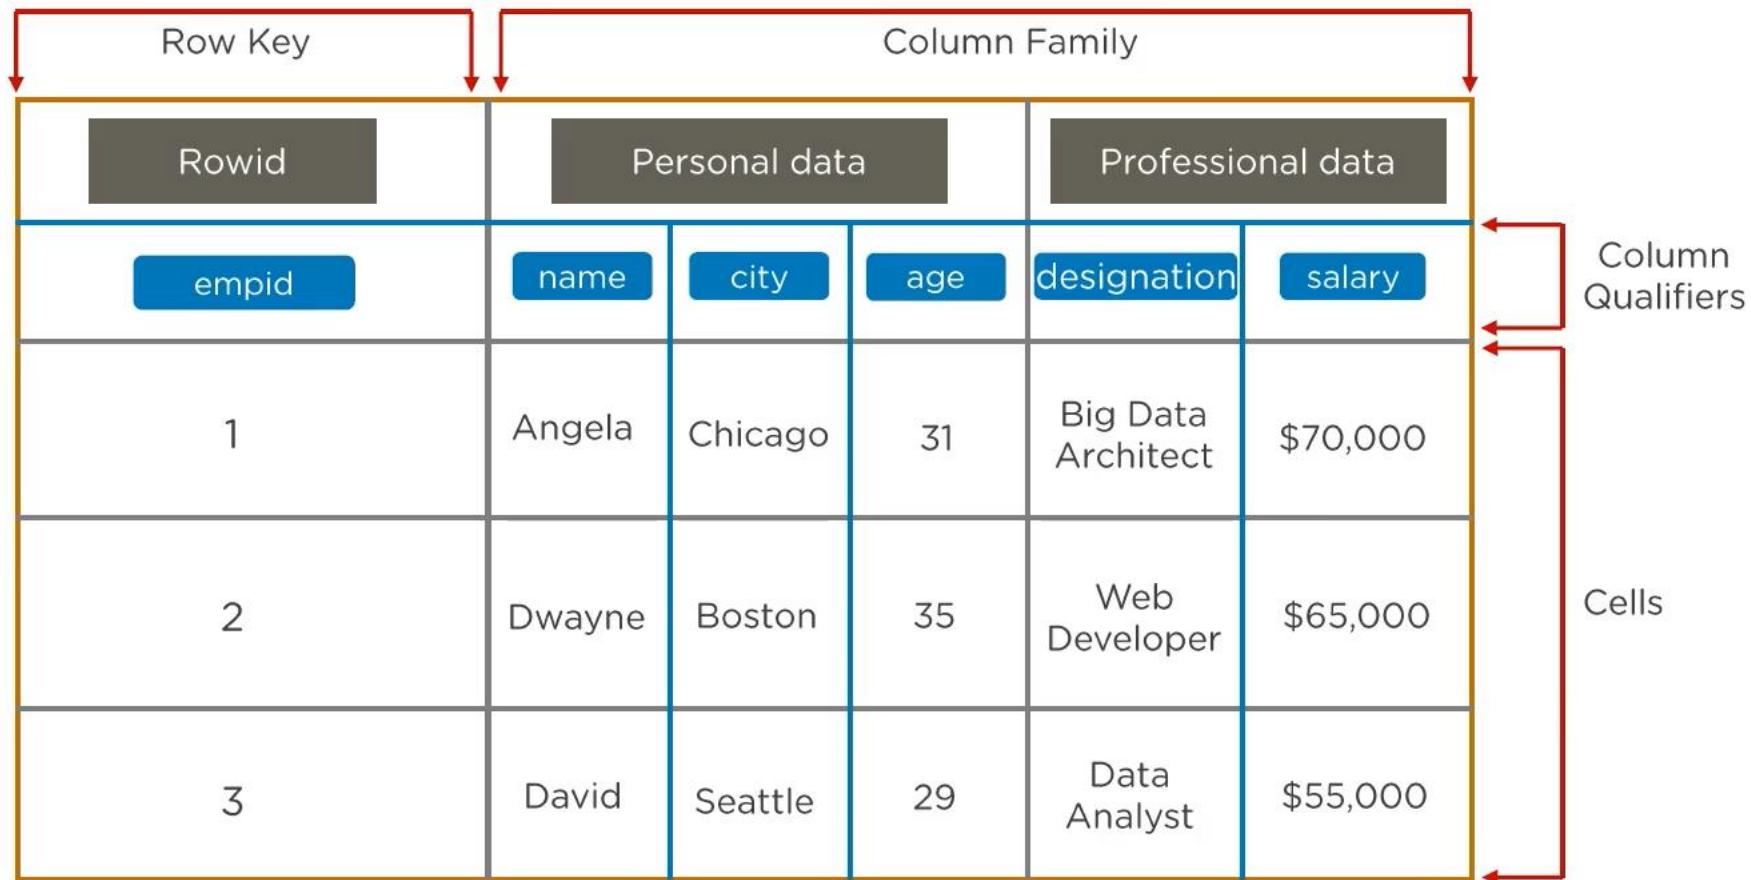
HBase Storage



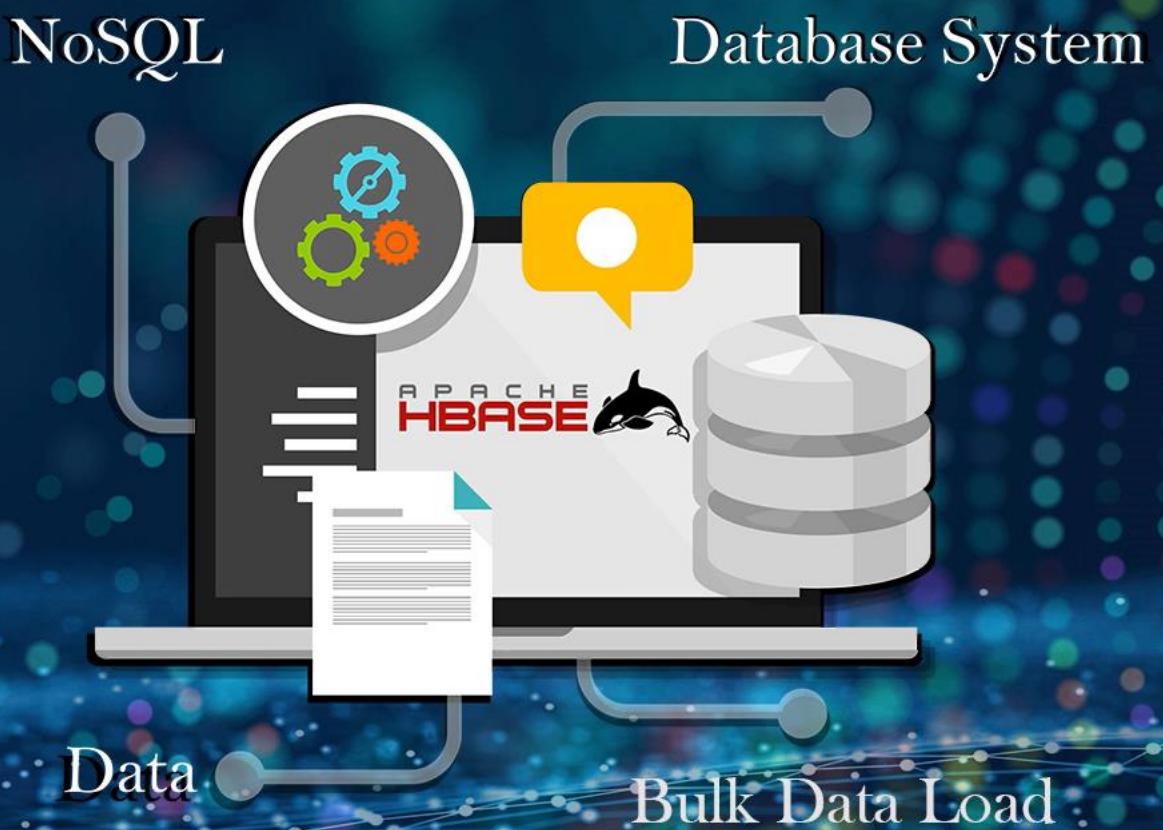
HBase column oriented storage



HBase column oriented storage



HBase Architecture



HBase Architectural Components



ZooKeeper is used for monitoring

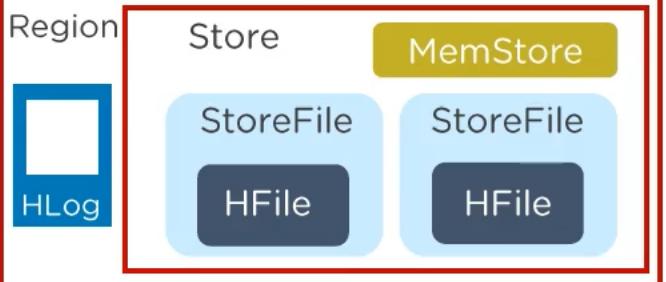


HMaster

HBase Master assigns regions and load balancing

Region server serves data for read and write

Region Server



Region Server



Region Server



HDFS

HBase Architectural Components - Regions



HBase tables are divided horizontally by row key range into “Regions”



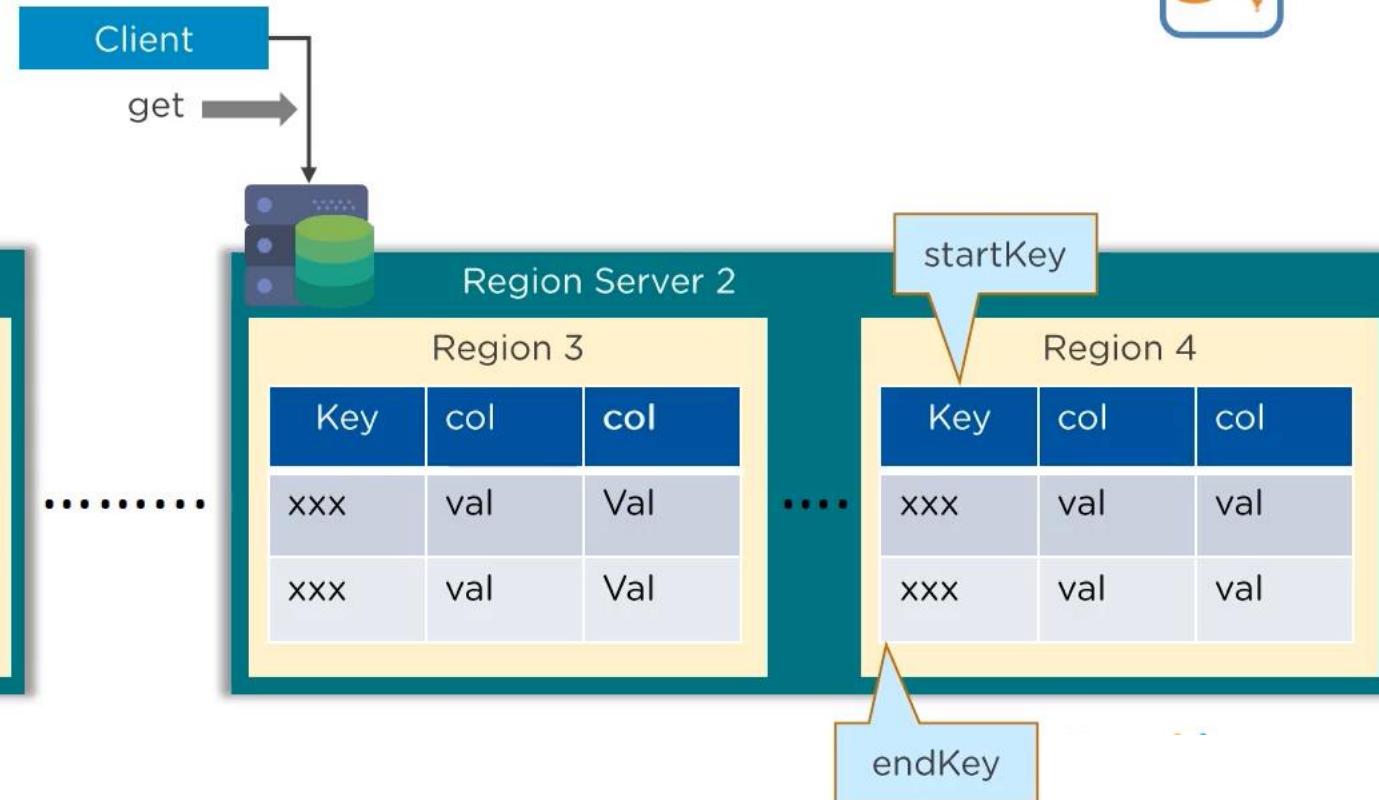
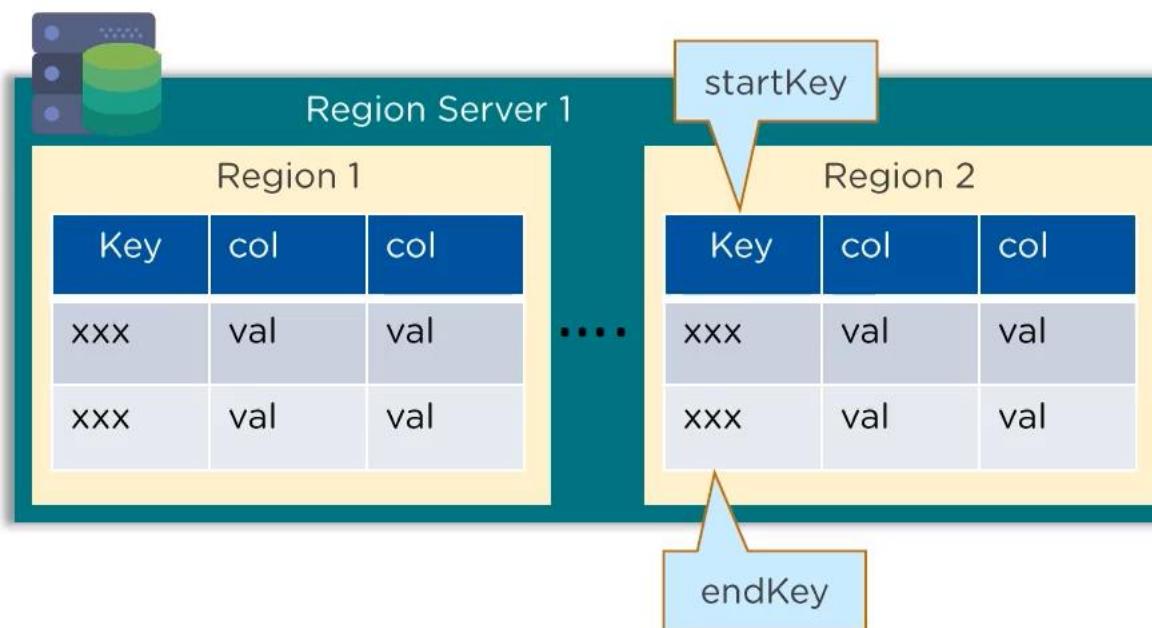
Regions are assigned to the nodes in the cluster, called “Region Servers”



A region contains all rows in the table between the region's **start key** and **end key**



These servers serve data for **read** and **write**



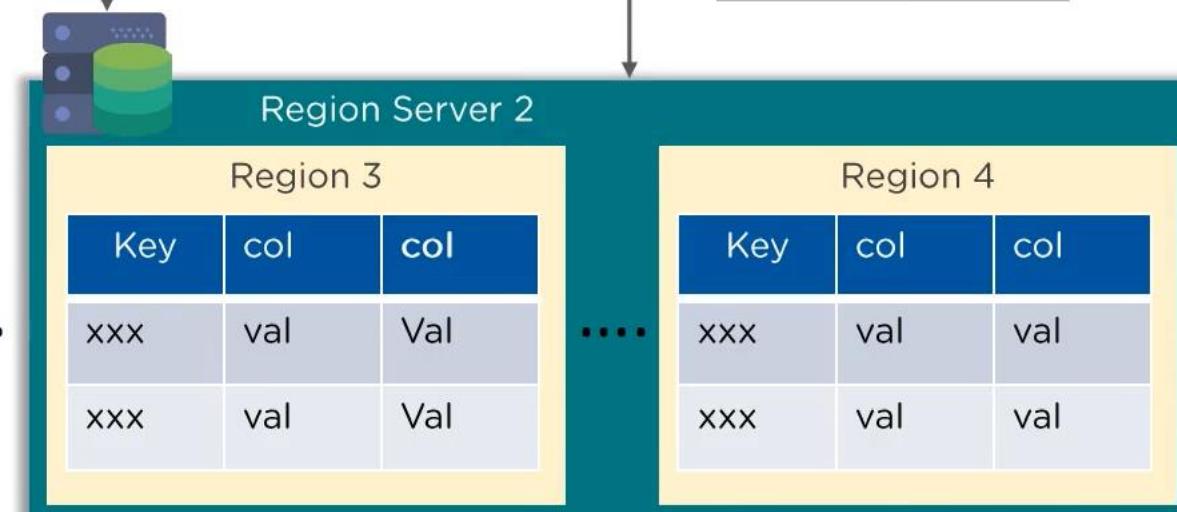
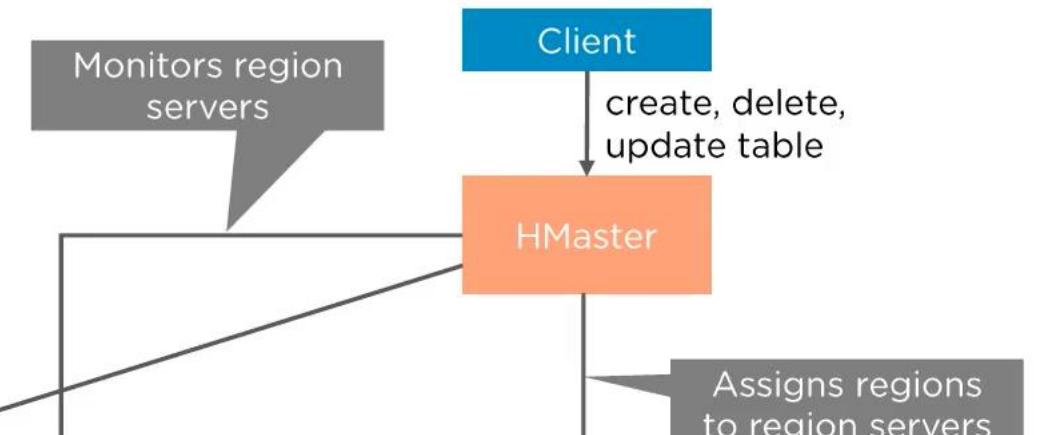
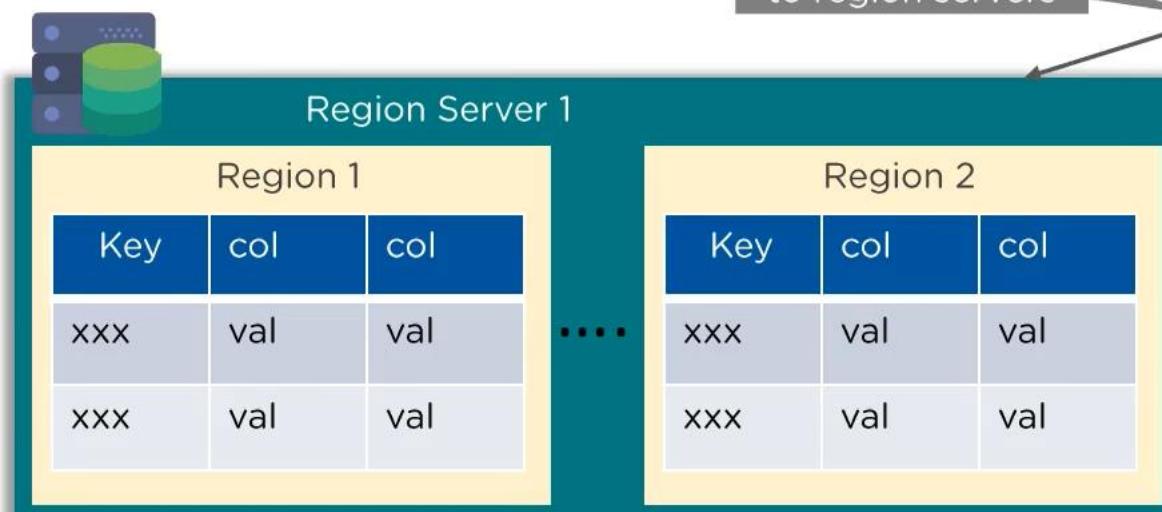
HBase Architectural Components - HMaster



Region assignment, Data Definition Language operation (create, delete) are handled by HMaster



Assigning and re-assigning regions for recovery or load balancing and monitoring all servers



HBase has a distributed environment where HMaster alone is not sufficient to manage everything. Hence, [ZooKeeper](#) was introduced

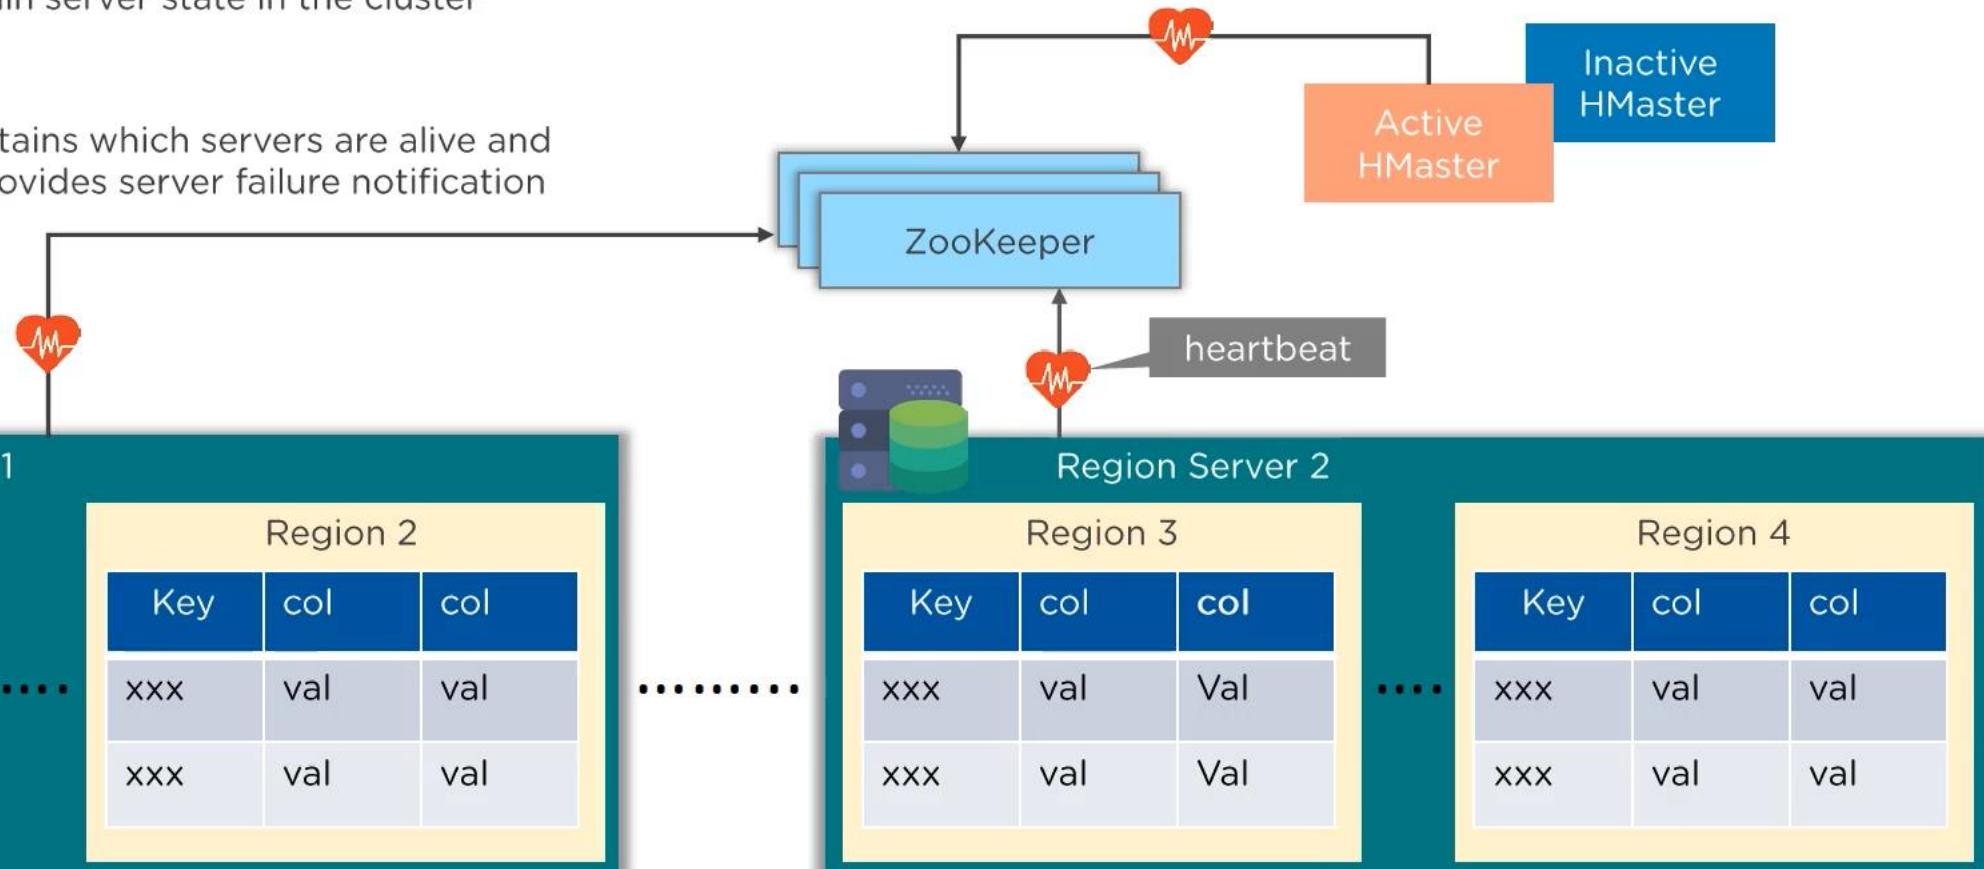
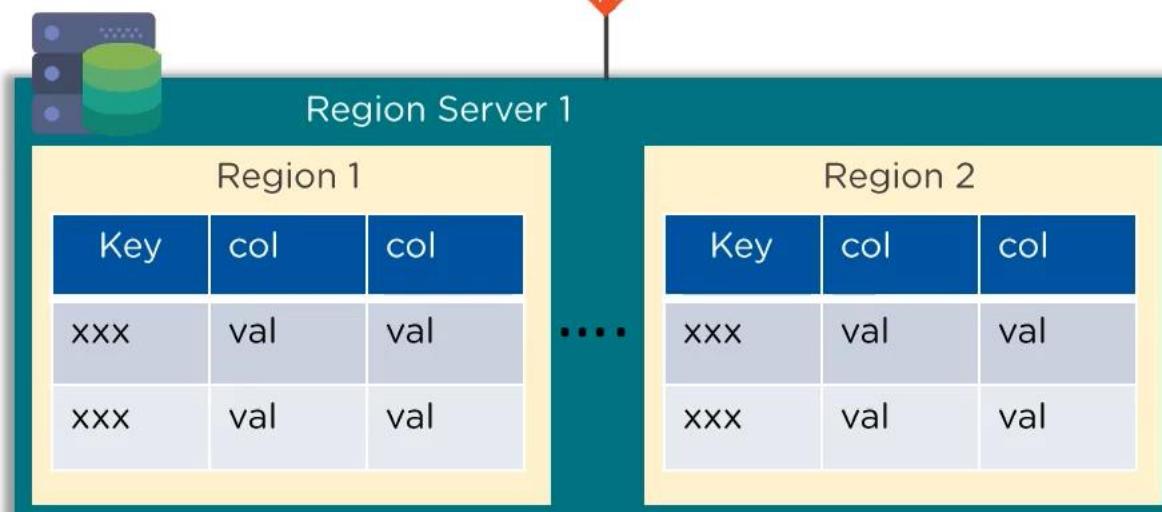
HBase Architectural Components - ZooKeeper



ZooKeeper is a distributed coordination service to maintain server state in the cluster



Zookeeper maintains which servers are alive and available, and provides server failure notification



Region servers send their status to ZooKeeper indicating they are ready for read and write operation

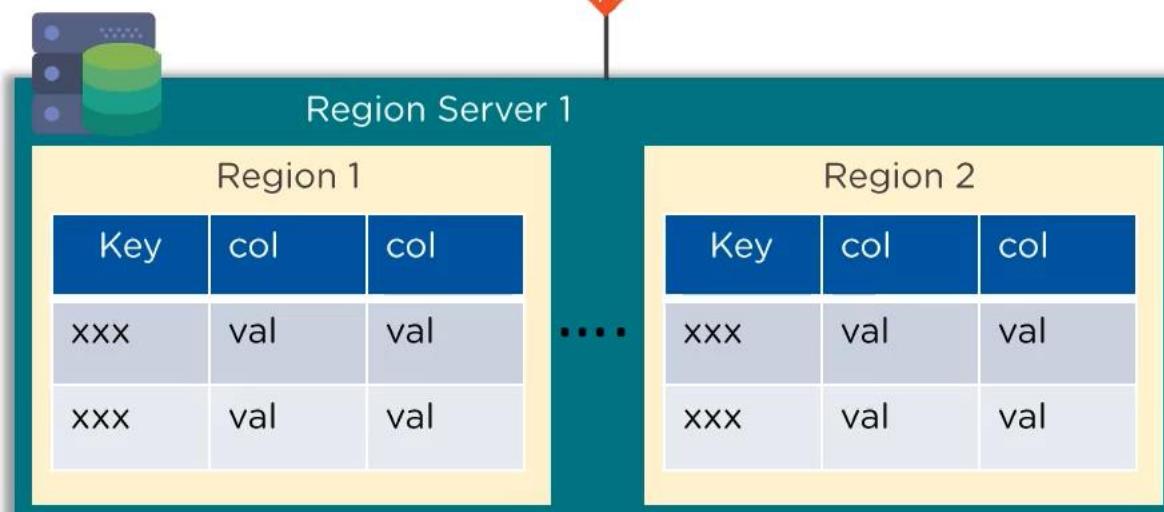
HBase Architectural Components - ZooKeeper



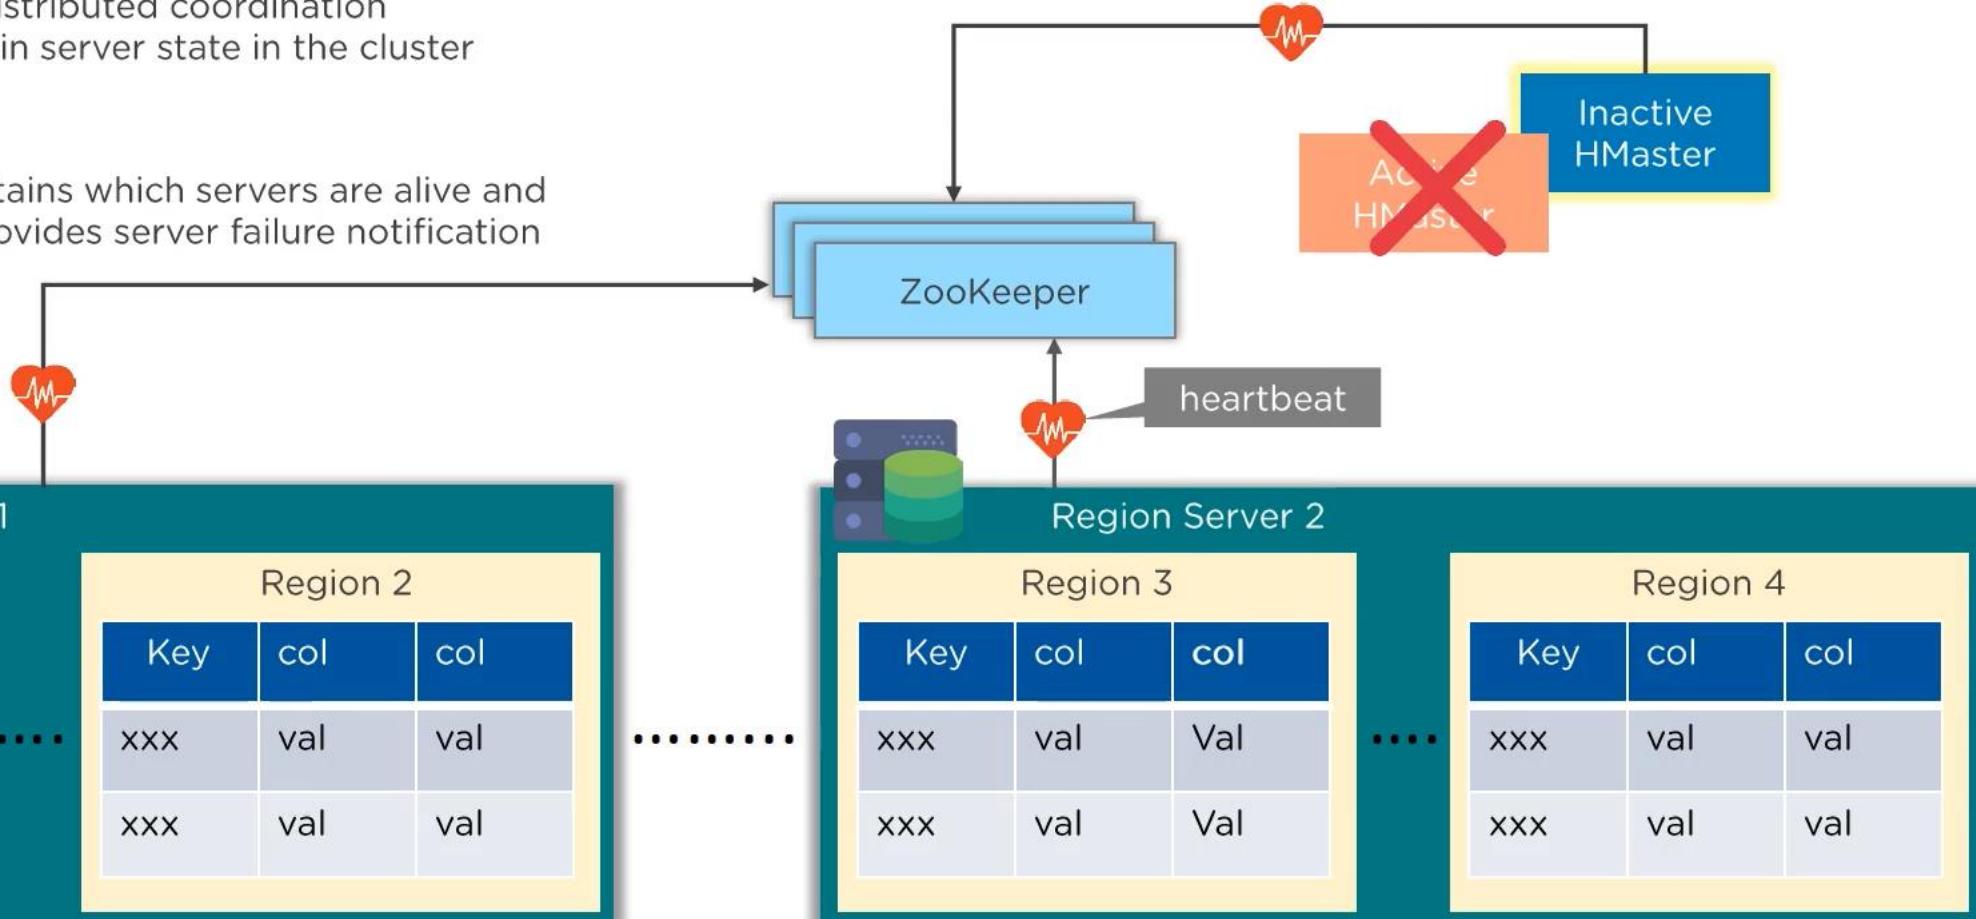
ZooKeeper is a distributed coordination service to maintain server state in the cluster



Zookeeper maintains which servers are alive and available, and provides server failure notification

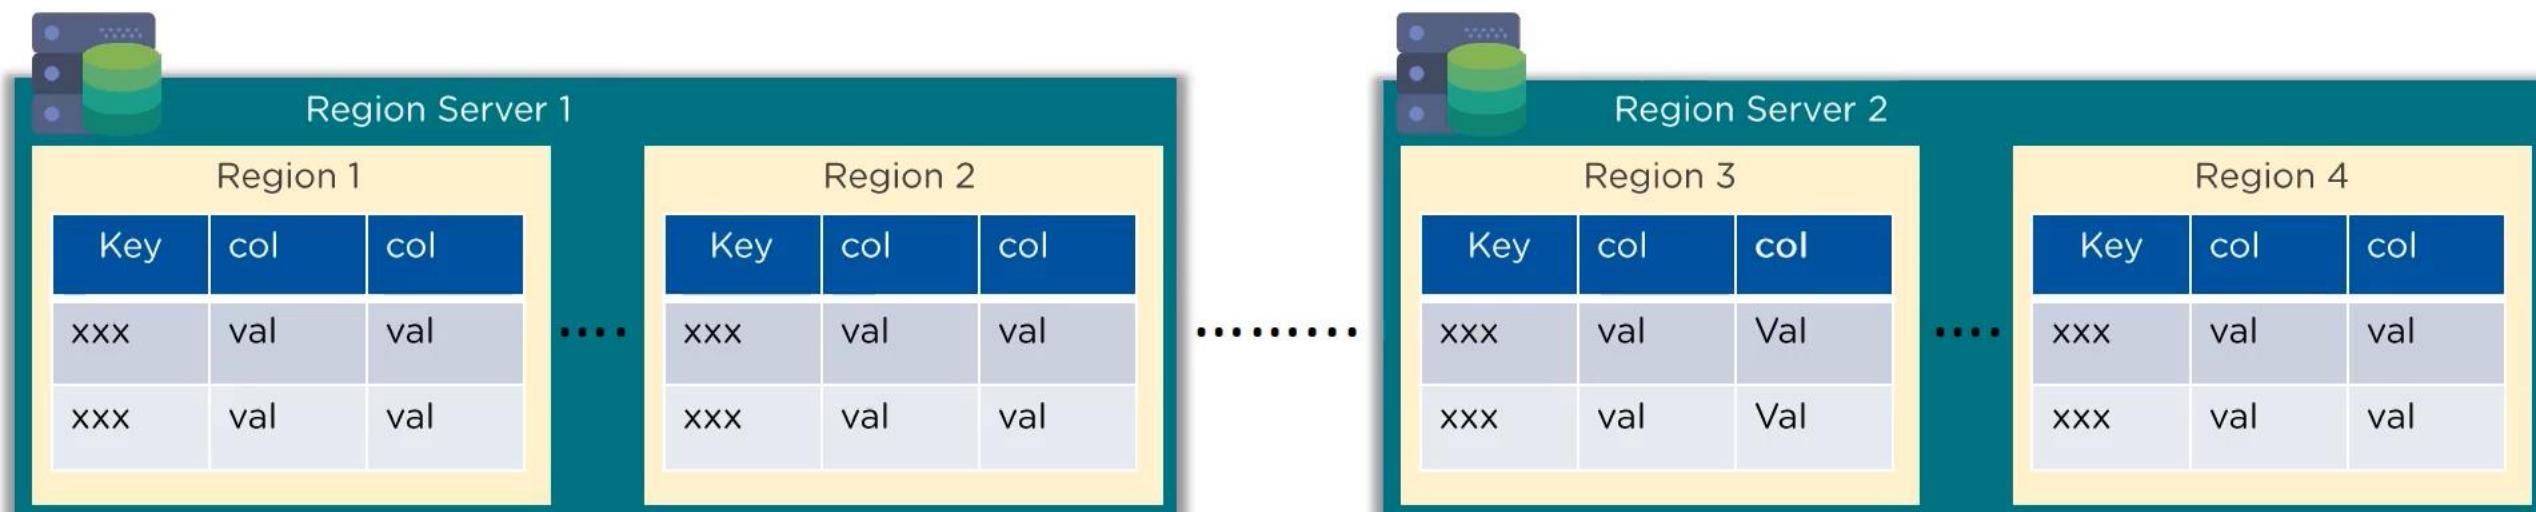
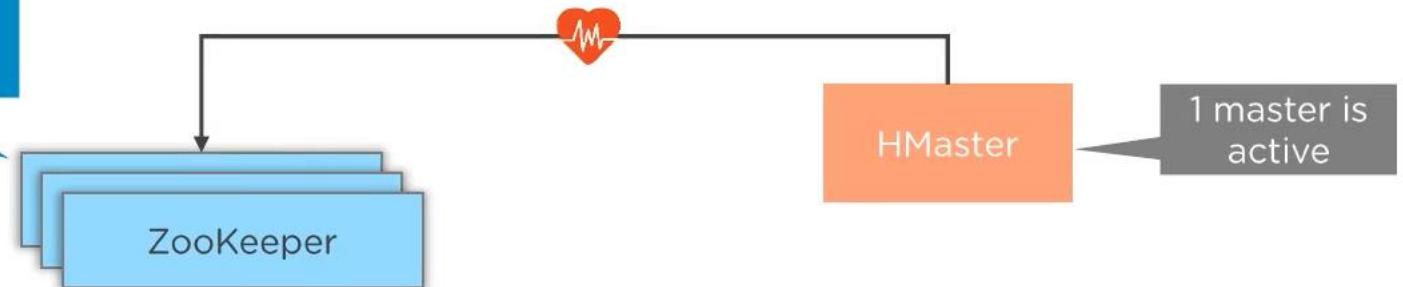


Inactive server acts as a backup. If the active HMaster fails, it will come to rescue



How the components work together?

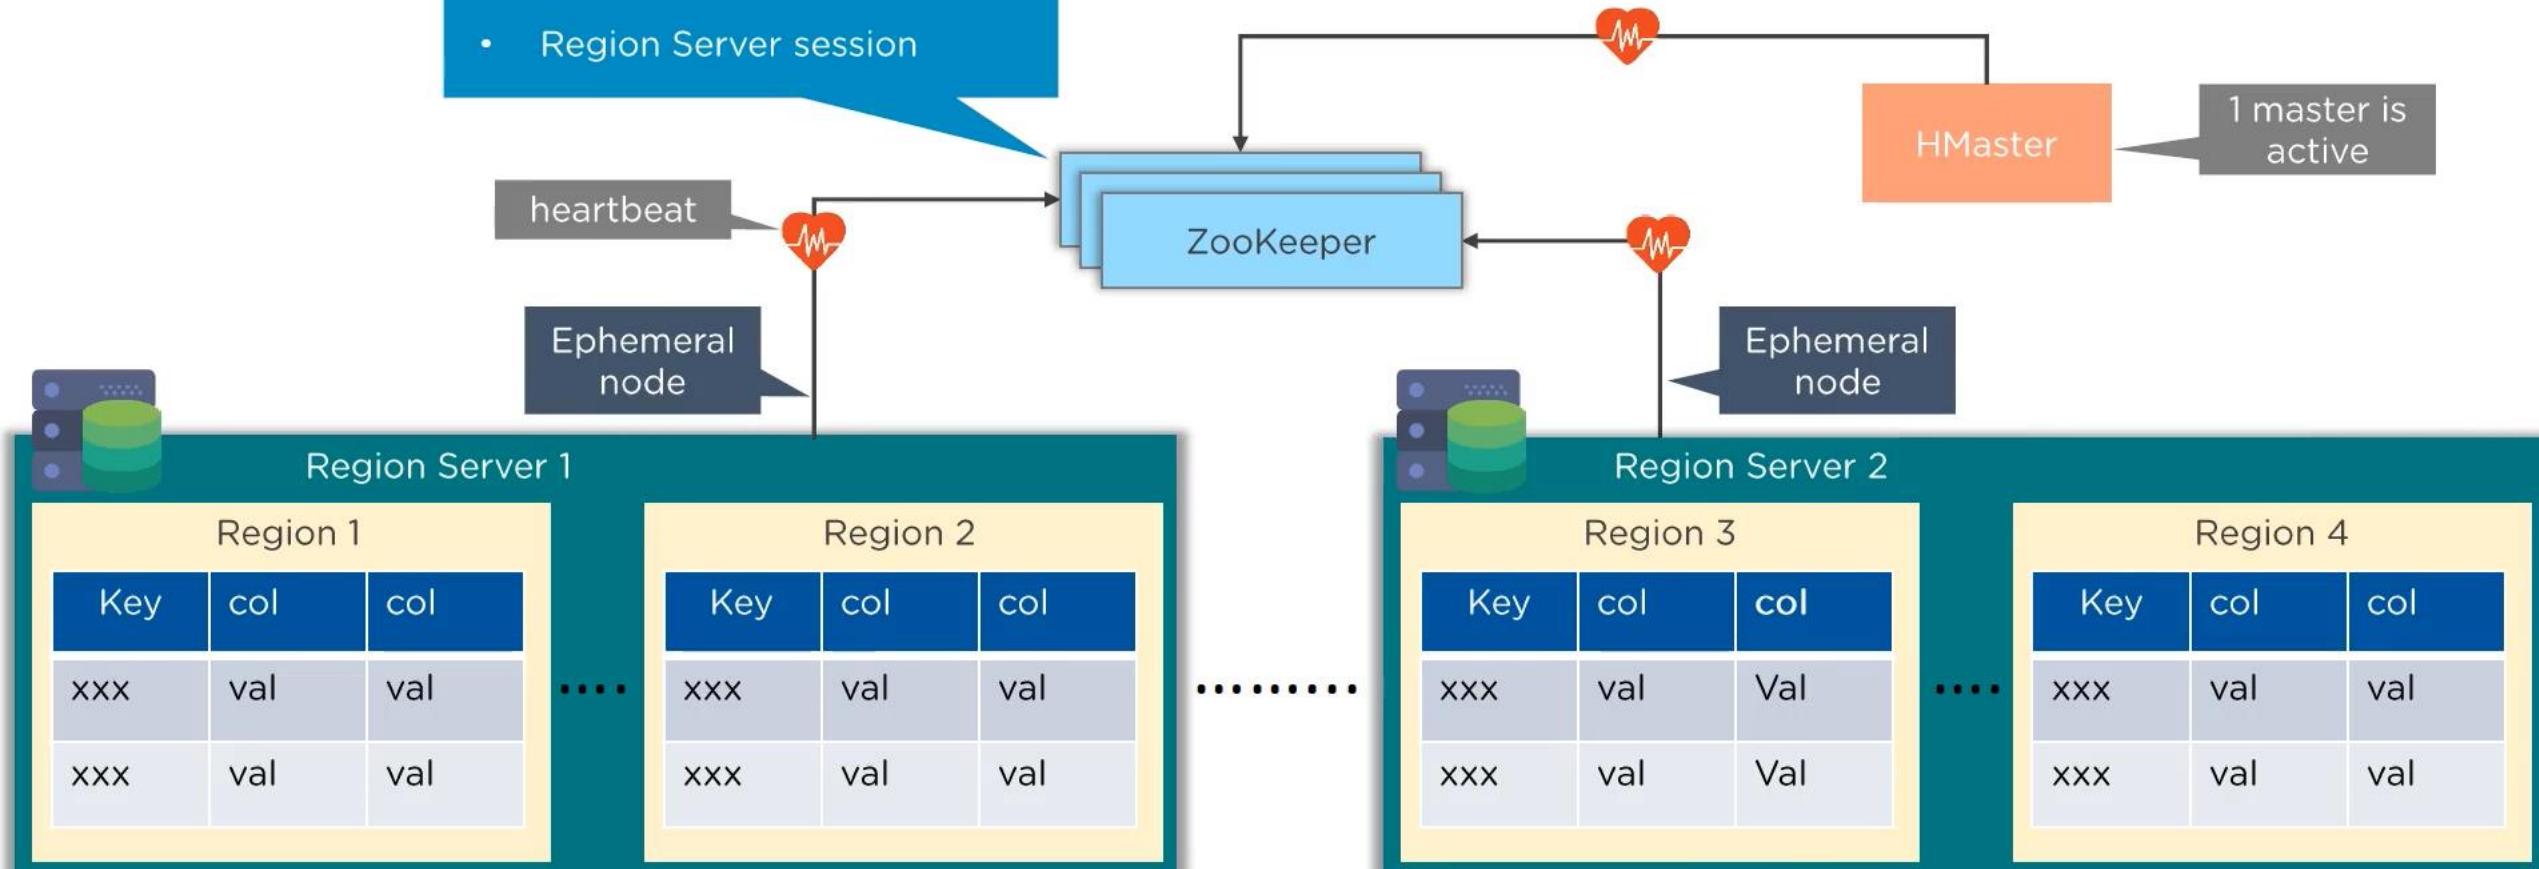
- Active HMaster selection
- Region Server session



Active HMaster and Region Servers connect with a session to ZooKeeper

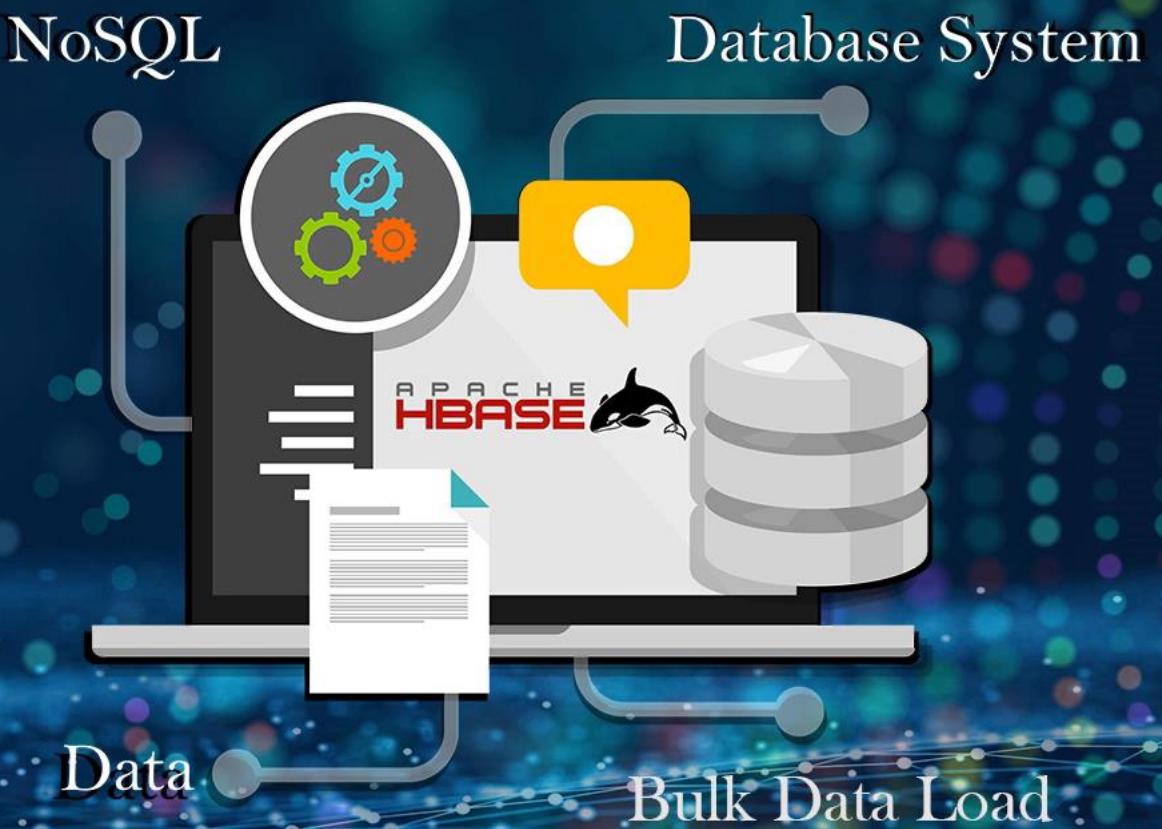
How the components work together?

- Active HMaster selection
- Region Server session



ZooKeeper maintains [ephemeral nodes](#) for active sessions via heartbeats to indicate that region servers are up and running

HBase Read or Write

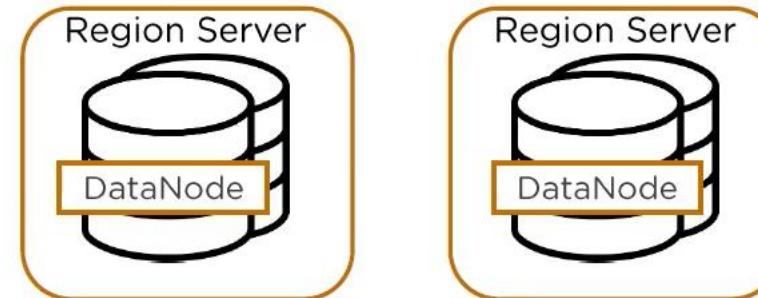
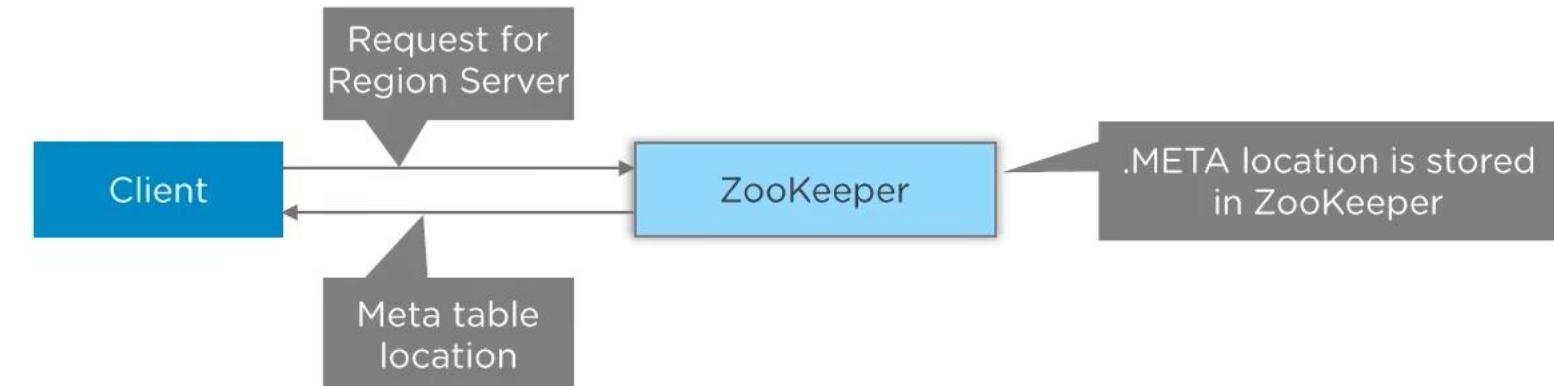


HBase Read or Write

There is a special HBase Catalog table called the META table, which holds the location of the regions in the cluster

Here is what happens the first time a client reads or writes data to HBase

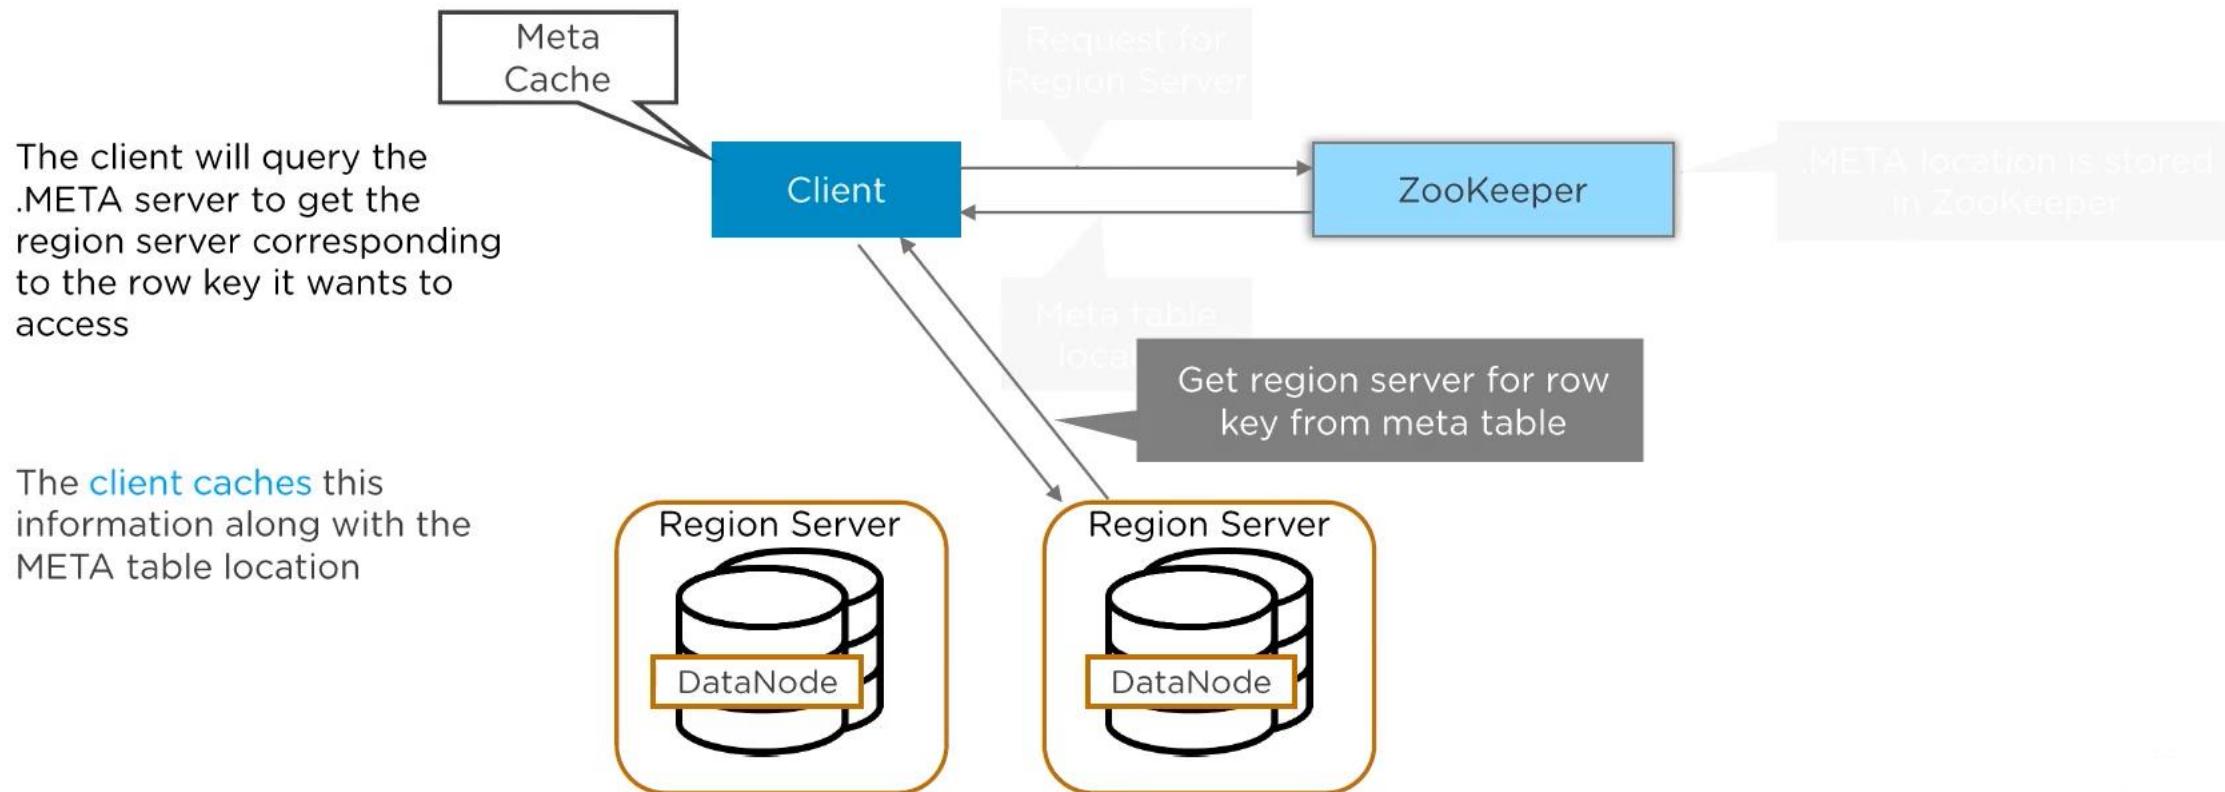
The client gets the Region Server that hosts the META table from ZooKeeper



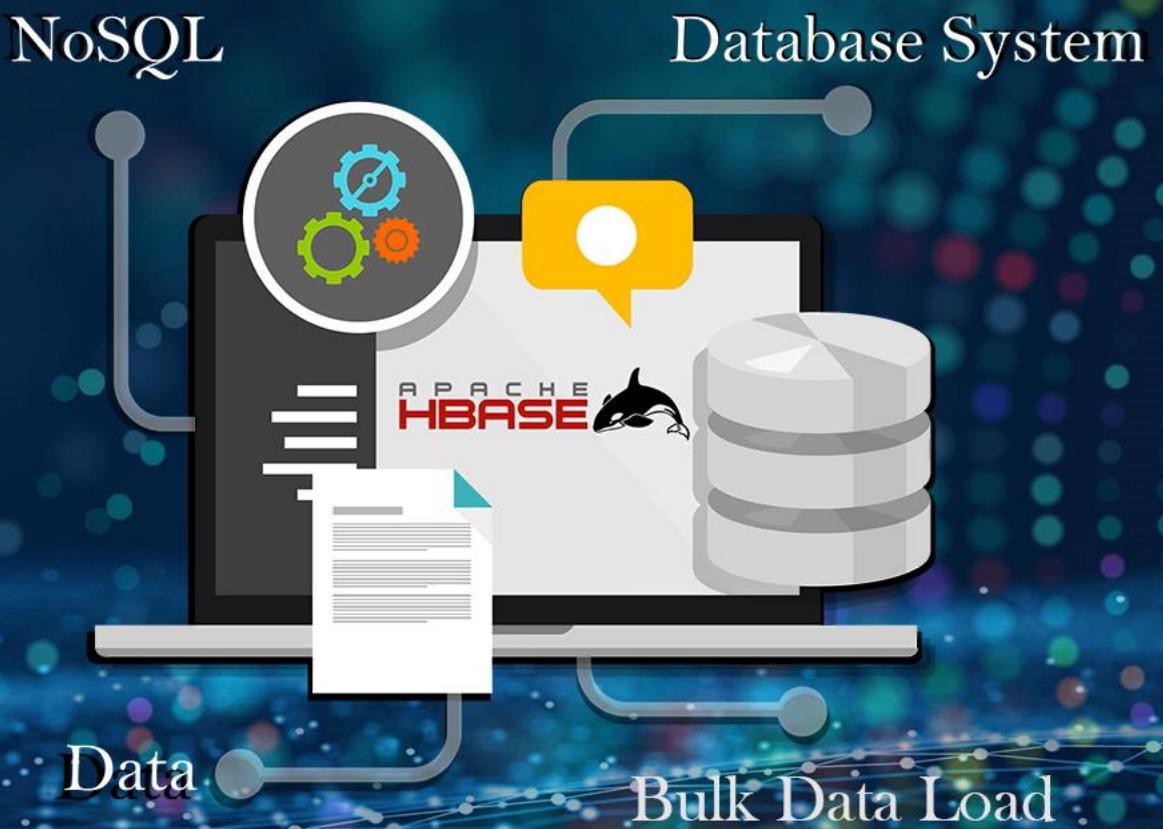
HBase Read or Write

There is a special HBase Catalog table called the META table, which holds the location of the regions in the cluster

Here is what happens the first time a client reads or writes data to HBase



HBase Meta Table

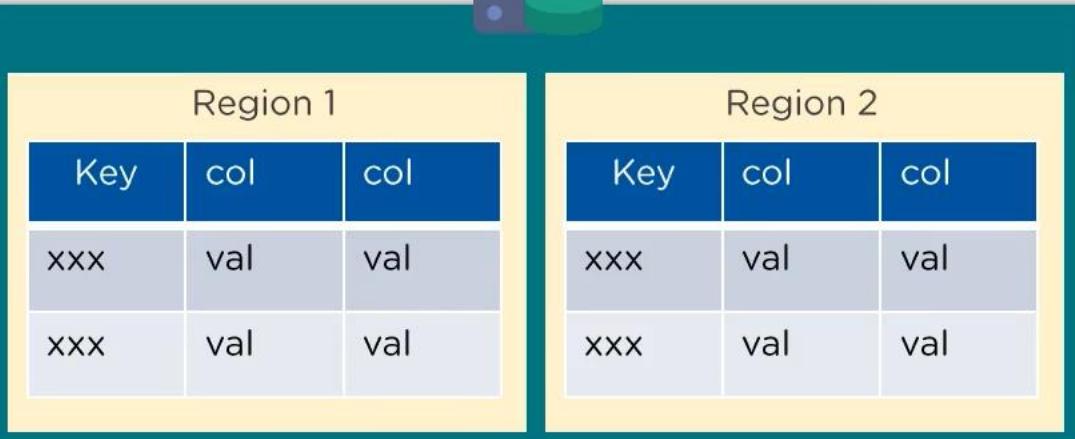


HBase Meta Table

Special HBase catalog table that maintains a list of all the Region Servers in the HBase storage system

Meta Table	
Row key	value
table, key, region	region server

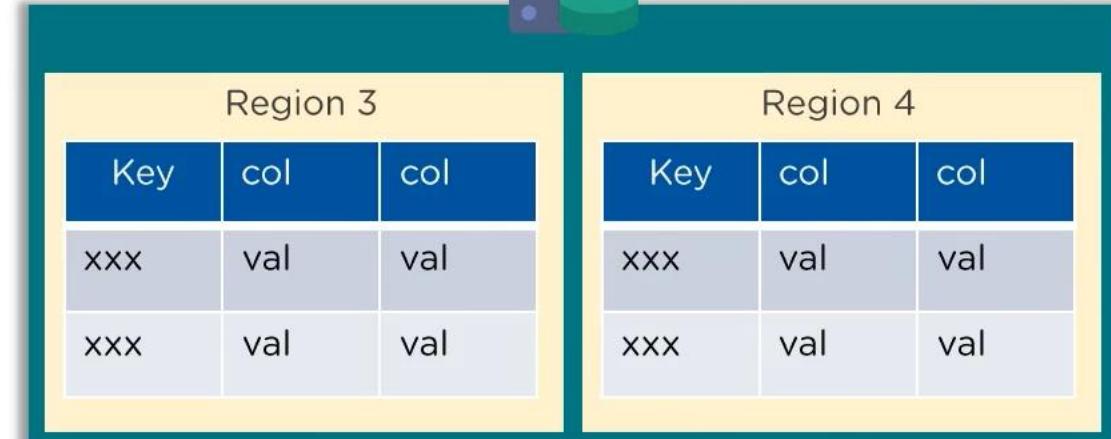
META table is used to find the Region for a given Table key



Region Server

Region 1		
Key	col	col
xxx	val	val
xxx	val	val

Region 2		
Key	col	col
xxx	val	val
xxx	val	val

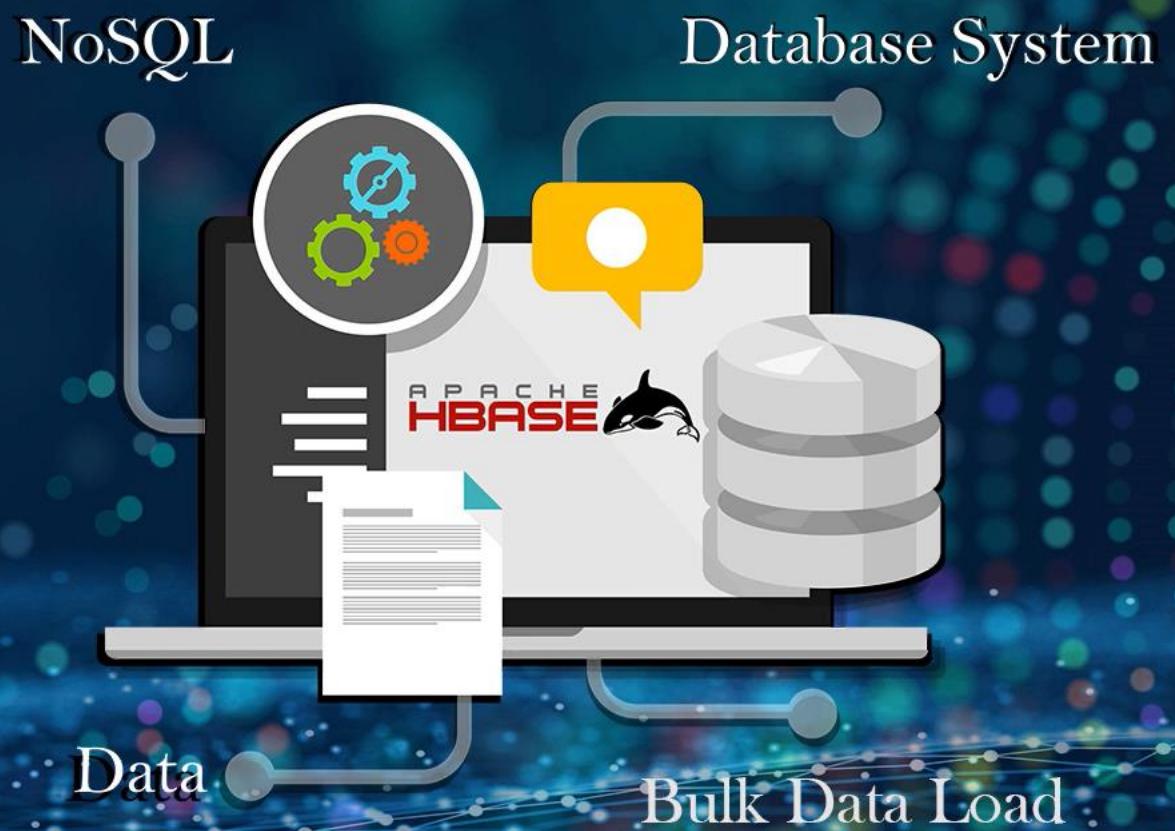


Region Server

Region 3		
Key	col	col
xxx	val	val
xxx	val	val

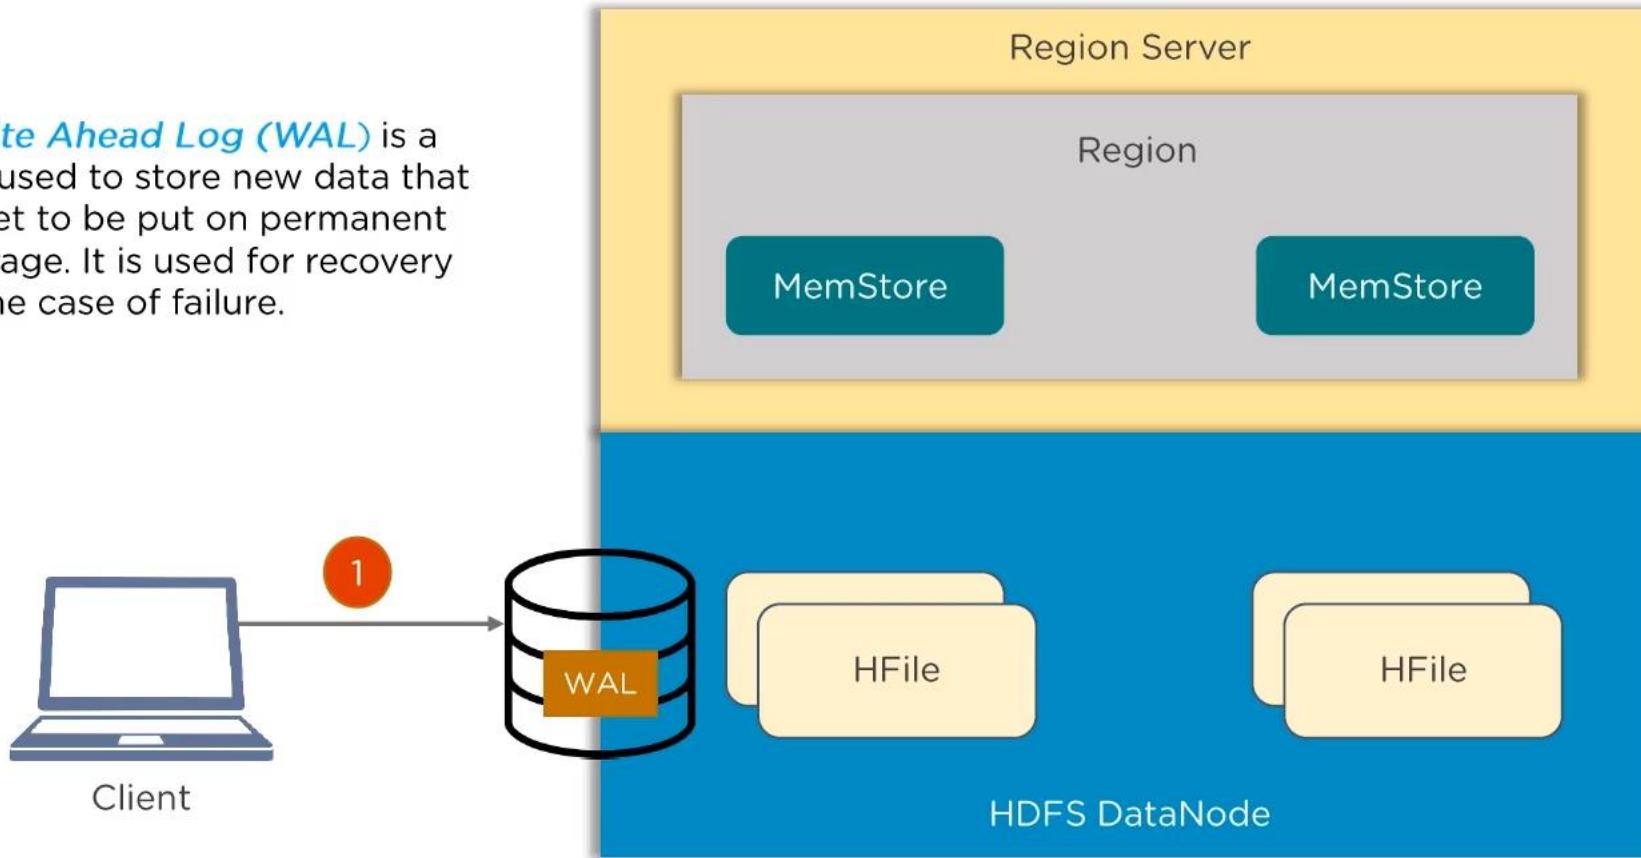
Region 4		
Key	col	col
xxx	val	val
xxx	val	val

HBase Write Mechanism



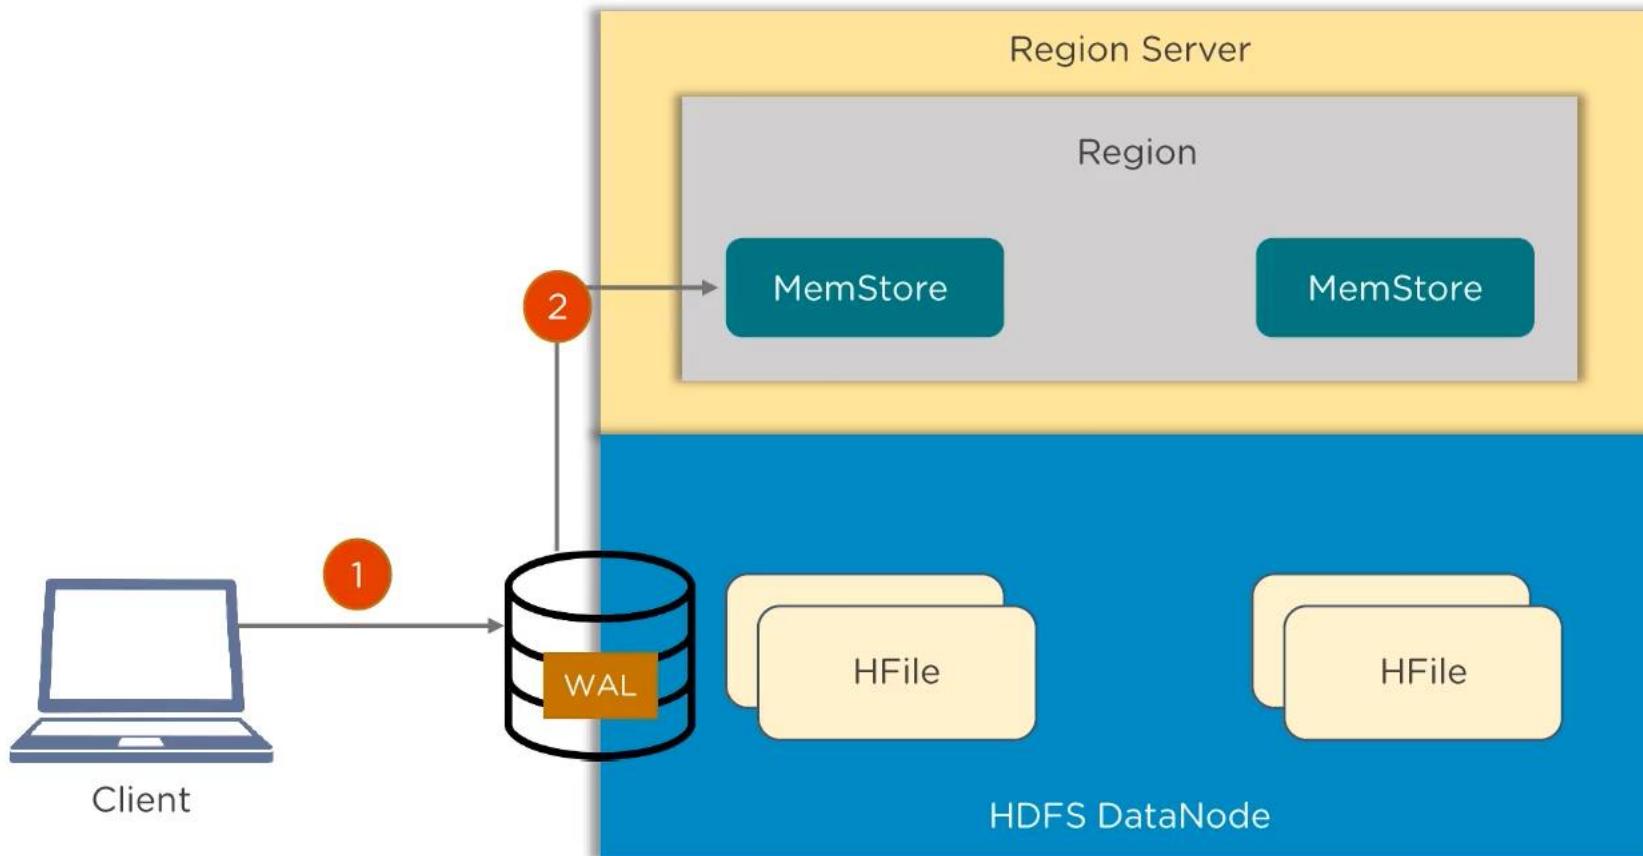
HBase Write Mechanism

Write Ahead Log (WAL) is a file used to store new data that is yet to be put on permanent storage. It is used for recovery in the case of failure.



- 1 When client issues a **put** request, it will write the data to the write-ahead log (WAL)

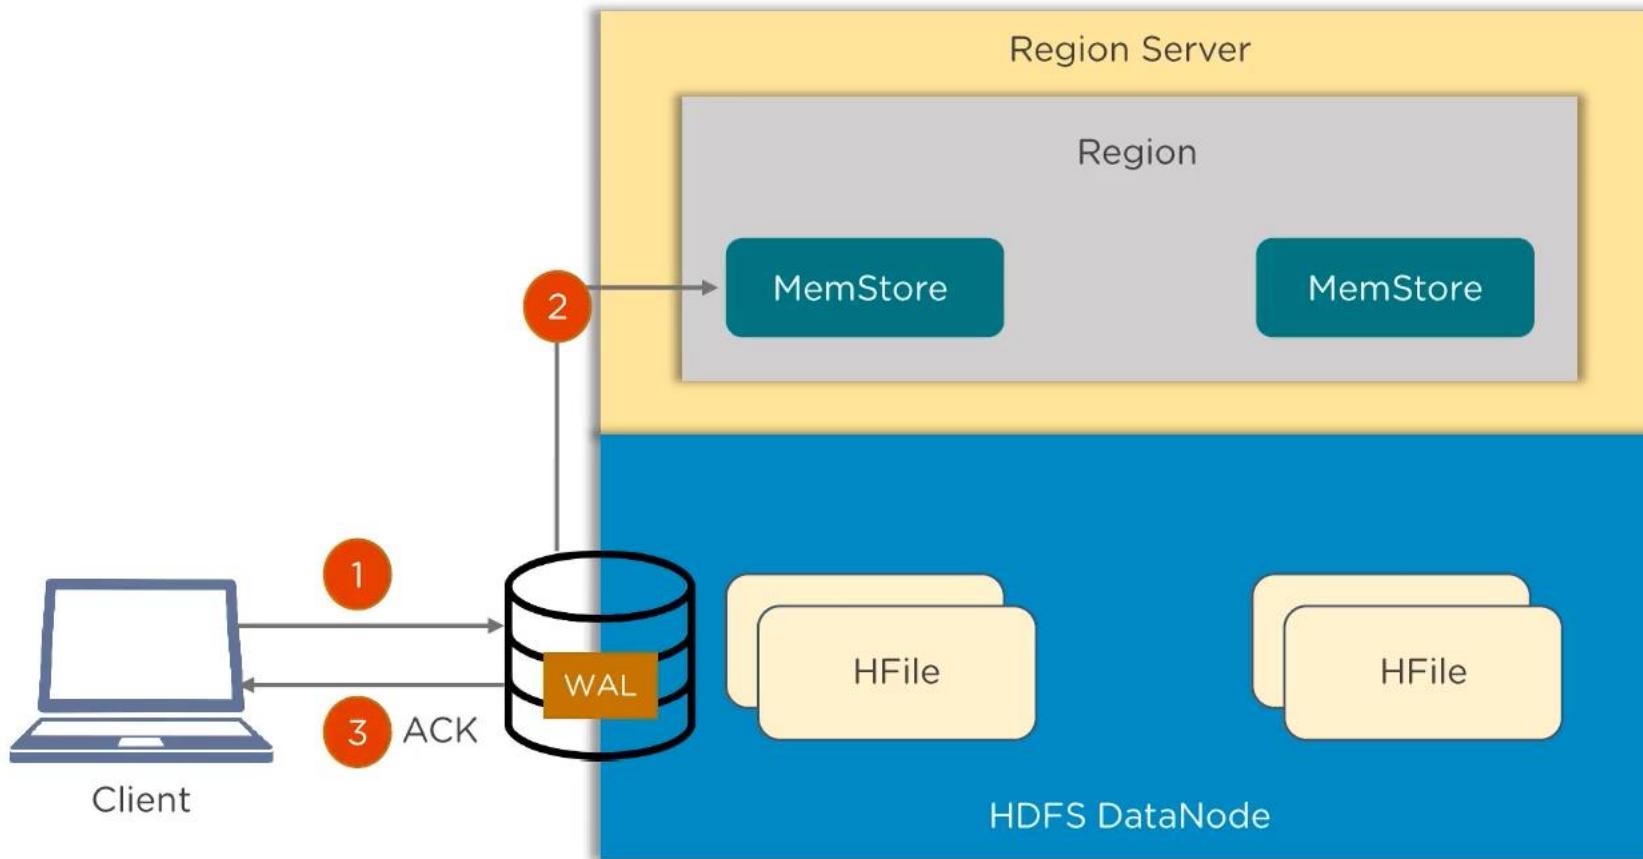
HBase Write Mechanism



MemStore is the write cache that stores new data that has not yet been written to disk. There is one MemStore per column family per region.

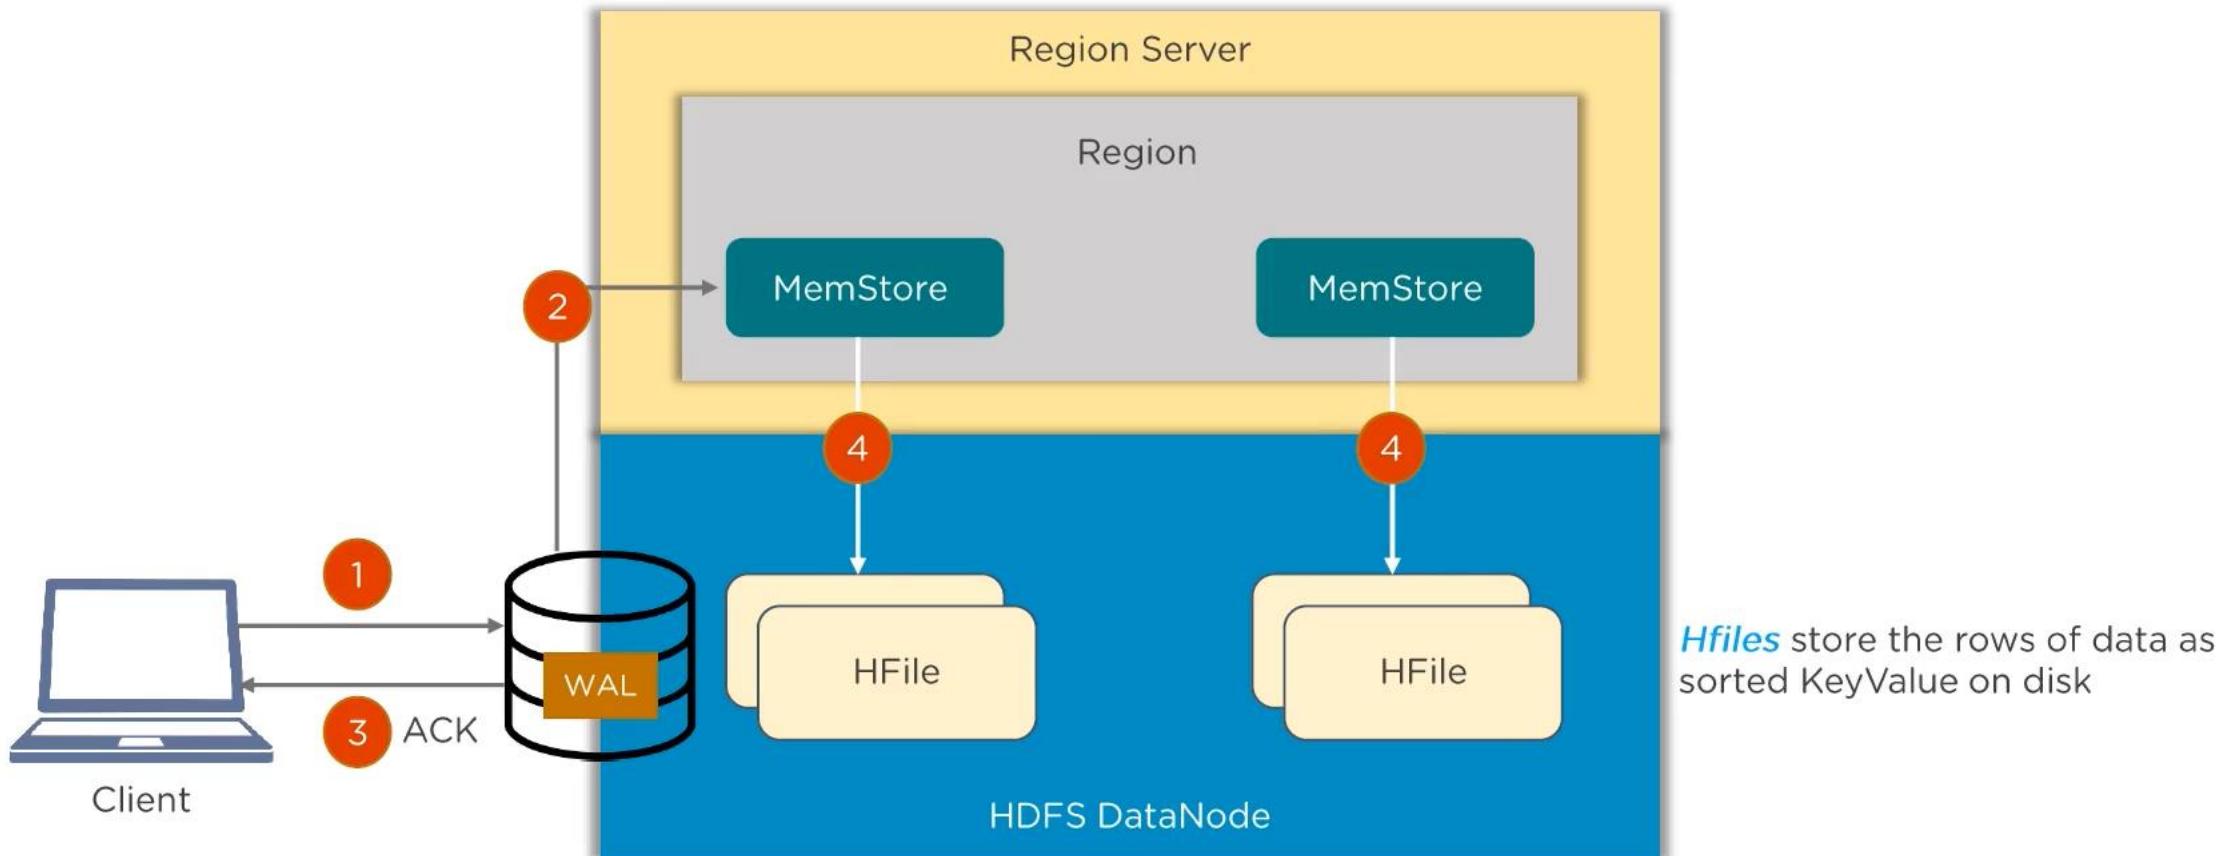
- Once data is written to the WAL, it is then copied to the **MemStore**

HBase Write Mechanism



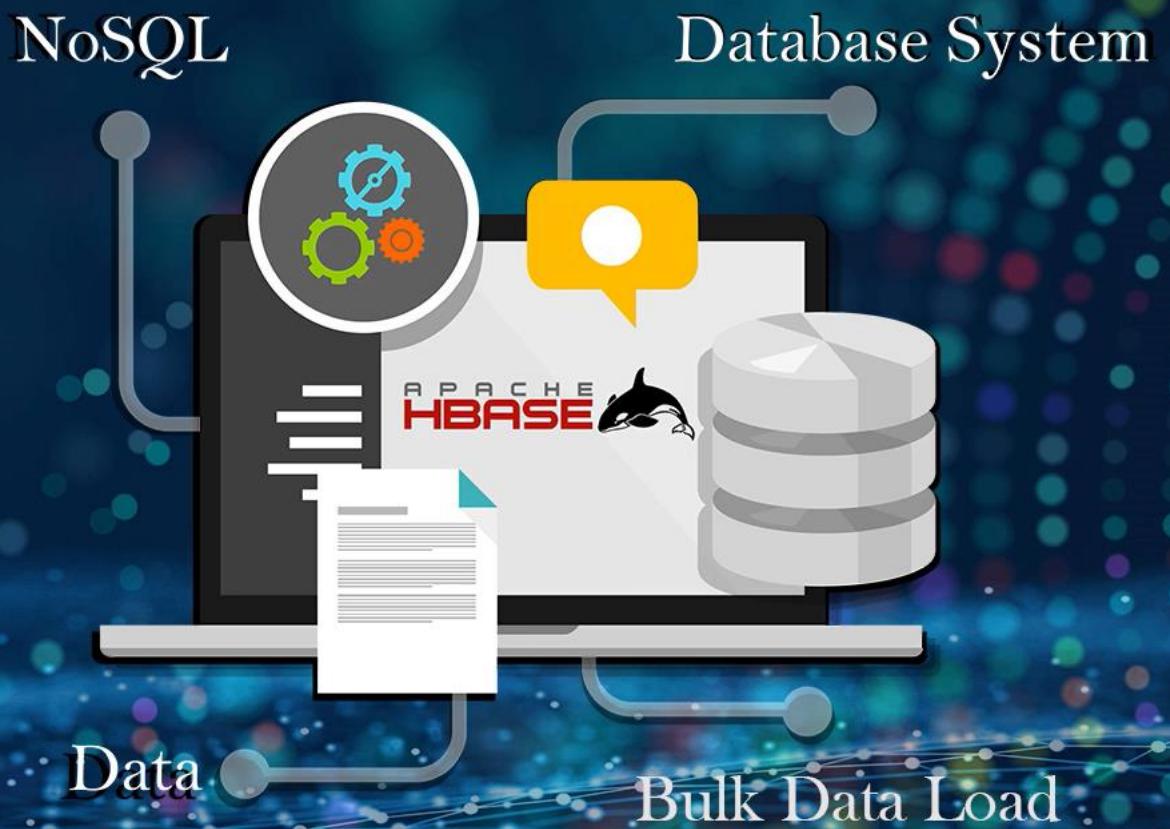
- 3 Once the data is placed in MemStore, the client then receives the **acknowledgment**

HBase Write Mechanism



- 4 When the MemStore reaches the threshold, it dumps or commits the data into a **HFile**

Demo on HBase



HBase - General Commands

- **Status**
This command returns the status of the system including the details of the servers running on the system.
- **Version**
This command returns the version of HBase used in your system
- **Table_help**
This command guides you what and how to use table-referenced commands.
- **Whoami**
This command returns the user details of HBase. If you execute this command, returns the current HBase user

Demo on HBase

Open cmd as Admin and write next commands:

```
cd \
cd hbase-2.3.7\bin
start-hbase.cmd
hbase shell
```

Demo on HBase

After a couple of seconds, you'll be inside the HBase shell where you can type the HBase commands. You can start off by typing the following commands:

```
list // Lists down all the tables present in HBase  
create 'newtbl', 'knowledge' // Creates a new table  
list // Lists down all the tables present in HBase  
describe 'newtbl' // Checks if the table was created  
status 'summary' // Checks the status of HBase
```

Demo on HBase

Now that we have created a new table, let's put some data into it.

```
put 'newtbl', 'r1', 'knowledge:sports', 'cricket'  
put 'newtbl', 'r1', 'knowledge:sciense', 'chemistry'  
put 'newtbl', 'r1', 'knowledge:sciense', 'physics'  
put 'newtbl', 'r2', 'knowledge:economics', 'macro economics'  
put 'newtbl', 'r2', 'knowledge:music', 'hard rock'
```

Let's now list the contents of the table by typing:

```
scan 'newtbl' //
```

We cannot see “chemistry” as it will only display the latest update which in this case is “physics”

Demo on HBase

Now we can type the following commands:

is_enabled 'newtbl' // Checks if the table is enabled

disable 'newtbl' // Disables the table

scan 'newtbl' // Lists the contents of the table.

Note that this will throw an error as the table is disabled.

Now before we move and enable the table,
let's do an alteration on it.

alter 'newtbl', 'test_info' // Updates column family in the table

enable 'newtbl' // Enables the table

describe 'newtbl' //Checks the column families after updating

scan 'newtbl'

Demo on HBase

Then extract values for one particular row and also see how to add new information to a row by using the following commands:

```
get 'newtbl', 'r1' // Extracts the values for r1 in the table
```

```
put 'newtbl', 'r1', 'knowledge:economics', 'macro economics' // Adds  
new  
information to r1 for  
economics.
```

Note that this will update the table but will not override the information

```
get 'newtbl', 'r1' // Displays the results for r1
```

Demo on HBase

Using the **delete** command, you can delete a specific cell in a table.

```
scan 'newtbl'  
delete 'newtbl', 'r2', 'knowledge:music' // Here is an example  
to  
cell  
deleteall 'newtbl', 'r1' // You can delete all the cells in a row  
scan 'newtbl'
```