



Introduction to Hadoop Yarn

BIG DATA
HADOOP & SPARK

What is YARN?

YARN= Yet Another Resource Negotiator



YARN is a resource manager



Created by separating the processing engine and the management function of MapReduce



Monitors and manages workloads, maintains a multi-tenant environment, manages the high availability features of Hadoop, and implements security controls

Need for YARN

Before 2012

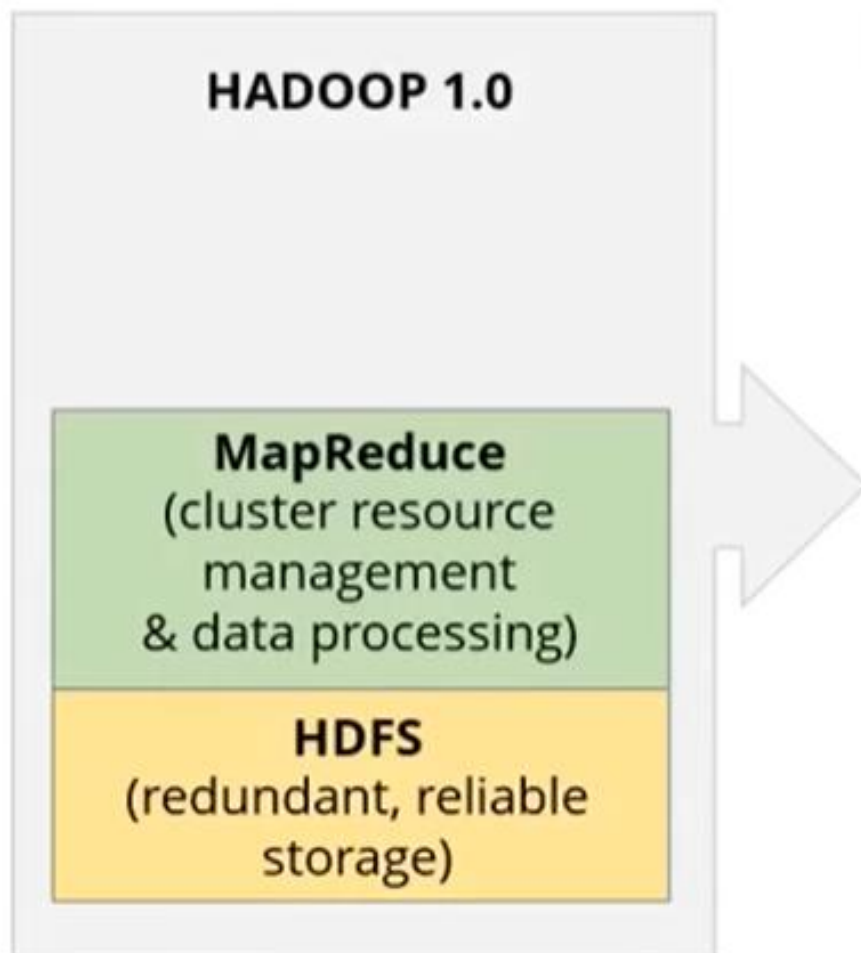
Users could write MapReduce programs using scripting languages



Need for YARN

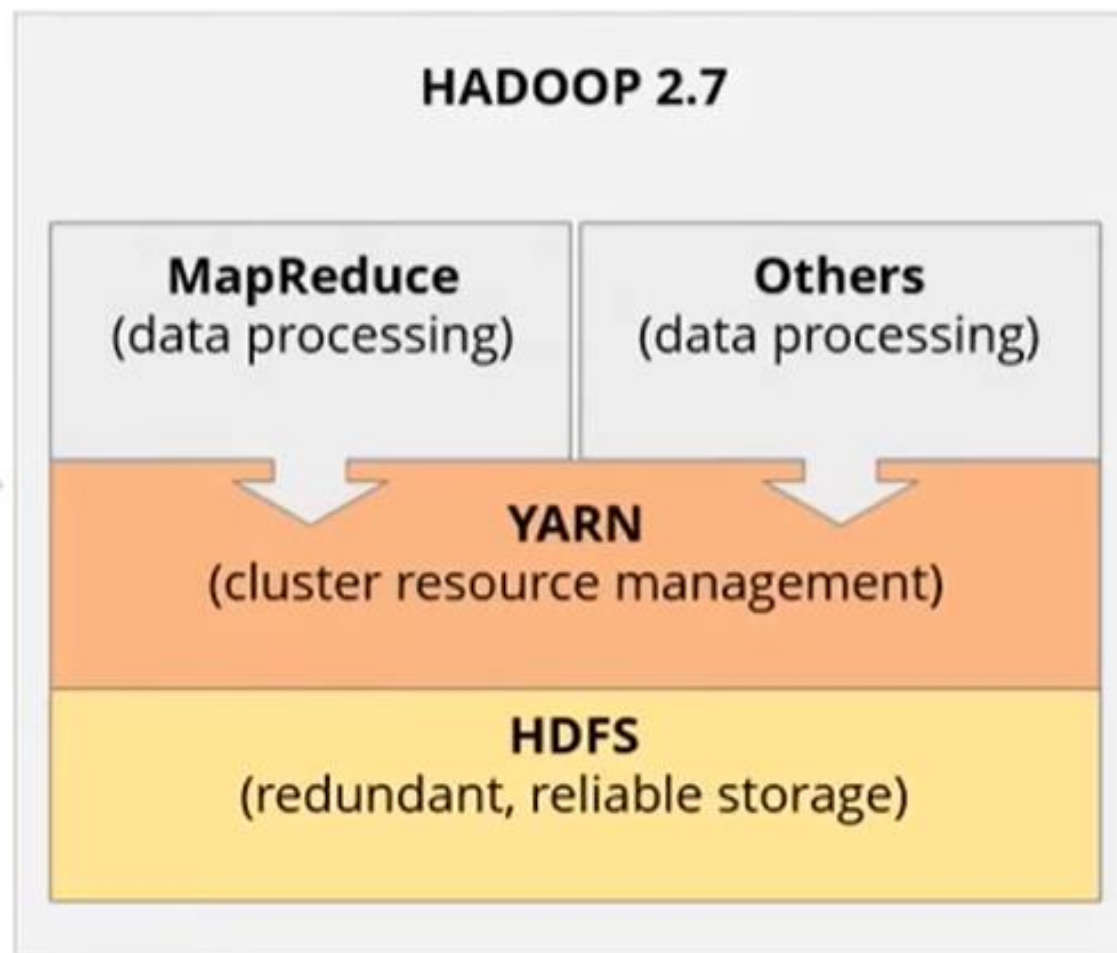
Before 2012

Users could write MapReduce programs using scripting languages



Since 2012

Users could work on multiple processing models in addition to MapReduce



YARN—Use Case

YAHOO!

- Yahoo was the first company to embrace Hadoop in a big way, and it is a trendsetter within the Hadoop ecosystem. In late 2012, it struggled to handle iterative and stream processing of data on Hadoop infrastructure due to MapReduce limitations.
- After implementing YARN in the first quarter of 2013, Yahoo has installed more than 30,000 production nodes on
 - Spark for iterative processing
 - Storm for stream processing
 - Hadoop for batch processing

YARN—Advantages

The single-cluster approach provides a number of advantages, including:

There's no need to move data between Hadoop YARN and systems running on different clusters of computers

Reduced data motion

Higher cluster utilization

Resources unutilized by a framework can be consumed by another

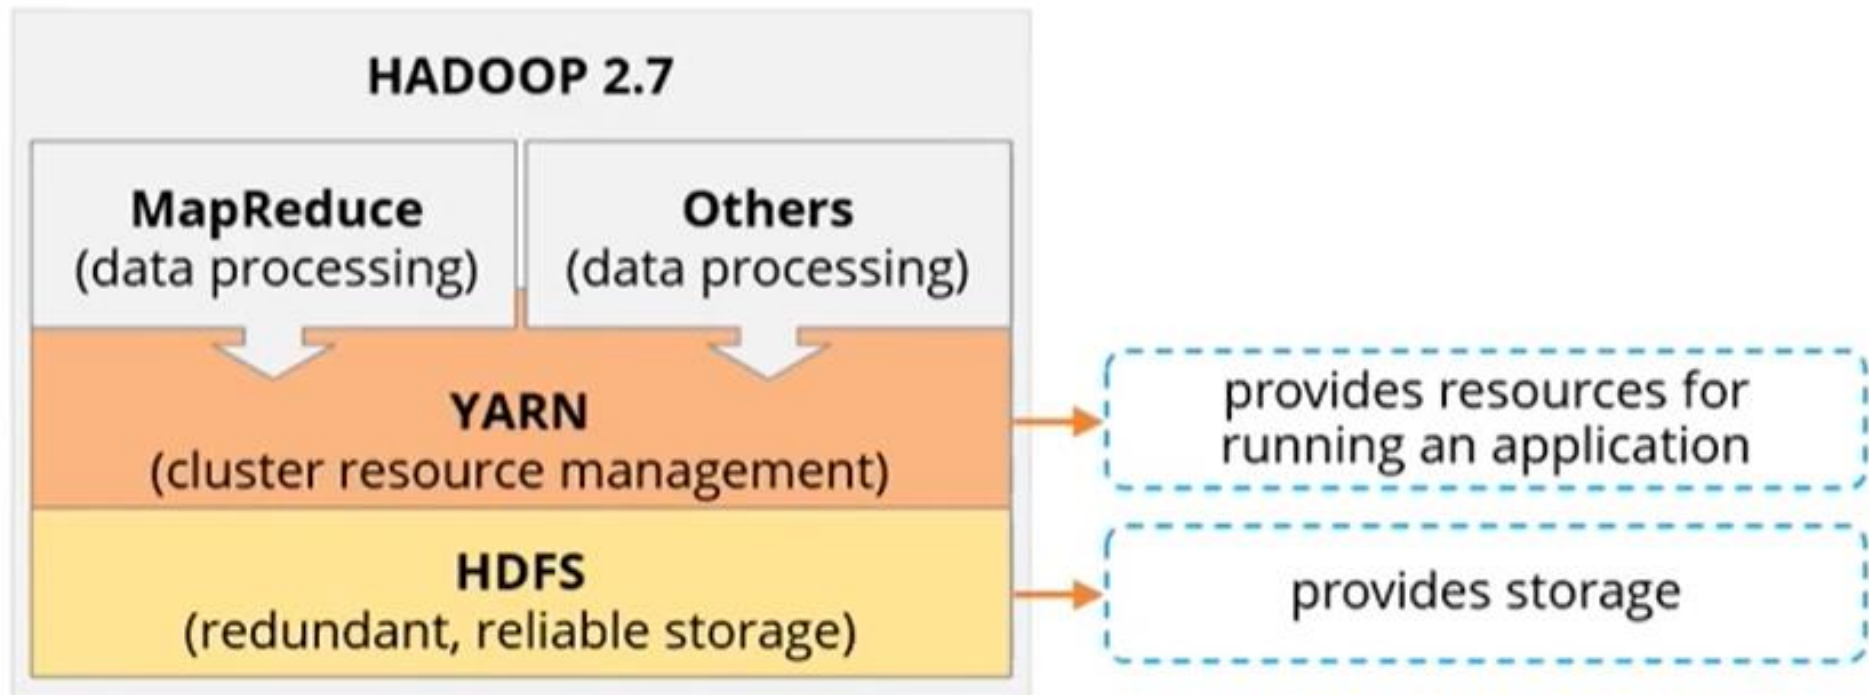
Advantages of the single-cluster approach

Lower operational costs

Only one "do-it-all" cluster needs to be managed

YARN Infrastructure

The YARN Infrastructure is responsible for providing computational resources for application executions.

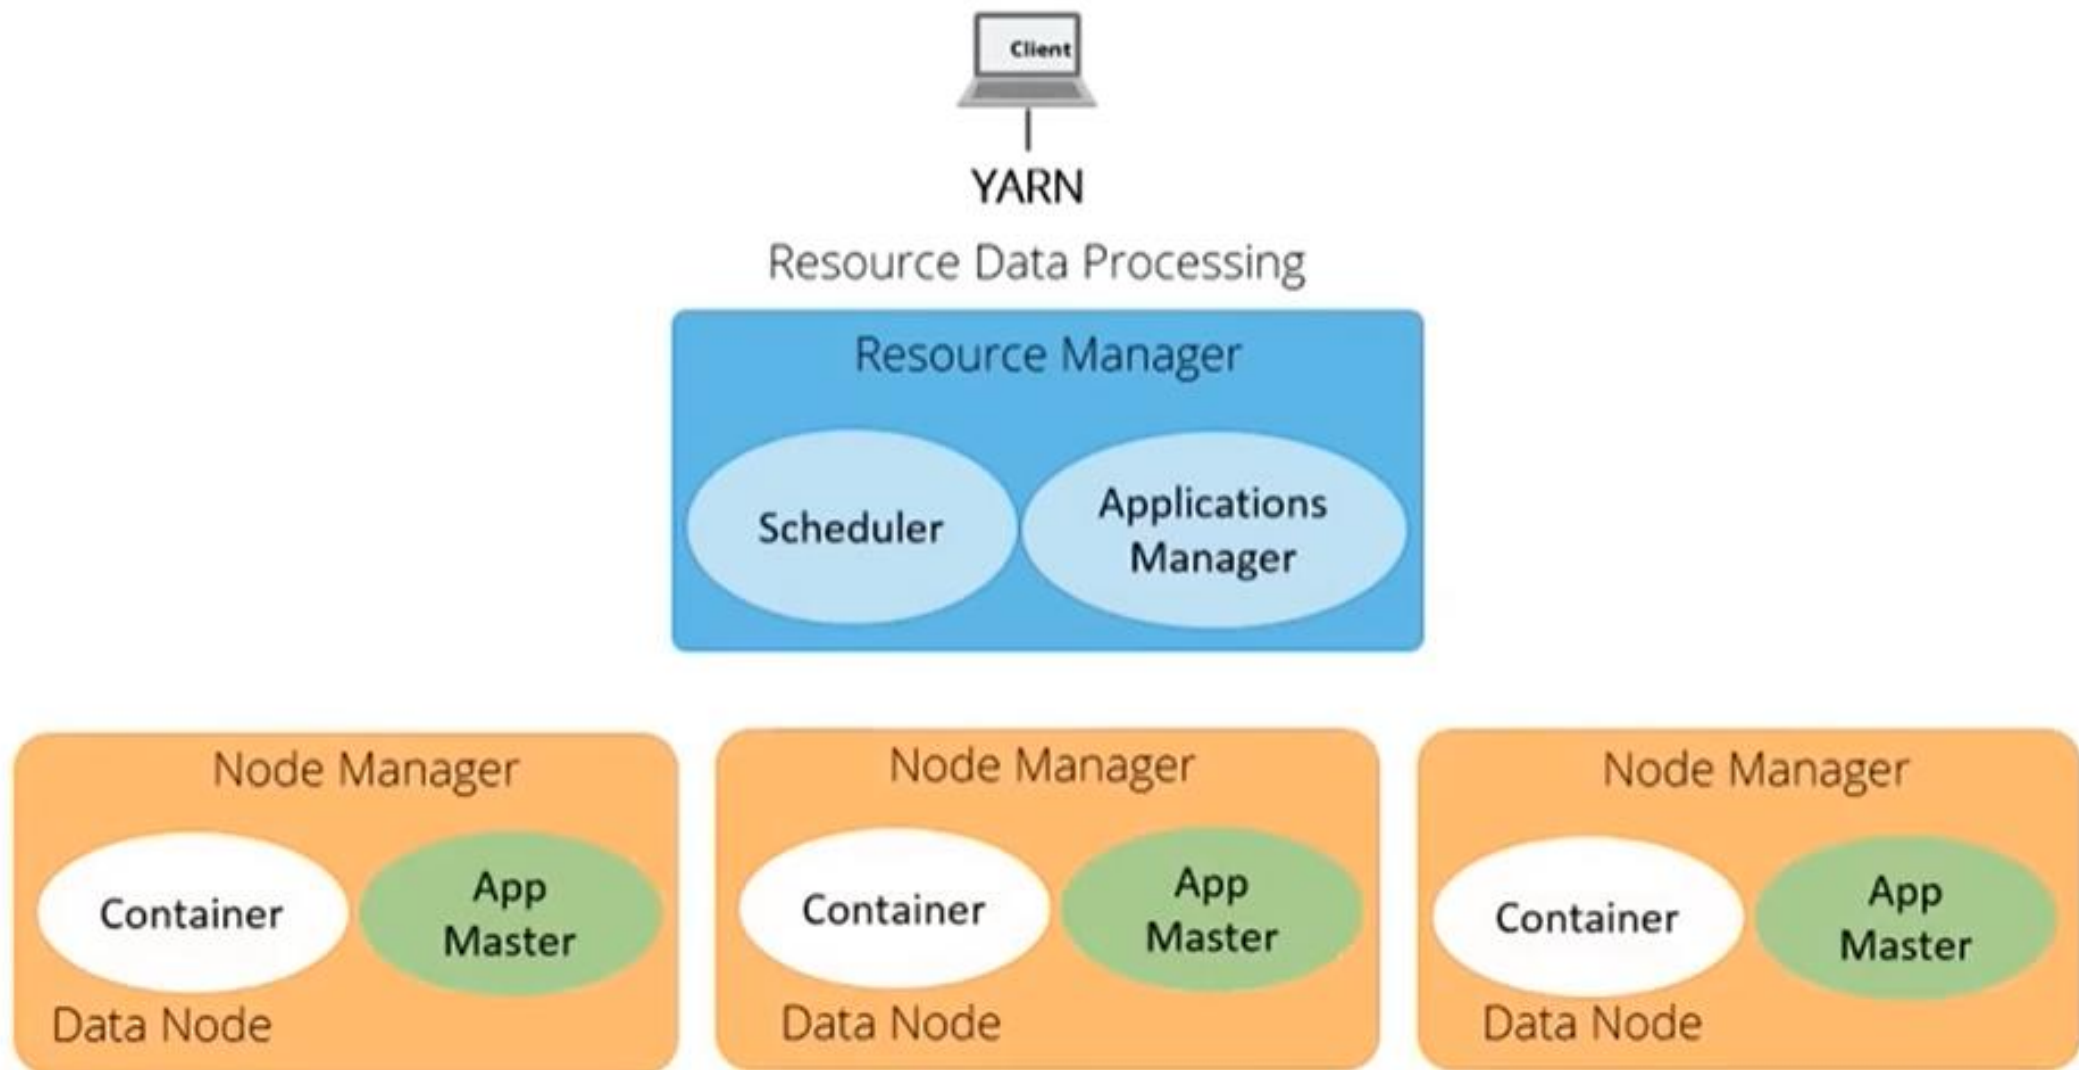


HDFS and YARN

YARN and its Architecture

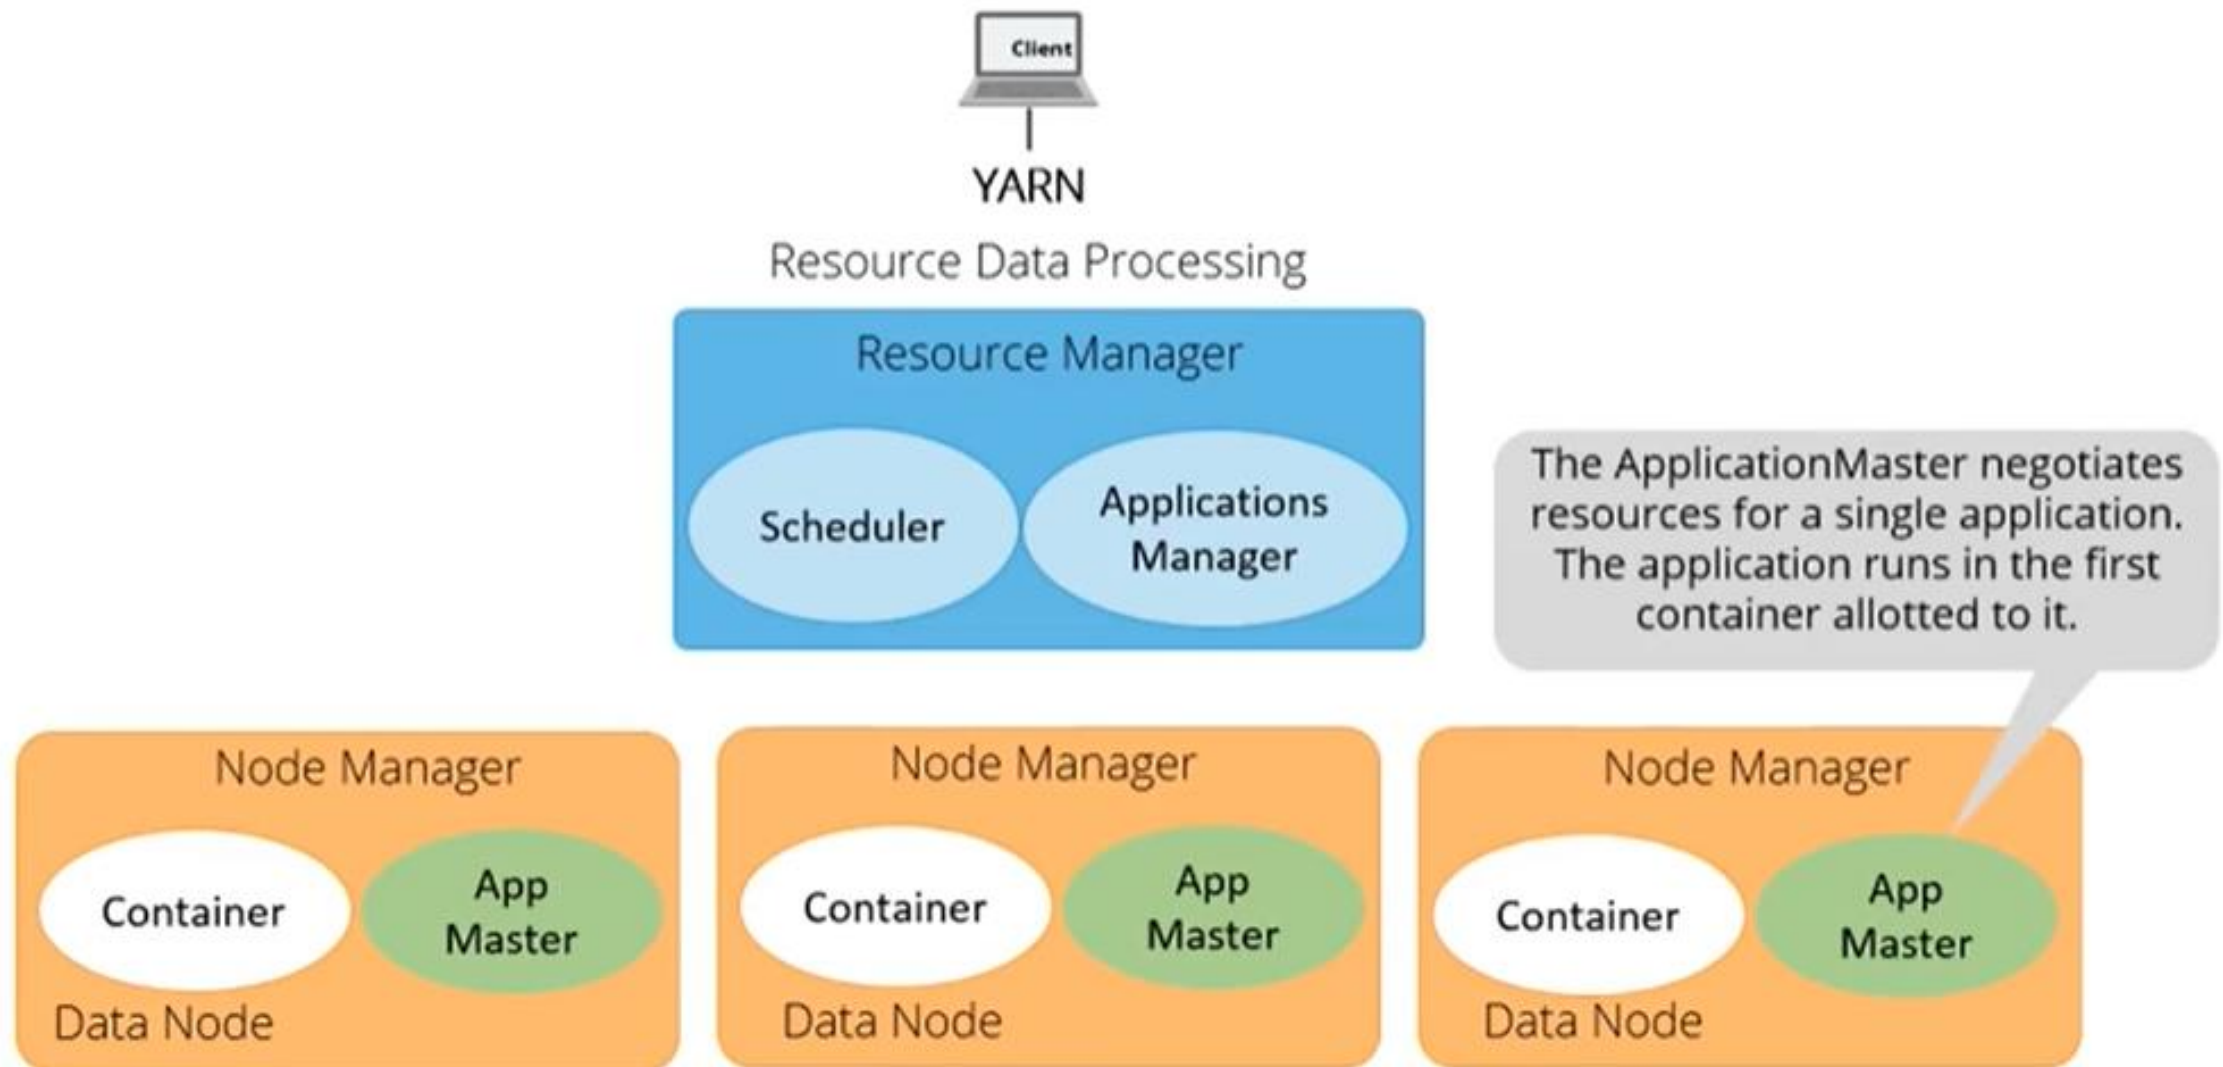
Three Elements of YARN Architecture

The three important elements of the YARN architecture are the Resource Manager, Application Master, and Node Manager.



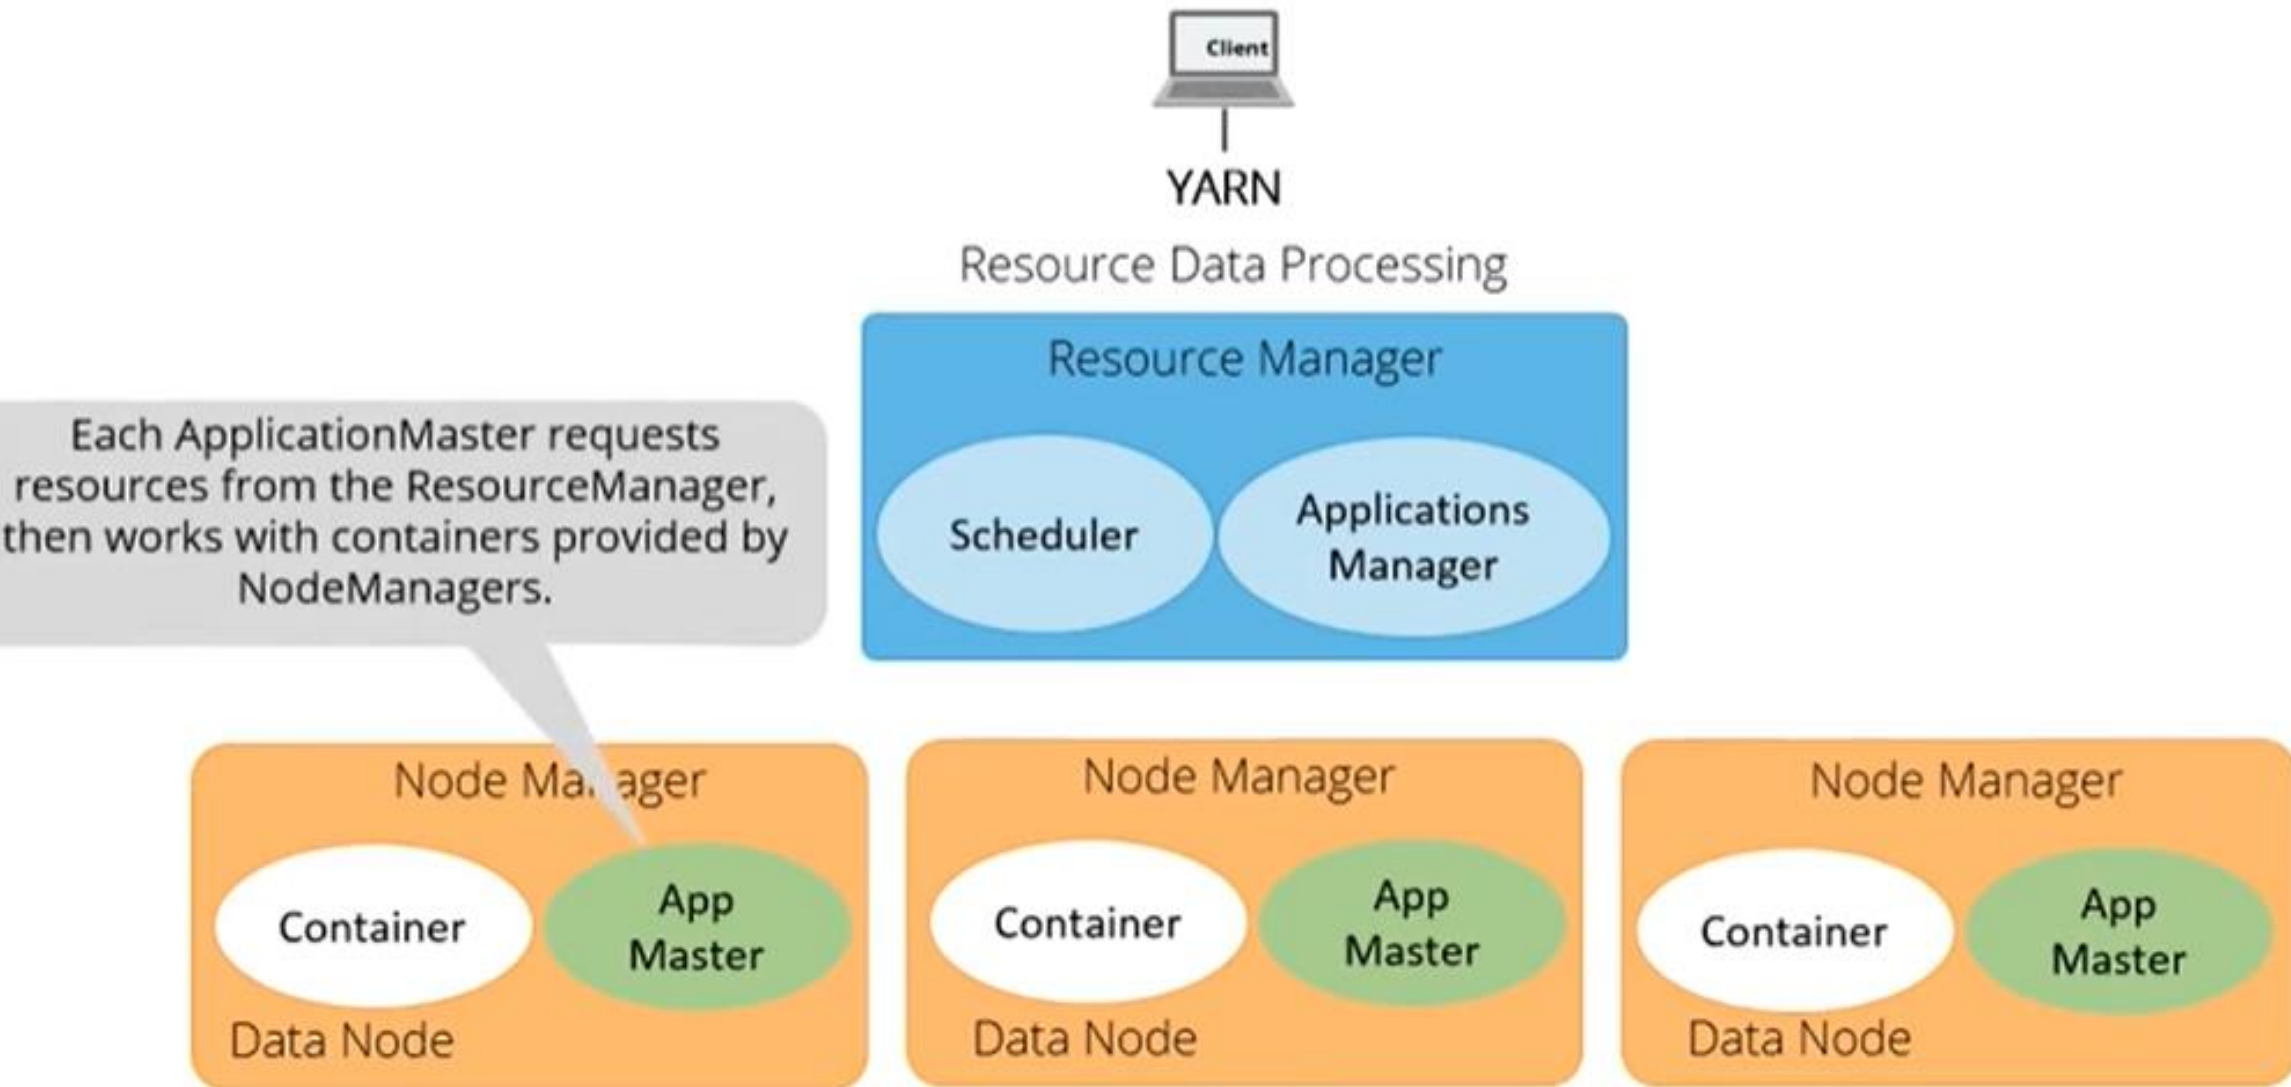
Three Elements of YARN Architecture

The three important elements of the YARN architecture are the Resource Manager, Application Master, and Node Manager.



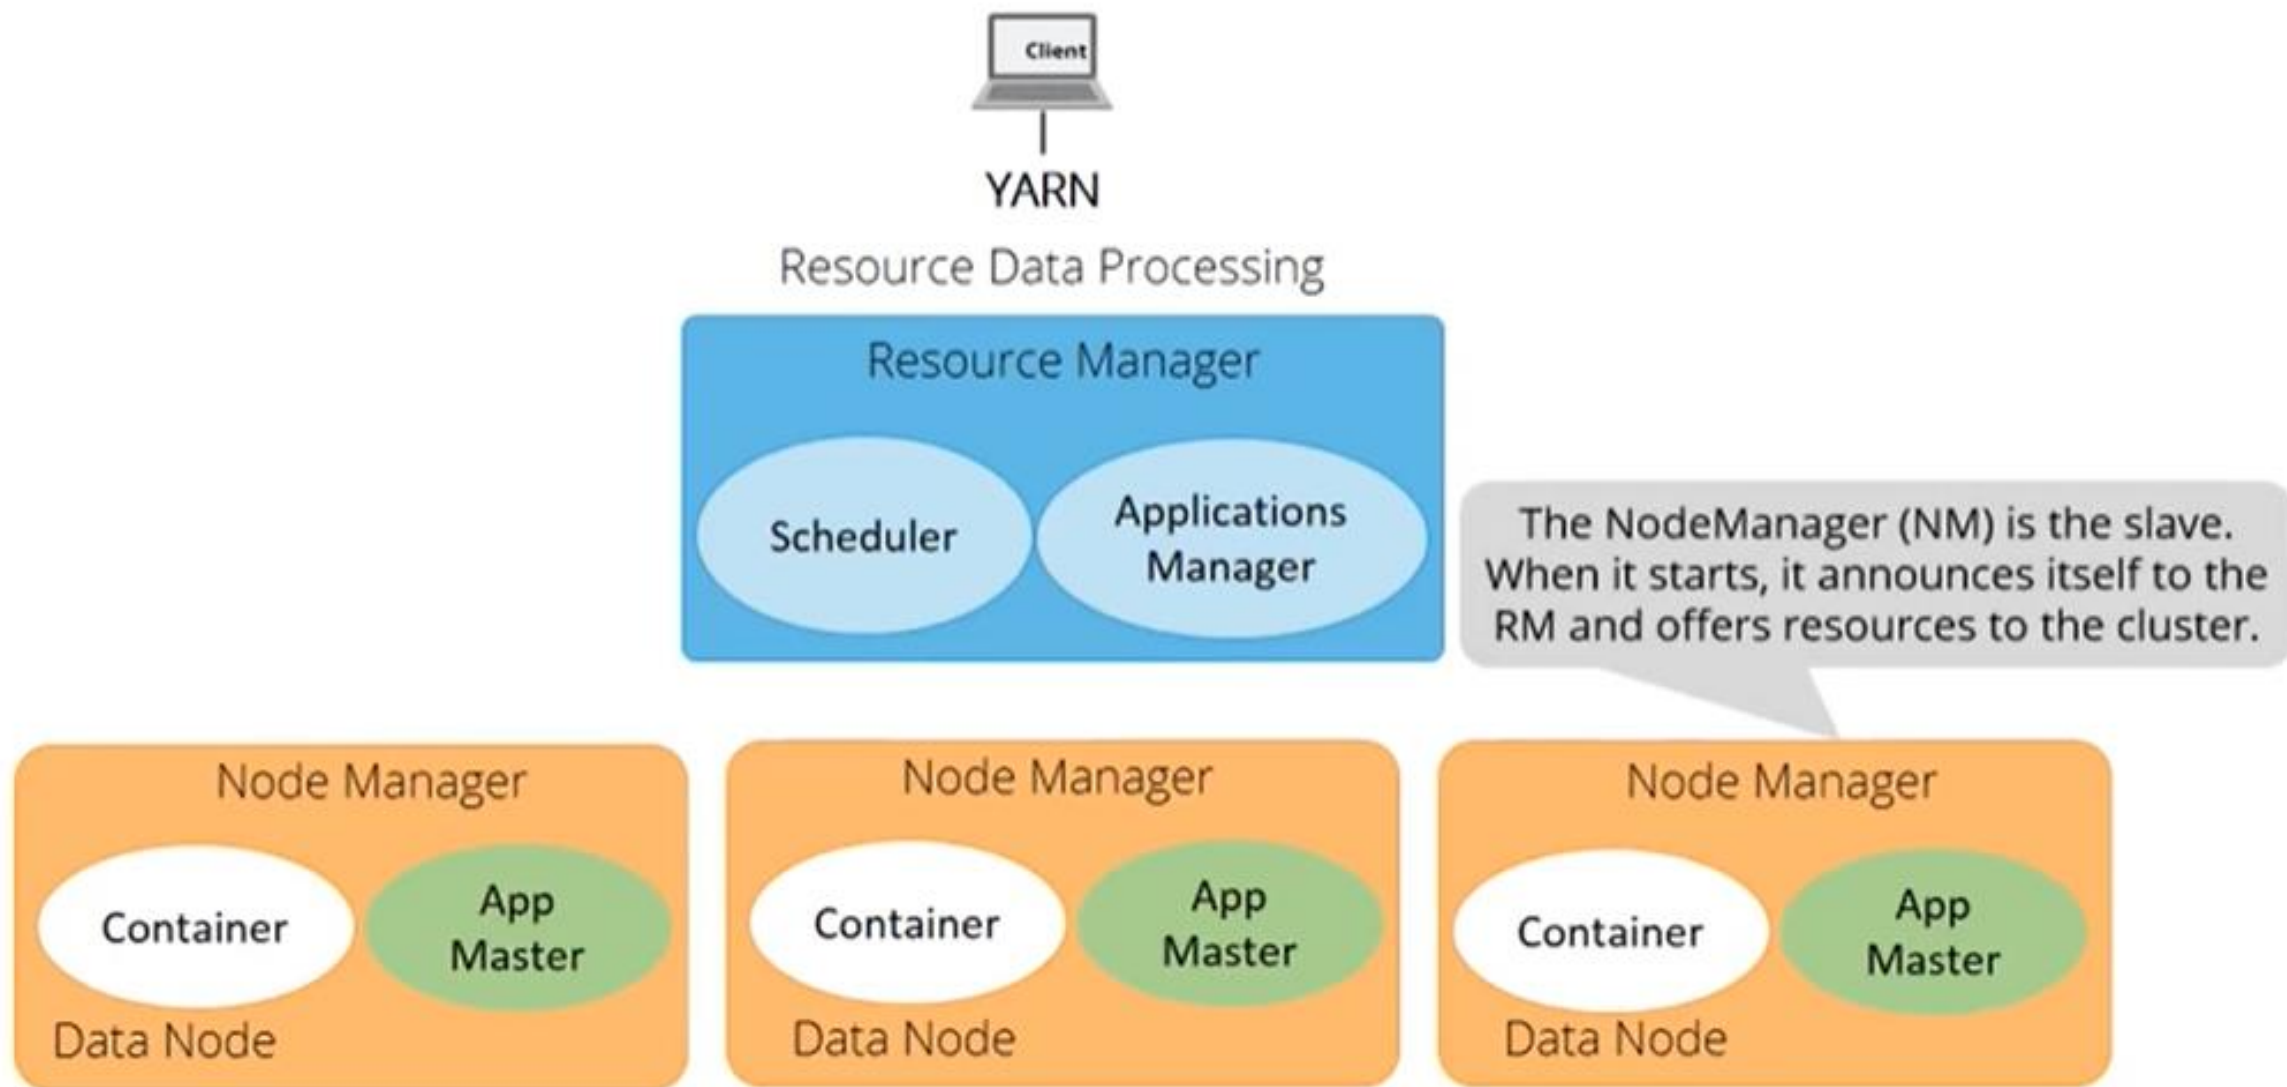
Three Elements of YARN Architecture

The three important elements of the YARN architecture are the ResourceManager, ApplicationMaster, and NodeManager.



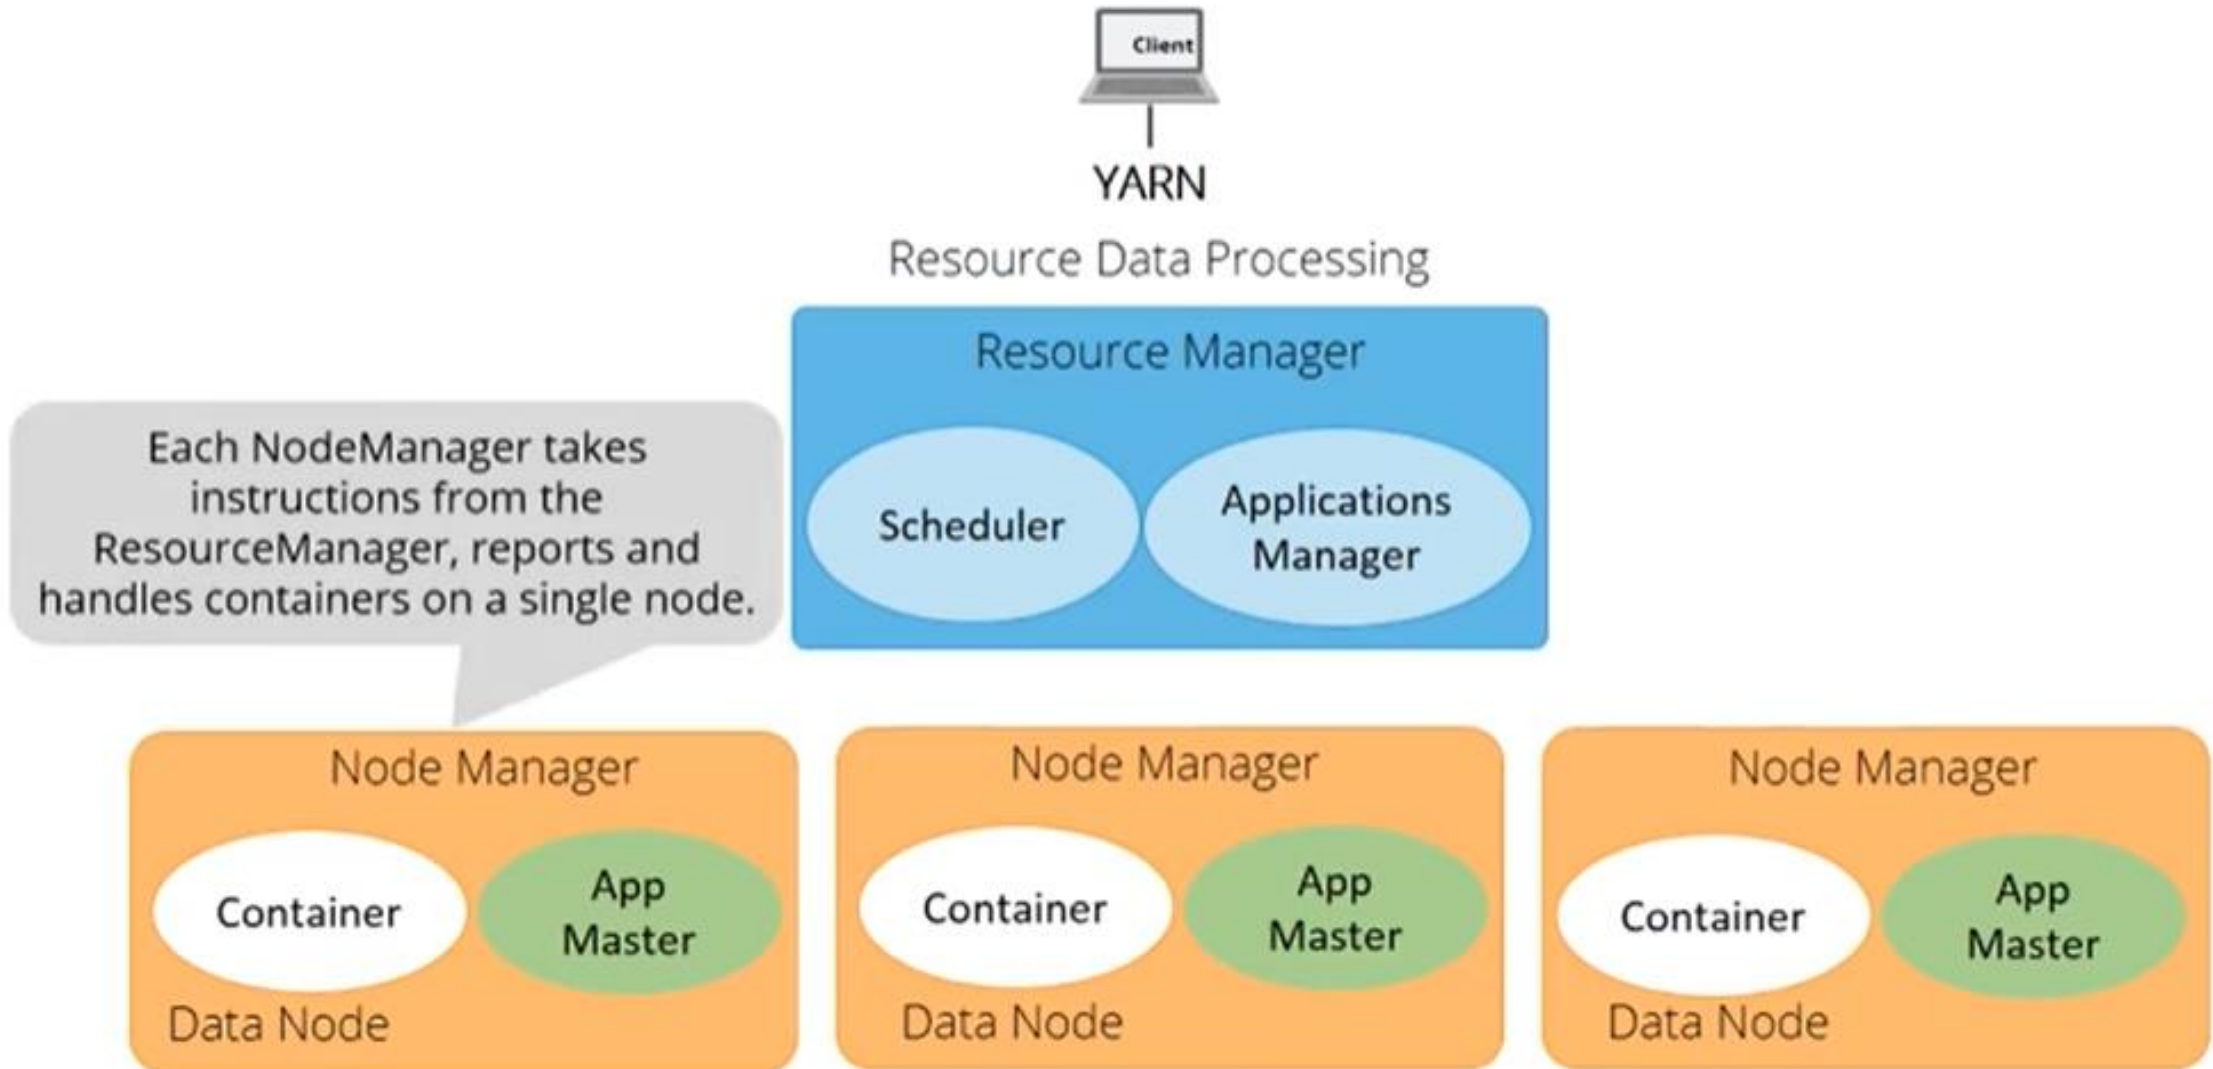
Three Elements of YARN Architecture

The three important elements of the YARN architecture are the Resource Manager, Application Master, and Node Manager.



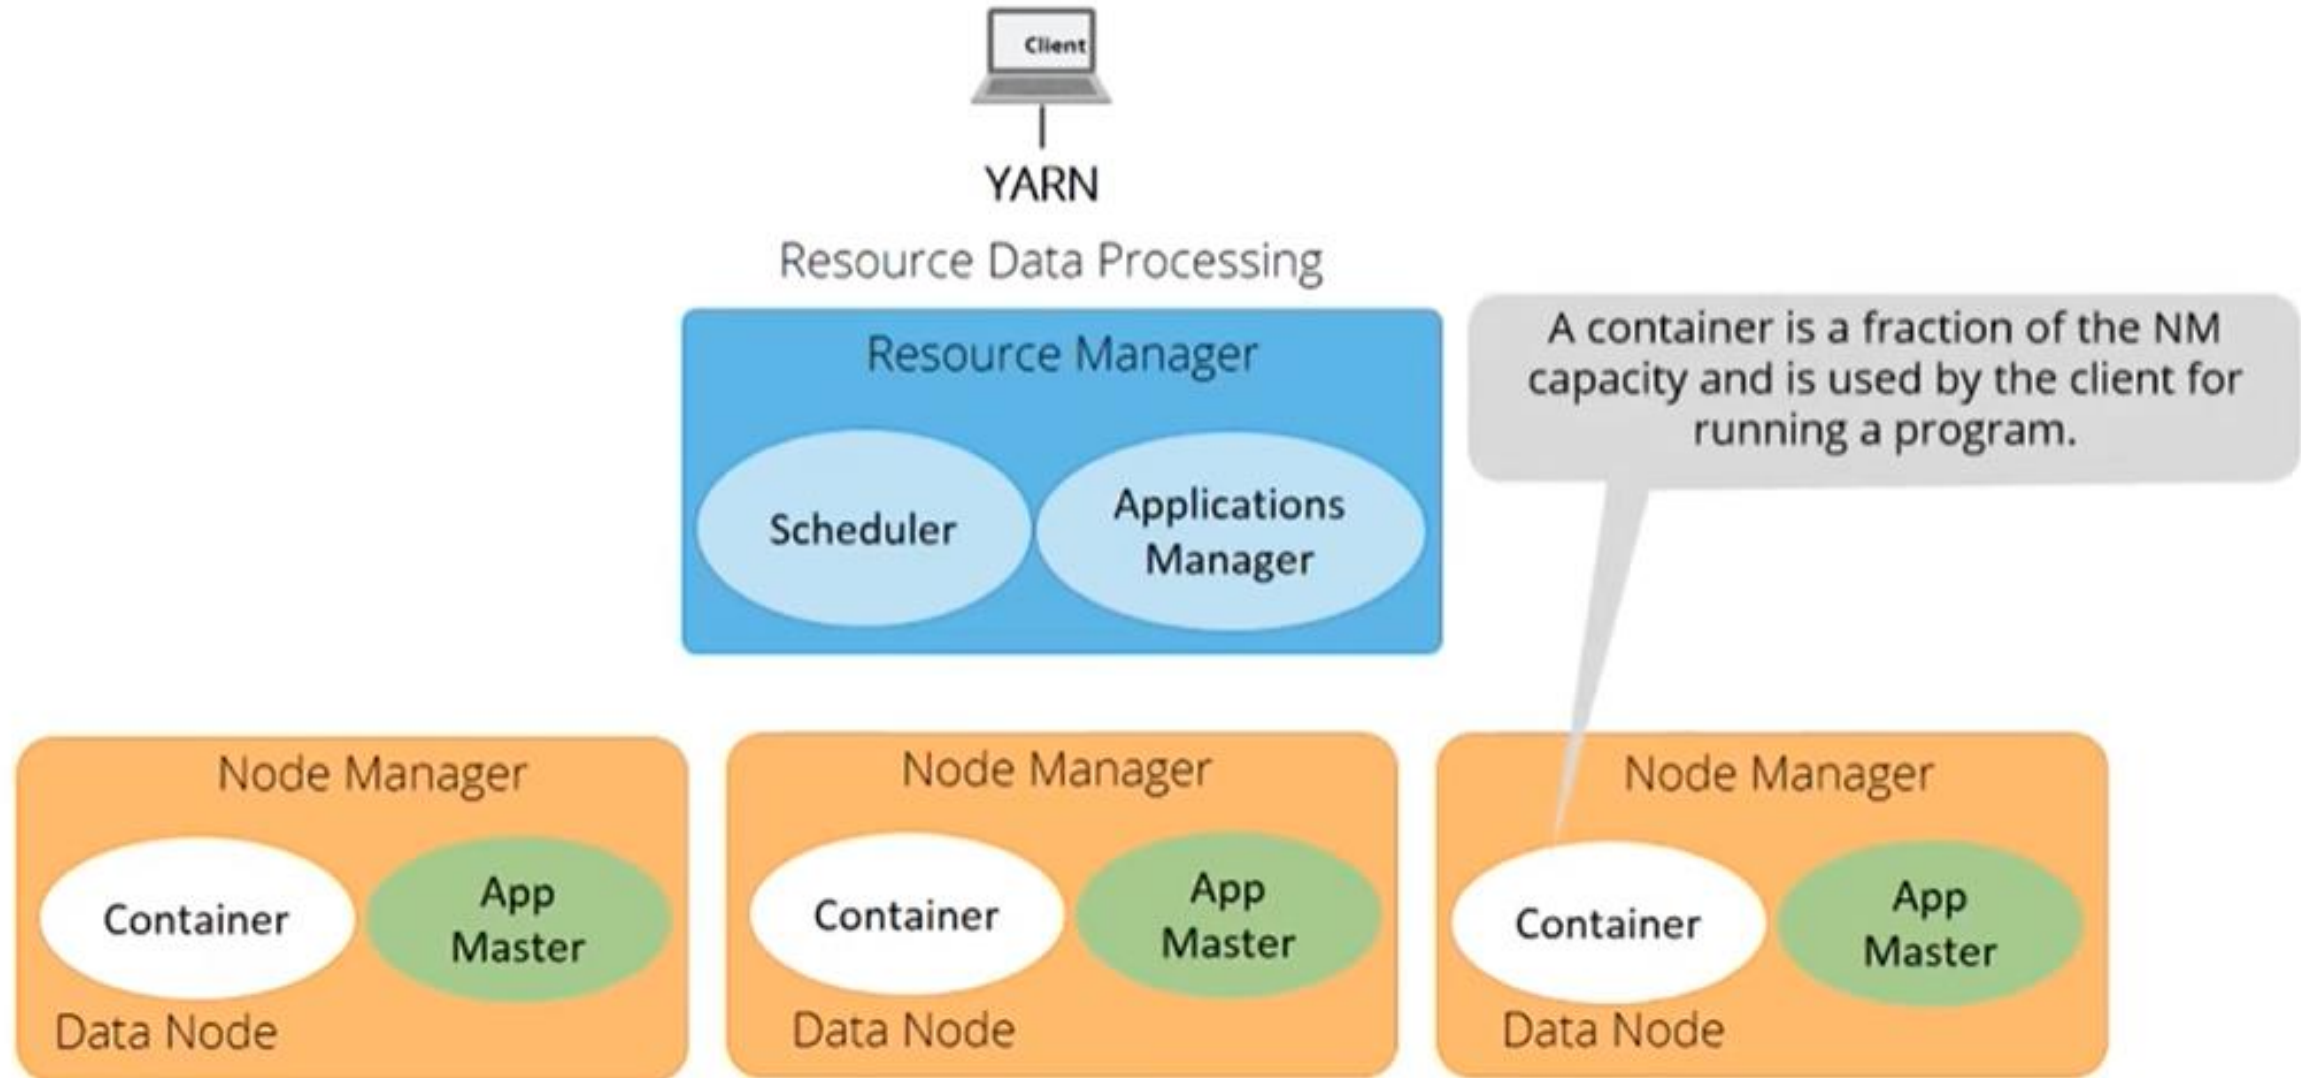
Three Elements of YARN Architecture

The three important elements of the YARN architecture are the ResourceManager, ApplicationMaster, and NodeManager.




Three Elements of YARN Architecture

The three important elements of the YARN architecture are the Resource Manager, Application Master, and Node Manager.



YARN Architecture Element—ResourceManager

The RM mediates the available resources in the cluster among competing applications—to maximum cluster utilization.




It includes a pluggable scheduler called the YarnScheduler, which allows different policies for managing constraints such as capacity, fairness, and Service Level Agreements.

ResourceManager Component—Scheduler

The Scheduler is responsible for allocating resources to various running applications.


The Scheduler does not monitor or track the status of the application; nor does it restart failed tasks.

The Scheduler has a policy plugin to partition cluster resources among various applications. Examples: CapacityScheduler, FairScheduler.



ResourceManager Component—ApplicationManager

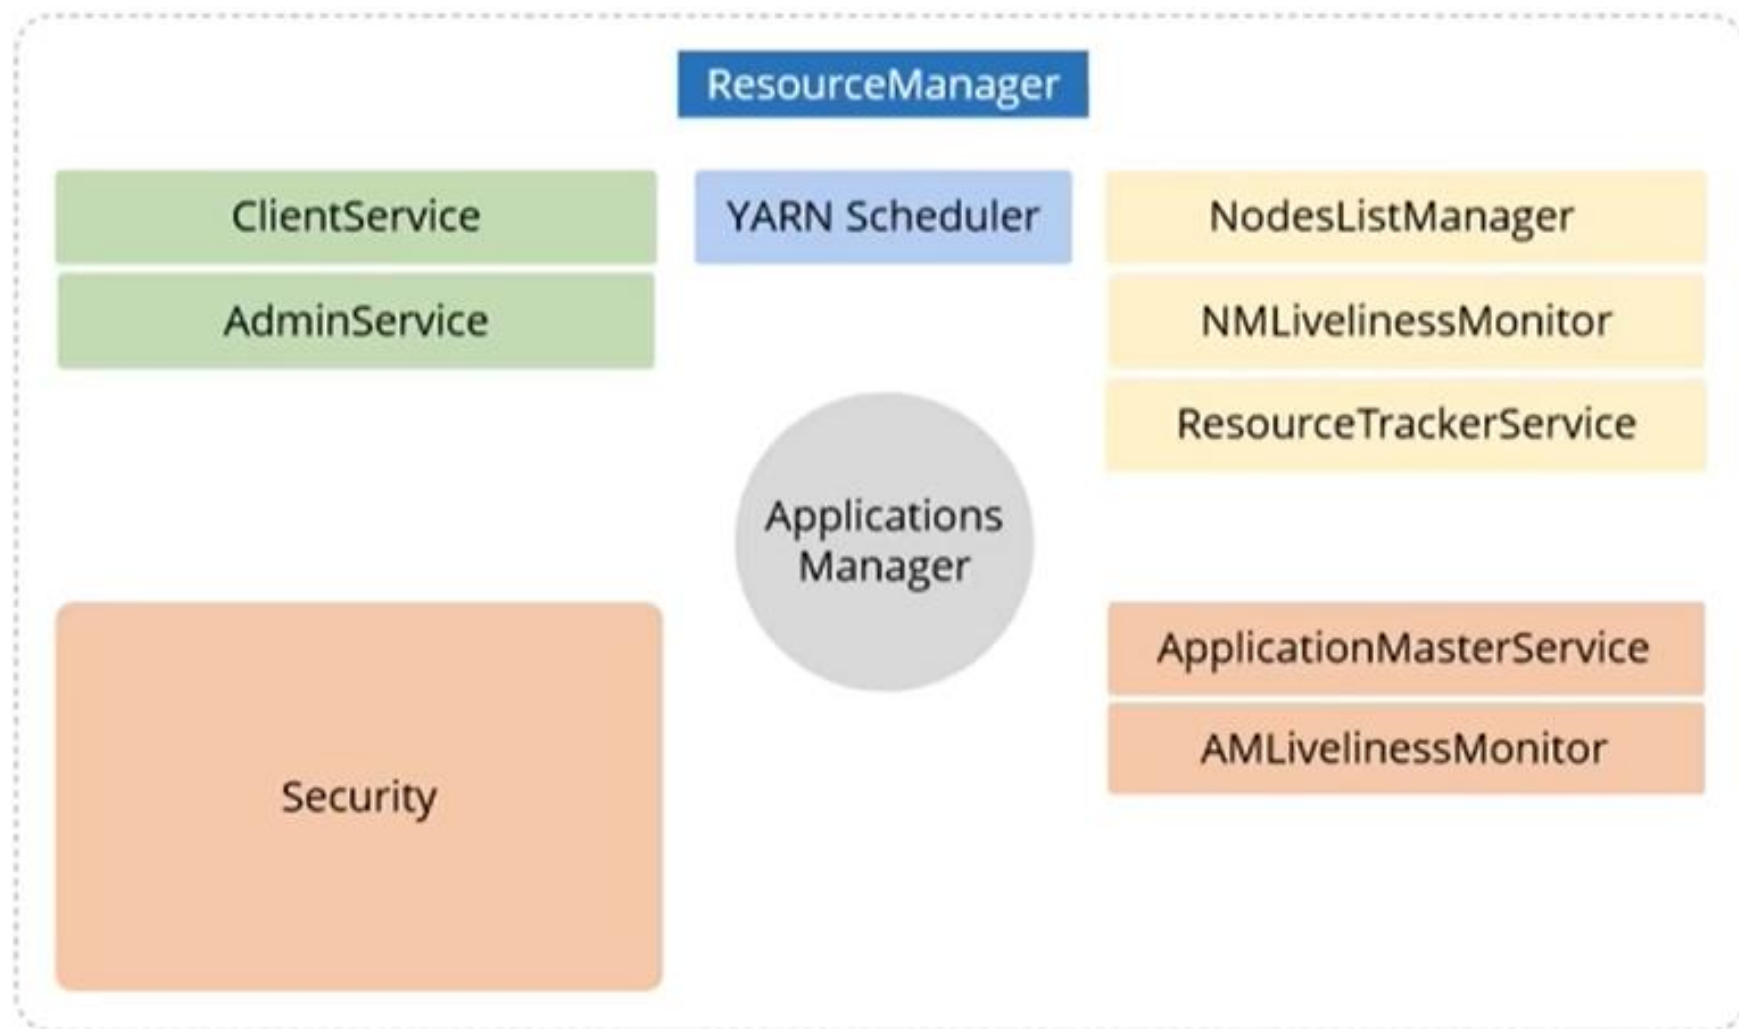
The ApplicationsManager is an interface which maintains a list of applications that have been submitted, currently running, or completed.



The ApplicationsManager accepts job submissions, negotiates the first container for executing the application, and restarts the ApplicationMaster container on failure.

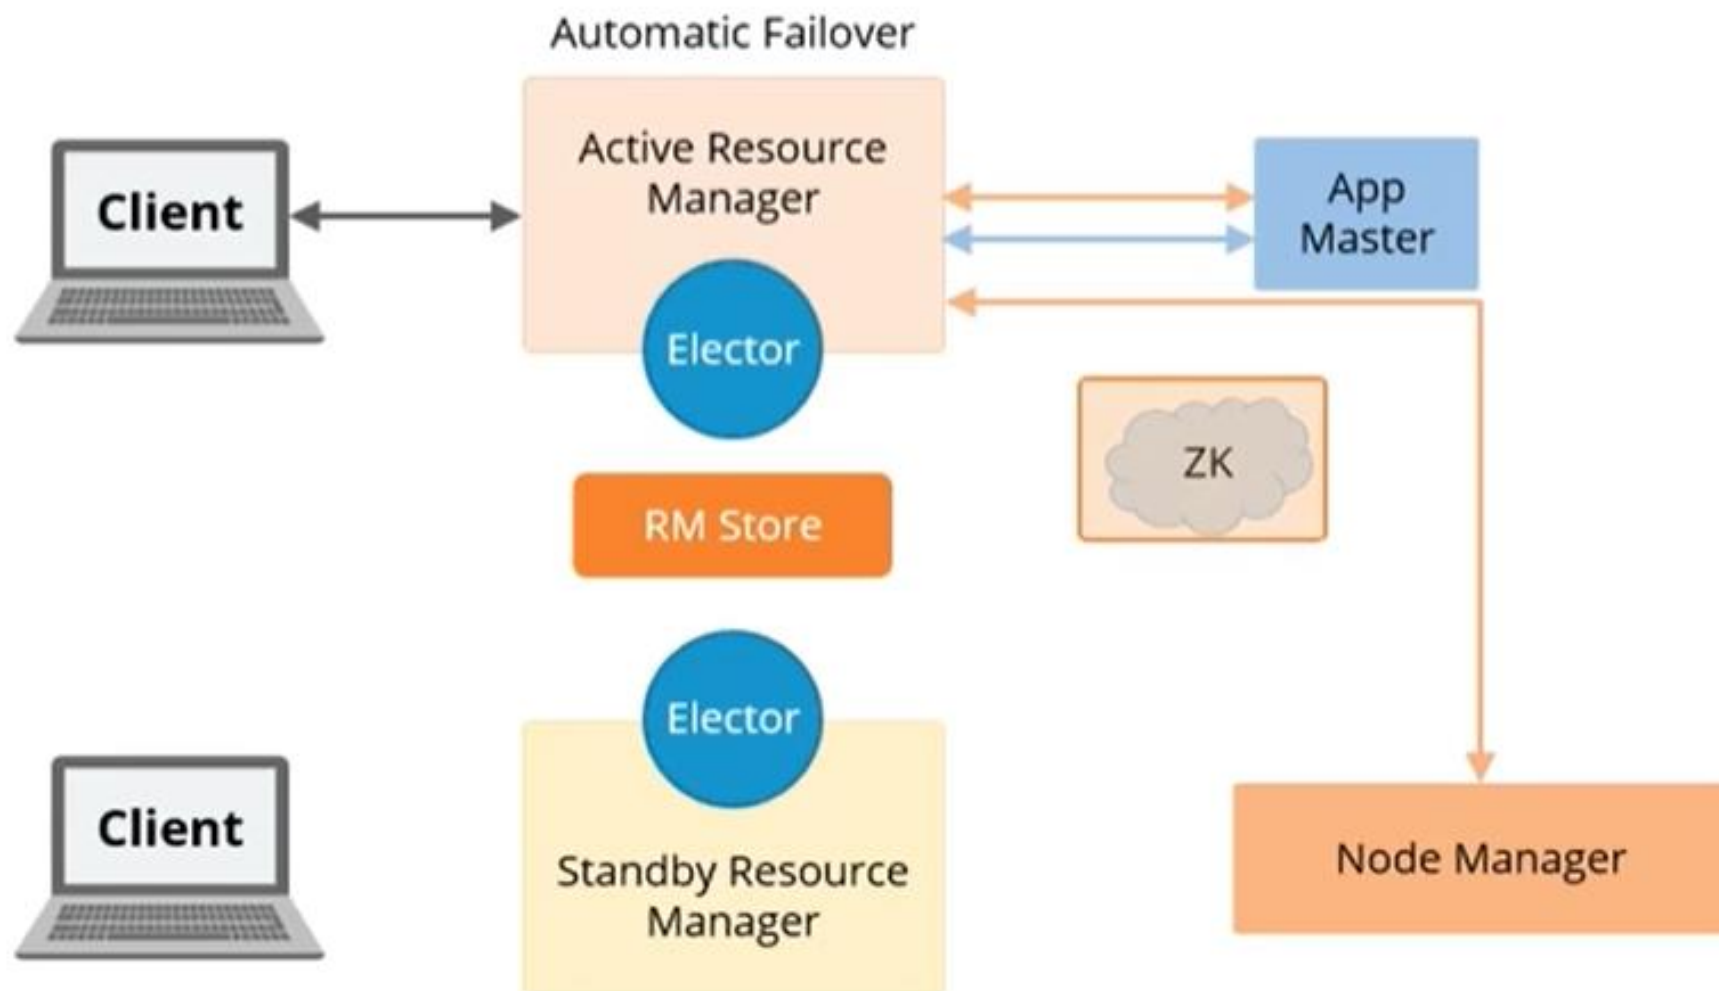
How a ResourceManager operates

The figure shown here displays all the internal components of the ResourceManager.




ResourceManager in High Availability Mode

Before Hadoop 2.4, the ResourceManager was the single point of failure in a YARN cluster. The High Availability, or HA, feature an Active/Standby ResourceManager pair to remove this single point of failure.



YARN Architecture Element—ApplicationMaster

The ApplicationMaster in YARN is a framework-specific library, which negotiates resources from the RM and works with the NodeManager or Managers to execute and monitor containers and their resource consumption.




The ApplicationMaster:

- manages the application lifecycle
- makes dynamic adjustments to resource consumption
- manages execution flow
- manages faults
- provides status and metrics to the RM
- interacts with NodeManager and RM using extensible communication protocols
- Is not run as a trusted service

While every application has its own instance of an AppMaster, it is possible to implement an AppMaster for a set of applications as well.

YARN Architecture Element—NodeManager

When a container is leased to an application, the NodeManager sets up the container's environment, including the resource constraints specified in the lease and any dependencies.




The NodeManager runs on each node and manages the following:

- Container lifecycle management
- Container dependencies
- Container leases
- Node and container resource usage
- Node health
- Log management
- Reporting node and container status to the RM

YARN Container

A YARN container is a result of a successful resource allocation, that is, the RM has granted an application a lease to use specified resources on a specific node.

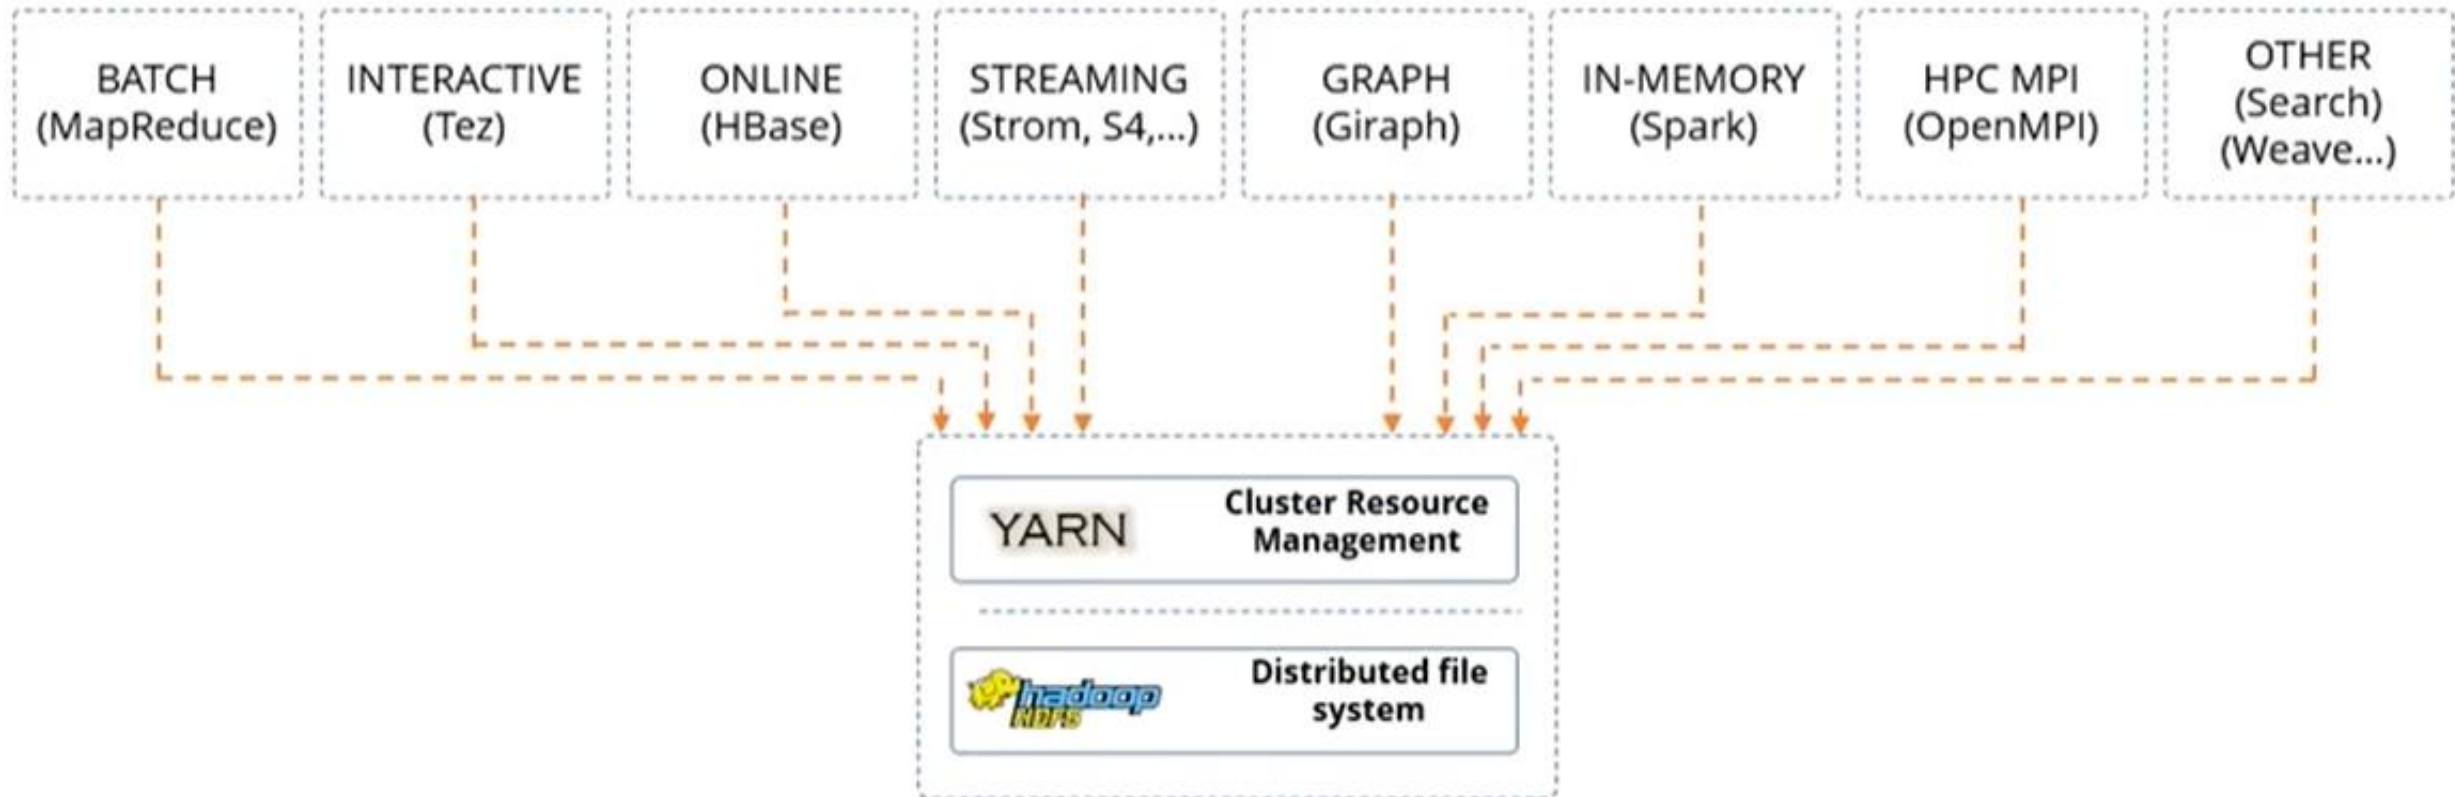


To launch the container, the ApplicationMaster must provide a container launch context (CLC) that includes the following information:

- Environment variables
- Dependencies, that is, local resources such as data files or shared objects needed prior to launch
- Security tokens
- The command necessary to create the process the application wants to launch






Applications on YARN

There can be many different workloads running on a Hadoop YARN cluster.



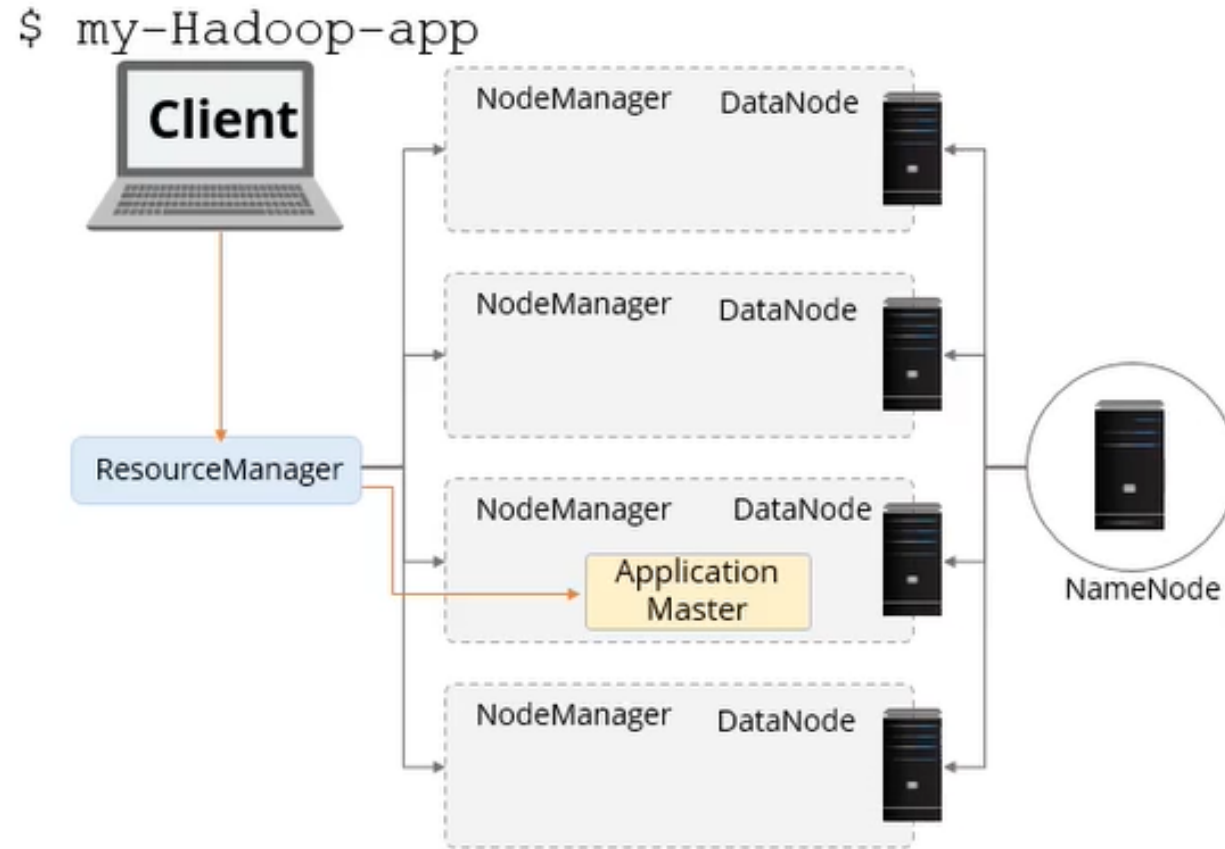
How YARN Runs an Application

There are five steps involved in YARN to run an application:

- 01 The client submits an application to the ResourceManager
- 02 The ResourceManager allocates a container
- 03 The ApplicationMaster contacts the related NodeManager
- 04 The NodeManager launches the container
- 05 The container executes the ApplicationMaster

Step1—Application Submitted to ResourceManager

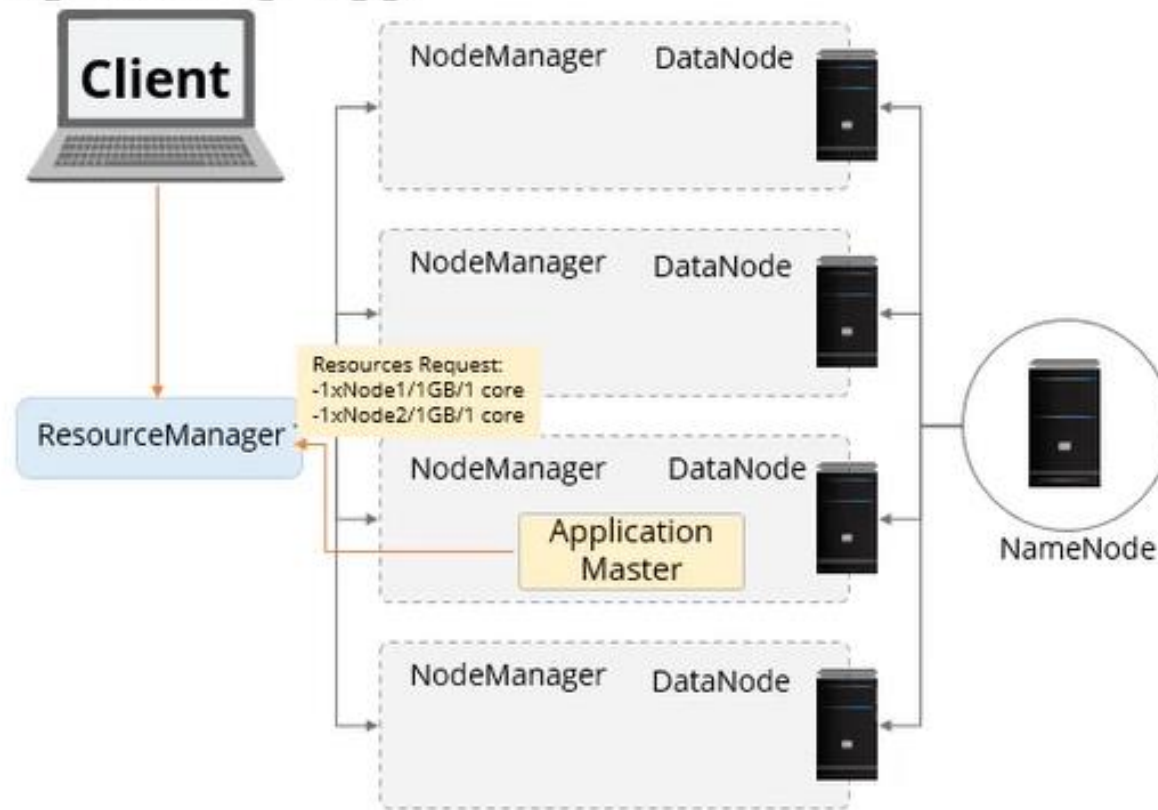
Users submit applications to the ResourceManager by typing the `hadoop jar` command.



Step2—ResourceManager Allocates a Container

When the ResourceManager accepts a new application submission, one of the first decisions the Scheduler makes is selecting a container.

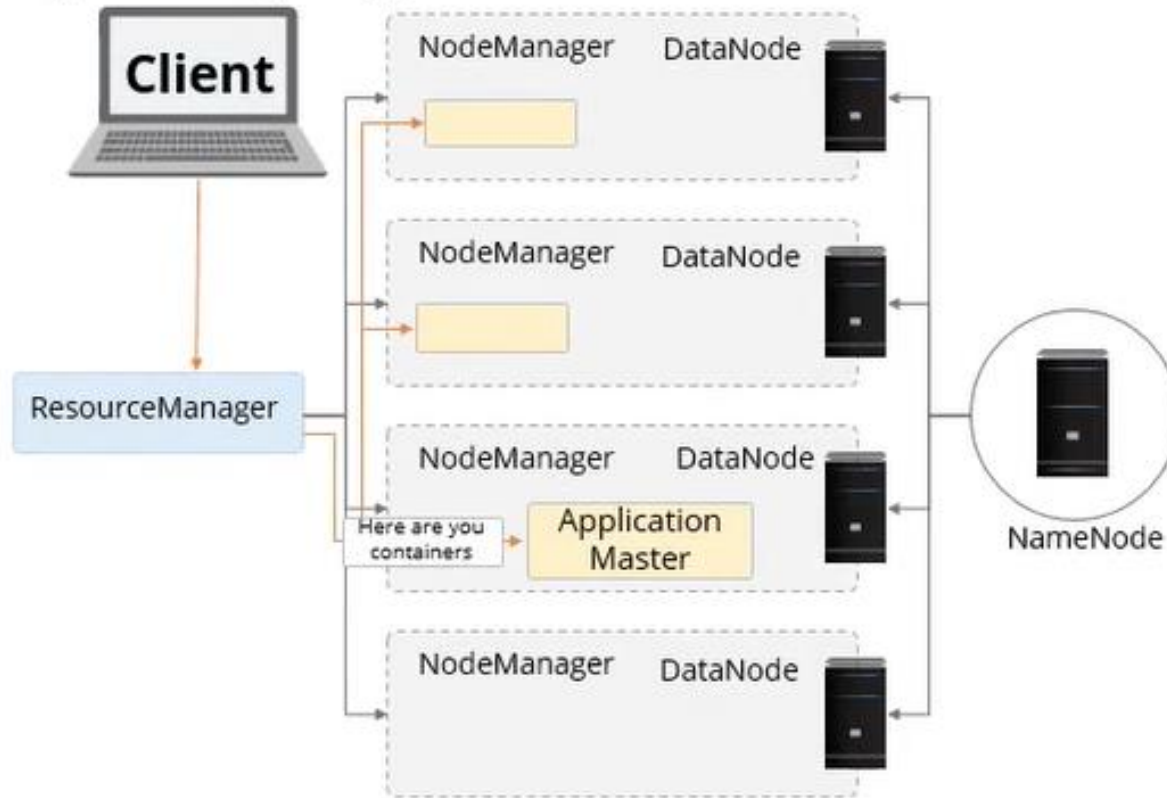
```
$ my-Hadoop-app
```



Step3 —ApplicationMaster Contacts NodeManager

After a container is allocated, the ApplicationMaster asks the NodeManager managing the host on which the container was allocated to use these resources to launch an application-specific task.

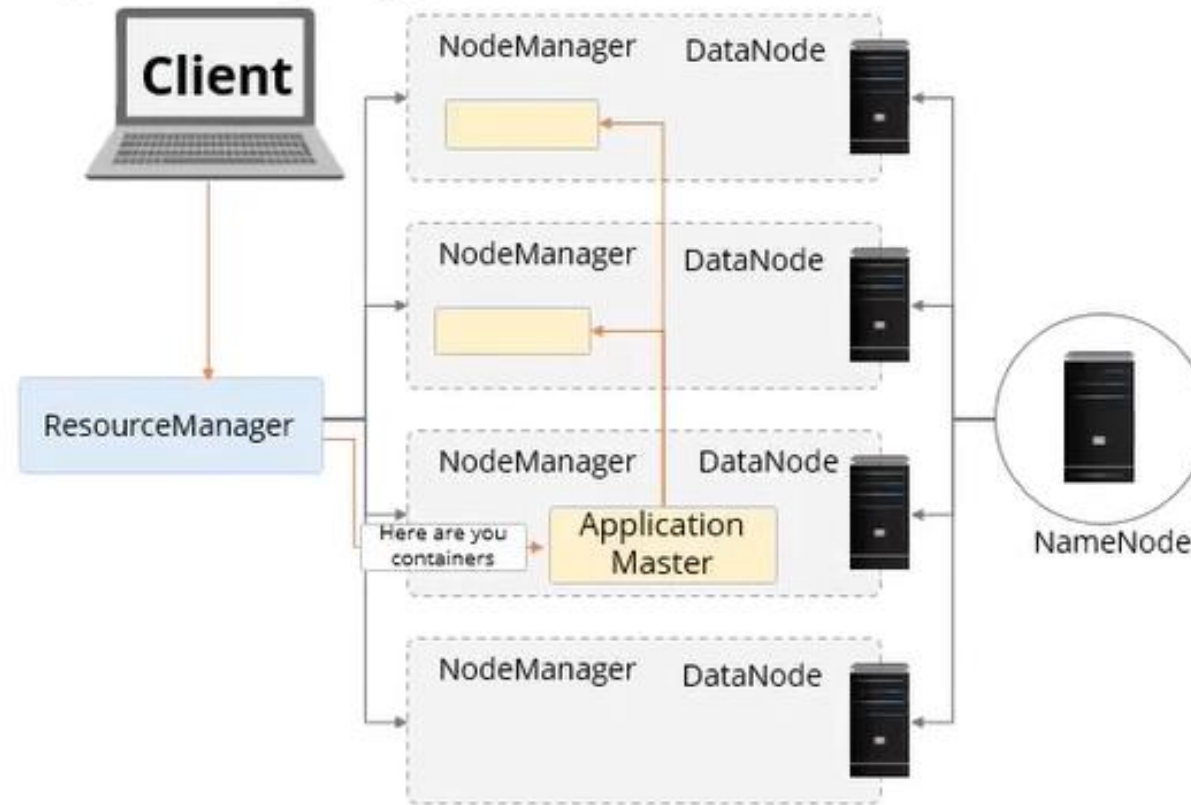
\$ my-Hadoop-app



Step4—ResourceManager Launches a Container

The NodeManager does not monitor tasks; it only monitors the resource usage in the containers.

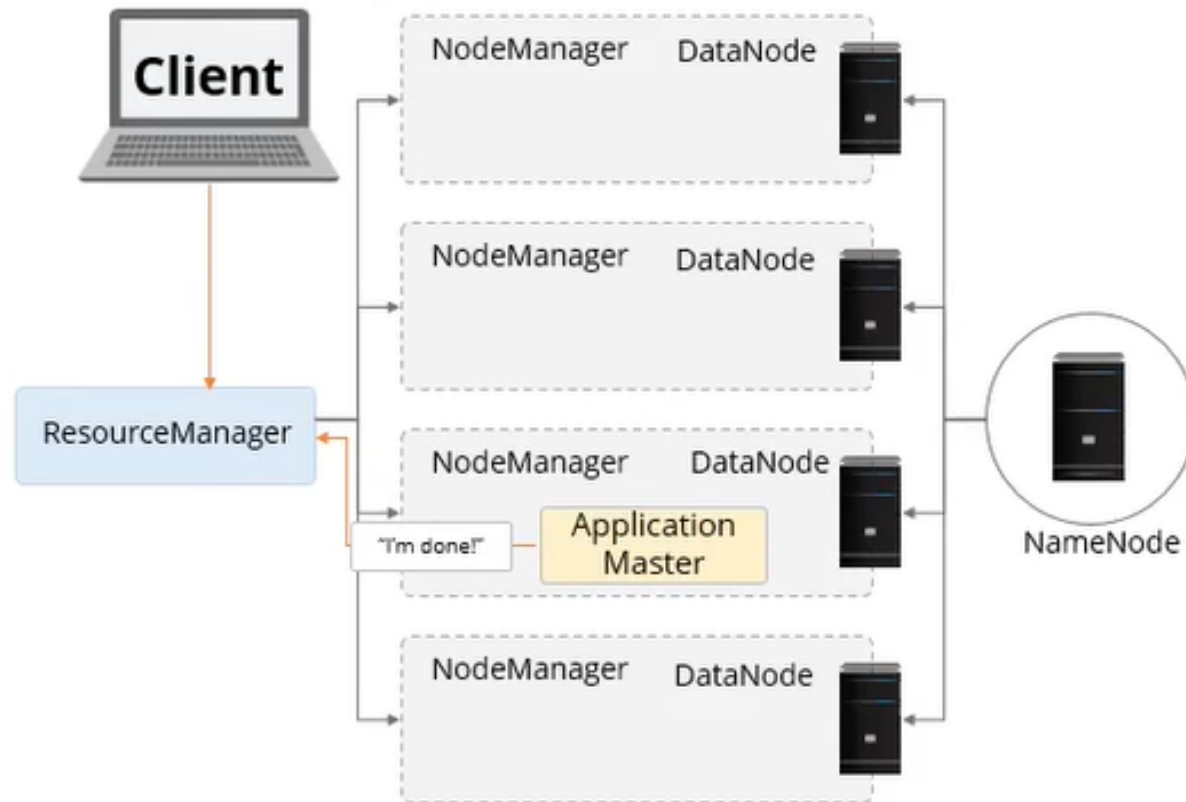
```
$ my-Hadoop-app
```



Step5—Container Executes the ApplicationMaster

After the application is complete, the ApplicationMaster shuts itself and releases its own container.

```
$ my-Hadoop-app
```



Three Tools for YARN Developers

Hadoop includes three tools for YARN developers:



YARN Web UI



YARN Web UI

YARN web UI runs on 8088 port, by default.

It also provides a better view than Hue; however you can't control or configure from YARN web UI.

Hue Job Browser

Hue Job
browser

The Hue Job Browser allows you to monitor status of job, kill a running job, and view logs.

The screenshot displays the Hue Job Browser interface. At the top, there's a navigation bar with 'HUE' logo and tabs for 'Query Editors', 'Data Browsers', and 'Workflows'. Below this, the 'Job Browser' title is visible. A search bar contains 'training' under 'Username' and 'Search for text' under 'Text'. To the right, there are status filters: 'Succeeded' (green), 'Running' (orange), 'Failed' (red), and 'Killed' (grey). The main area features a table of jobs:

Logs	ID	Name	Application Type	Status	User	Maps	Reduces	Queue	Priority	Duration	Submitted	
	1469722441471_0001	PythonWordCount	SPARK	RUNNING	training	0/0%	0/0%	root.training	N/A		07/28/16 09:24:06	Kill

At the bottom, it says 'Showing 1 to 1 of 1 entries' and has pagination controls: '- Previous', '1', 'Next -'.

YARN Command Line



Most of the YARN commands are for administrator rather than developer

Few useful commands for developer:

- yarn -help
list all command of yarn

- yarn -version
print the version

- yarn logs -applicationId <app-id>
views logs of specified application ID



Running MapReduce Examples on Hadoop YARN

Run Cmd as administrator

```
hdfs namenode -format
```

Run Cmd as administrator

```
cd\  
cd hadoop\sbin  
start-dfs  
start-yarn
```

Run Cmd as administrator

```
cd\  
cd hadoop\share\hadoop\mapreduce  
yarn jar hadoop-mapreduce-examples-3.3.1.jar
```

Running MapReduce Examples on Hadoop YARN

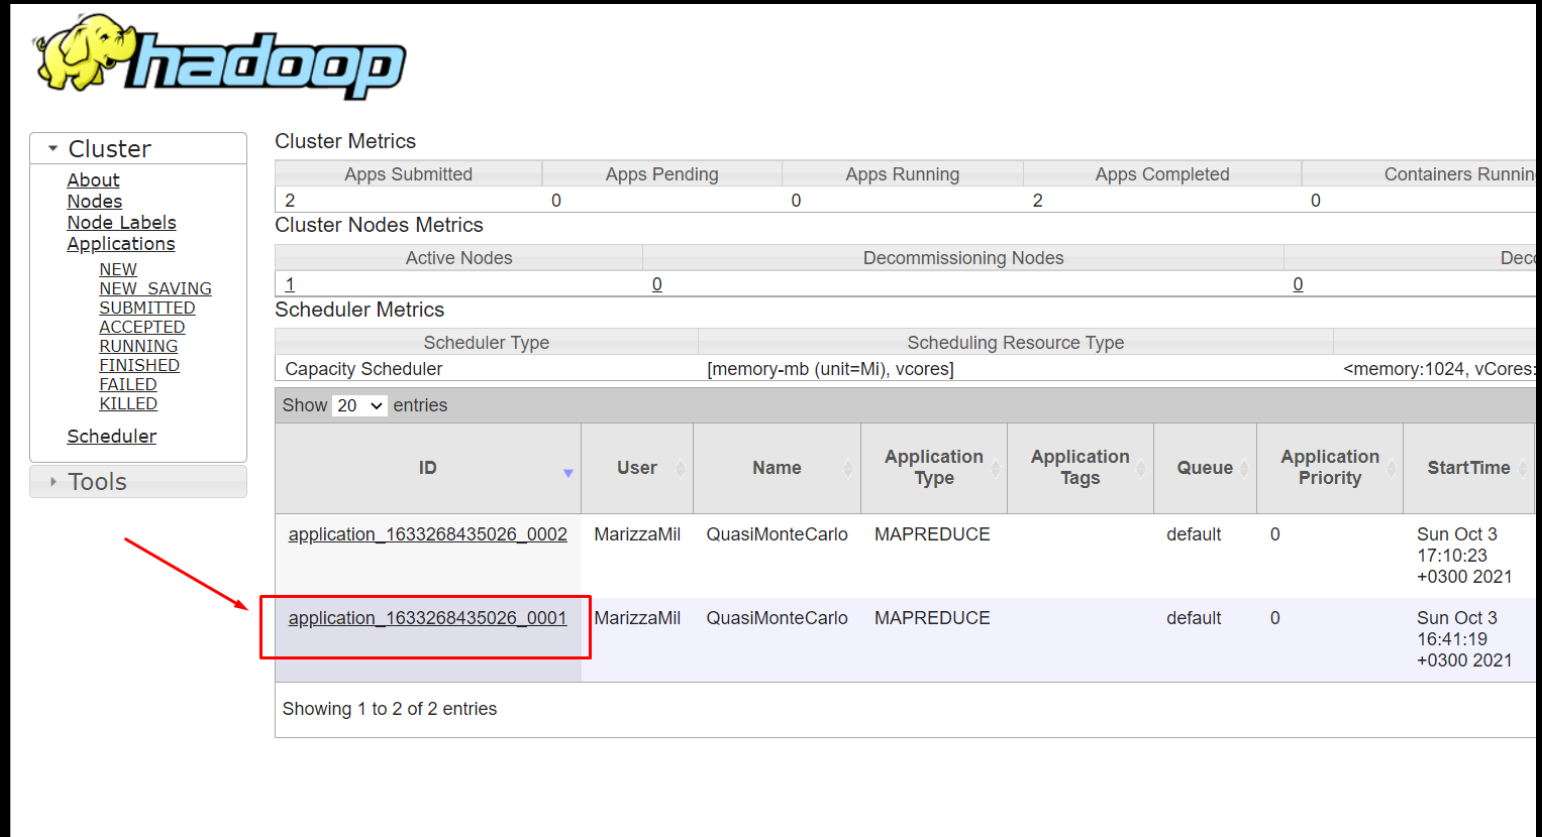
Running the pi Example

To run the pi example with 16 maps and 100000 samples, run the following command:
Run Cmd as administrator

```
cd\  
cd hadoop\share\hadoop\mapreduce  
yarn jar hadoop-mapreduce-examples-3.3.1.jar pi 16 100000
```


Using the Web GUI to Monitor Examples

The following figure shows the main YARN web interface (<http://hostname:8088>)



The screenshot displays the Hadoop YARN web interface. On the left is a navigation menu with a 'Cluster' dropdown containing links for 'About', 'Nodes', 'Node Labels', 'Applications' (with sub-links for NEW, NEW_SAVING, SUBMITTED, ACCEPTED, RUNNING, FINISHED, FAILED, KILLED), and 'Scheduler'. Below this is a 'Tools' button. A red arrow points from the 'Applications' link to a table of running applications. The table has columns for ID, User, Name, Application Type, Application Tags, Queue, Application Priority, and StartTime. Two applications are listed, with the second one highlighted by a red box. Above the table are sections for 'Cluster Metrics' and 'Cluster Nodes Metrics'. The 'Scheduler Metrics' section shows 'Capacity Scheduler' as the type and '[memory-mb (unit=M), vcores]' as the resource type.

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running
2	0	0	2	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes
1	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type
Capacity Scheduler	[memory-mb (unit=M), vcores]

Showing 20 entries

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime
application_1633268435026_0002	MarizzaMil	QuasiMonteCarlo	MAPREDUCE		default	0	Sun Oct 3 17:10:23 +0300 2021
application_1633268435026_0001	MarizzaMil	QuasiMonteCarlo	MAPREDUCE		default	0	Sun Oct 3 16:41:19 +0300 2021

Showing 1 to 2 of 2 entries

This page provides information similar to that on the Running Applications page, but only for the selected job.



Logged in as: dr.who

Application application_1633268435026_0001

Cluster

[About](#)
[Nodes](#)
[Node Labels](#)
[Applications](#)
 [NEW](#)
 [NEW SAVING](#)
 [SUBMITTED](#)
 [ACCEPTED](#)
 [RUNNING](#)
 [FINISHED](#)
 [FAILED](#)
 [KILLED](#)
[Scheduler](#)

Tools

Application Overview

User: [MarizzaMil](#)
Name: QuasiMonteCarlo
Application Type: MAPREDUCE
Application Tags:
Application Priority: 0 (Higher Integer value indicates higher priority)
YarnApplicationState: FINISHED
Queue: [default](#)
FinalStatus Reported by AM: SUCCEEDED
Started: Sun Oct 03 16:41:19 +0300 2021
Launched: Sun Oct 03 16:41:20 +0300 2021
Finished: Sun Oct 03 16:42:02 +0300 2021
Elapsed: 43sec
Tracking URL: [History](#)
Log Aggregation Status: DISABLED
Application Timeout (Remaining Time): Unlimited
Diagnostics:
Unmanaged Application: false
Application Node Label expression: <Not set>
AM container Node Label expression: <DEFAULT_PARTITION>

Application Metrics

Total Resource Preempted: <memory:0, vCores:0>
Total Number of Non-AM Containers Preempted: 0
Total Number of AM Containers Preempted: 0
Resource Preempted from Current Attempt: <memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt: 0
Aggregate Resource Allocation: 212557 MB-seconds, 149 vcore-seconds
Aggregate Preempted Resource Allocation: 0 MB-seconds, 0 vcore-seconds

Show 20 entries

Search:

Attempt ID	Started	Node	Logs	Nodes blacklisted by the app	Nodes blacklisted by the system
appattempt_1633268435026_0001_000001	Sun Oct 3 16:41:19 +0300 2021	http://DESKTOP-KJ7307V:8042	Logs	0	0

Run wordcount MapReduce job

Run Cmd as administrator

```
hdfs namenode -format
```

Run Cmd as administrator

```
cd\  
cd hadoop\sbin  
start-dfs  
start-yarn
```

// 1. Create a text file with some content. We'll pass this file as input to the **wordcount** MapReduce job for counting words.

C:\file1.txt

// 2. Create a directory ('new_demo') in HDFS to keep all the text files ('file1.txt') to be used for counting words.

```
hadoop fs -mkdir /new_demo
```

// 3. Copy the text file(say 'file1.txt') from local disk to the newly created 'new_demo' directory in HDFS.

```
hdfs dfs -copyFromLocal C:\temp\file1.txt /new_demo
```

// 4. Check content of the copied file

```
hdfs dfs -ls /new_demo
```

Run Cmd as administrator

```
cd\  
cd hadoop\share\hadoop\mapreduce
```

// 5. Run the **wordcount** MapReduce job provided

```
yarn jar hadoop-mapreduce-examples-3.3.1.jar wordcount /new_demo output
```

// 6. Check output.

```
hdfs dfs -cat output/*
```

http://hostname:8088



FINISHED Applications

Cluster

[About](#)
[Nodes](#)
[Node Labels](#)
[Applications](#)

[NEW](#)
[NEW SAVING](#)
[SUBMITTED](#)
[ACCEPTED](#)
[RUNNING](#)
[FINISHED](#)
[FAILED](#)
[KILLED](#)

[Scheduler](#)

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Used Resources	Total Resources
3	0	0	3	0	<memory:0 B, vCores:0>	<memory:8 GB, vCores:4>

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes
1	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[memory-mb (unit=M), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>

Show 20 entries

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers
application_1633268435026_0004	MarizzaMil	word count	MAPREDUCE		default	0	Sun Oct 3 17:38:02 +0300 2021	Sun Oct 3 17:38:02 +0300 2021	Sun Oct 3 17:38:30 +0300 2021	FINISHED	SUCCEEDED	N/A
application_1633268435026_0002	MarizzaMil	QuasiMonteCarlo	MAPREDUCE		default	0	Sun Oct 3 17:10:23 +0300 2021	Sun Oct 3 17:10:23 +0300 2021	Sun Oct 3 17:11:07 +0300 2021	FINISHED	SUCCEEDED	N/A
application_1633268435026_0001	MarizzaMil	QuasiMonteCarlo	MAPREDUCE		default	0	Sun Oct 3 16:41:19 +0300 2021	Sun Oct 3 16:41:20 +0300 2021	Sun Oct 3 16:42:02 +0300 2021	FINISHED	SUCCEEDED	N/A

Showing 1 to 3 of 3 entries